



PANDAS FOUNDATIONS

pandas Foundations



What is pandas?

- Python library for data analysis
- High-performance containers for data analysis
- Data structures with a lot of functionality
 - Meaningful labels
 - Time series functionality
 - Handling missing data
 - Relational operations

What you will learn

- How to work with pandas
 - Data import & export in various formats
- Exploratory Data Analysis using pandas
 - Statistical & graphical methods
- Using pandas to model *time series*
 - Time indexes, resampling



PANDAS FOUNDATIONS

**See you in
the course!**



PANDAS FOUNDATIONS

Review of pandas DataFrames



pandas DataFrames

- Example: DataFrame of Apple Stock data

Date	Open	High	Low	Close	Volume	Adj Close
2014-09-16	99.80	101.26	98.89	100.86	66818200	100.86
2014-09-15	102.81	103.05	101.44	101.63	61216500	101.63
2014-09-12	101.21	102.19	101.08	101.66	62626100	101.66
...



Indexes and columns

```
In [1]: import pandas as pd
```

```
In [2]: type(AAPL)
```

```
Out[2]: pandas.core.frame.DataFrame
```

```
In [3]: AAPL.shape
```

```
Out[3]: (8514, 6)
```

```
In [4]: AAPL.columns
```

```
Out[4]:
```

```
Index(['Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close'],  
      dtype='object')
```

```
In [5]: type(AAPL.columns)
```

```
Out[5]: pandas.indexes.base.Index
```



Indexes and columns

```
In [6]: AAPL.index
```

```
Out[6]:
```

```
DatetimeIndex(['2014-09-16', '2014-09-15', '2014-09-12',  
               '2014-09-11', '2014-09-10', '2014-09-09',  
               '2014-09-08', '2014-09-05', '2014-09-04',  
               '2014-09-03',  
               ...  
               '1980-12-26', '1980-12-24', '1980-12-23',  
               '1980-12-22', '1980-12-19', '1980-12-18',  
               '1980-12-17', '1980-12-16', '1980-12-15',  
               '1980-12-12'],  
              dtype='datetime64[ns]', name='Date', length=8514,  
              freq=None)
```

```
In [7]: type(AAPL.index)
```

```
Out[7]: pandas.tseries.index.DatetimeIndex
```




Slicing

```
In [8]: AAPL.iloc[:5,:]
```

```
Out[8]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2014-09-16	99.80	101.26	98.89	100.86	66818200	100.86
2014-09-15	102.81	103.05	101.44	101.63	61216500	101.63
2014-09-12	101.21	102.19	101.08	101.66	62626100	101.66
2014-09-11	100.41	101.44	99.62	101.43	62353100	101.43
2014-09-10	98.01	101.11	97.76	101.00	100741900	101.00

```
In [9]: AAPL.iloc[-5:,:]
```

```
Out[9]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
1980-12-18	26.63	26.75	26.63	26.63	18362400	0.41
1980-12-17	25.87	26.00	25.87	25.87	21610400	0.40
1980-12-16	25.37	25.37	25.25	25.25	26432000	0.39
1980-12-15	27.38	27.38	27.25	27.25	43971200	0.42
1980-12-12	28.75	28.87	28.75	28.75	117258400	0.45



head()

```
In [10]: AAPL.head(5)
```

```
Out[10]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2014-09-16	99.80	101.26	98.89	100.86	66818200	100.86
2014-09-15	102.81	103.05	101.44	101.63	61216500	101.63
2014-09-12	101.21	102.19	101.08	101.66	62626100	101.66
2014-09-11	100.41	101.44	99.62	101.43	62353100	101.43
2014-09-10	98.01	101.11	97.76	101.00	100741900	101.00

```
In [11]: AAPL.head(2)
```

```
Out[11]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2014-09-16	99.80	101.26	98.89	100.86	66818200	100.86
2014-09-15	102.81	103.05	101.44	101.63	61216500	101.63



tail()

```
In [12]: AAPL.tail()
```

```
Out[12]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
1980-12-18	26.63	26.75	26.63	26.63	18362400	0.41
1980-12-17	25.87	26.00	25.87	25.87	21610400	0.40
1980-12-16	25.37	25.37	25.25	25.25	26432000	0.39
1980-12-15	27.38	27.38	27.25	27.25	43971200	0.42
1980-12-12	28.75	28.87	28.75	28.75	117258400	0.45

```
In [13]: AAPL.tail(3)
```

```
Out[13]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
1980-12-16	25.37	25.37	25.25	25.25	26432000	0.39
1980-12-15	27.38	27.38	27.25	27.25	43971200	0.42
1980-12-12	28.75	28.87	28.75	28.75	117258400	0.45



info()

```
In [14]: AAPL.info()
Out[14]:
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8514 entries, 2014-09-16 to 1980-12-12
Data columns (total 6 columns):
Open           8514 non-null float64
High           8514 non-null float64
Low            8514 non-null float64
Close          8514 non-null float64
Volume         8514 non-null int64
Adj Close      8514 non-null float64
dtypes: float64(5), int64(1)
memory usage: 465.6 KB
```



Broadcasting

```
In [15]: import numpy as np
```

```
In [16]: AAPL.iloc[:, 3, -1] = np.nan
```

← Assigning scalar value to column slice *broadcasts* value to each row.

```
In [17]: AAPL.head(6)
```

```
Out[17]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2014-09-16	99.80	101.26	98.89	100.86	66818200	NaN
2014-09-15	102.81	103.05	101.44	101.63	61216500	101.63
2014-09-12	101.21	102.19	101.08	101.66	62626100	101.66
2014-09-11	100.41	101.44	99.62	101.43	62353100	NaN
2014-09-10	98.01	101.11	97.76	101.00	100741900	101.00
2014-09-09	99.08	103.08	96.14	97.99	189560600	97.99
2014-09-08	99.30	99.31	98.05	98.36	46277800	NaN



Broadcasting

```
In [18]: AAPL.info()
Out[18]:
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8514 entries, 2014-09-16 to 1980-12-12
Data columns (total 6 columns):
Open                8514 non-null float64
High                8514 non-null float64
Low                 8514 non-null float64
Close               8514 non-null float64
Volume              8514 non-null int64
Adj Close           5676 non-null float64
dtypes: float64(5), int64(1)
memory usage: 465.6 KB
```



Series

```
In [19]: low = AAPL['Low']
```

```
In [20]: type(low)
```

```
Out[20]: pandas.core.series.Series
```

```
In [21]: low.head()
```

```
Out[21]:
```

Date

2014-09-16	98.89
------------	-------

2014-09-15	101.44
------------	--------

2014-09-12	101.08
------------	--------

2014-09-11	99.62
------------	-------

2014-09-10	97.76
------------	-------

Name: Low, dtype: float64

```
In [22]: lows = low.values
```

```
In [23]: type(lows)
```

```
Out[23]: numpy.ndarray
```



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Building DataFrames from scratch



DataFrames from CSV files

```
In [1]: import pandas as pd
```

```
In [2]: users = pd.read_csv('datasets/users.csv', index_col=0)
```

```
In [3]: print(users)
```

```
Out[3]:
```

	weekday	city	visitors	signups
0	Sun	Austin	139	7
1	Sun	Dallas	237	12
2	Mon	Austin	326	3
3	Mon	Dallas	456	5



DataFrames from dict (1)

```
In [1]: import pandas as pd
```

```
In [2]: data = {'weekday': ['Sun', 'Sun', 'Mon', 'Mon'],  
...:           'city': ['Austin', 'Dallas', 'Austin', 'Dallas'],  
...:           'visitors': [139, 237, 326, 456],  
...:           'signups': [7, 12, 3, 5]}
```

```
In [3]: users = pd.DataFrame(data)
```

```
In [4]: print(users)
```

```
Out[4]:
```

	weekday	city	visitors	signups
0	Sun	Austin	139	7
1	Sun	Dallas	237	12
2	Mon	Austin	326	3
3	Mon	Dallas	456	5



DataFrames from dict (2)

```
In [1]: import pandas as pd
```

```
In [2]: cities = ['Austin', 'Dallas', 'Austin', 'Dallas']
```

```
In [3]: signups = [7, 12, 3, 5]
```

```
In [4]: visitors = [139, 237, 326, 456]
```

```
In [5]: weekdays = ['Sun', 'Sun', 'Mon', 'Mon']
```

```
In [6]: list_labels = ['city', 'signups', 'visitors', 'weekday']
```

```
In [7]: list_cols = [cities, signups, visitors, weekdays]
```

```
In [8]: zipped = list(zip(list_labels, list_cols))
```



DataFrames from dict (3)

```
In [9]: print(zipped)
```

```
Out[9]:
```

```
[('city', ['Austin', 'Dallas', 'Austin', 'Dallas']), ('signups',  
[7, 12, 3, 5]), ('visitors', [139, 237, 326, 456]), ('weekday',  
['Sun', 'Sun', 'Mon', 'Mon'])]
```

```
In [10]: data = dict(zipped)
```

```
In [11]: users = pd.DataFrame(data)
```

```
In [12]: print(users)
```

```
Out[12]:
```

	weekday	city	visitors	signups
0	Sun	Austin	139	7
1	Sun	Dallas	237	12
2	Mon	Austin	326	3
3	Mon	Dallas	456	5



Broadcasting

```
In [13]: users['fees'] = 0 # Broadcasts to entire column
```

```
In [14]: print(users)
```

```
Out[14]:
```

	city	signups	visitors	weekday	fees
0	Austin	7	139	Sun	0
1	Dallas	12	237	Sun	0
2	Austin	3	326	Mon	0
3	Dallas	5	456	Mon	0



Broadcasting with a dict

```
In [1]: import pandas as pd
```

```
In [2]: heights = [ 59.0, 65.2, 62.9, 65.4, 63.7, 65.7, 64.1 ]
```

```
In [3]: data = {'height': heights, 'sex': 'M'}
```

```
In [4]: results = pd.DataFrame(data)
```

```
In [5]: print(results)
```

```
Out[5]:
```

	height	sex
0	59.0	M
1	65.2	M
2	62.9	M
3	65.4	M
4	63.7	M
5	65.7	M
6	64.1	M



Index and columns

```
In [6]: results.columns = ['height (in)', 'sex']
```

```
In [7]: results.index = ['A', 'B', 'C', 'D', 'E', 'F', 'G']
```

```
In [8]: print(results)
```

```
Out[8]:
```

	height (in)	sex
A	59.0	M
B	65.2	M
C	62.9	M
D	65.4	M
E	63.7	M
F	65.7	M
G	64.1	M



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Importing & exporting data



Original CSV file

- Dataset: Sunspot observations collected from SILSO

```
1818,01,01,1818.004, -1,1
1818,01,02,1818.007, -1,1
1818,01,03,1818.010, -1,1
1818,01,04,1818.012, -1,1
1818,01,05,1818.015, -1,1
1818,01,06,1818.018, -1,1
...
```



Datasets from CSV files

```
In [1]: import pandas as pd
```

```
In [2]: filepath = 'ISSN_D_tot.csv'
```

```
In [3]: sunspots = pd.read_csv(filepath)
```

```
In [4]: sunspots.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 71921 entries, 0 to 71920
```

```
Data columns (total 6 columns):
```

1818	71921 non-null	int64
01	71921 non-null	int64
01.1	71921 non-null	int64
1818.004	71921 non-null	float64
-1	71921 non-null	int64
1	71921 non-null	int64

```
dtypes: float64(1), int64(5)
```

```
memory usage: 3.3 MB
```



Datasets from CSV files

```
In [5]: sunspots.iloc[10:20, :]
```

```
Out[5]:
```

	1818	01	01.1	1818.004	-1	1
10	1818	1	12	1818.034	-1	1
11	1818	1	13	1818.037	22	1
12	1818	1	14	1818.040	-1	1
13	1818	1	15	1818.042	-1	1
14	1818	1	16	1818.045	-1	1
15	1818	1	17	1818.048	46	1
16	1818	1	18	1818.051	59	1
17	1818	1	19	1818.053	63	1
18	1818	1	20	1818.056	-1	1
19	1818	1	21	1818.059	-1	1

Problems

- CSV file has no column headers
 - Columns 0-2: Gregorian date (year, month, day)
 - Column 3: Date as fraction as year
 - Column 4: Daily total sunspot number
 - Column 5: Definitive/provisional indicator (1 or 0)
- Missing values in column 4: indicated by -1
- Dates representation inconvenient



Using header keyword

```
In [6]: sunspots = pd.read_csv(filepath, header=None)
```

```
In [7]: sunspots.iloc[10:20, :]
```

```
Out[7]:
```

	0	1	2	3	4	5
10	1818	1	11	1818.031	-1	1
11	1818	1	12	1818.034	-1	1
12	1818	1	13	1818.037	22	1
13	1818	1	14	1818.040	-1	1
14	1818	1	15	1818.042	-1	1
15	1818	1	16	1818.045	-1	1
16	1818	1	17	1818.048	46	1
17	1818	1	18	1818.051	59	1
18	1818	1	19	1818.053	63	1
19	1818	1	20	1818.056	-1	1



Using names keyword

```
In [8]: col_names = ['year', 'month', 'day', 'dec_date',  
....:                  'sunspots', 'definite']
```

```
In [9]: sunspots = pd.read_csv(filepath, header=None,  
....:                           names=col_names)
```

```
In [10]: sunspots.iloc[10:20, :]
```

```
Out[10]:
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	-1	1
11	1818	1	12	1818.034	-1	1
12	1818	1	13	1818.037	22	1
13	1818	1	14	1818.040	-1	1
14	1818	1	15	1818.042	-1	1
15	1818	1	16	1818.045	-1	1
16	1818	1	17	1818.048	46	1
17	1818	1	18	1818.051	59	1
18	1818	1	19	1818.053	63	1
19	1818	1	20	1818.056	-1	1



Using na_values keyword (1)

```
In [11]: sunspots = pd.read_csv(filepath, header=None,  
    ....:                        names=col_names, na_values='-1')
```

```
In [12]: sunspots.iloc[10:20, :]
```

```
Out[12]:
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	-1	1
11	1818	1	12	1818.034	-1	1
12	1818	1	13	1818.037	22	1
13	1818	1	14	1818.040	-1	1
14	1818	1	15	1818.042	-1	1
15	1818	1	16	1818.045	-1	1
16	1818	1	17	1818.048	46	1
17	1818	1	18	1818.051	59	1
18	1818	1	19	1818.053	63	1
19	1818	1	20	1818.056	-1	1



Using na_values keyword (2)

```
In [13]: sunspots = pd.read_csv(filepath, header=None,  
    ....:                        names=col_names, na_values=' -1')
```

```
In [14]: sunspots.iloc[10:20, :]
```

```
Out[14]:
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	NaN	1
11	1818	1	12	1818.034	NaN	1
12	1818	1	13	1818.037	22.0	1
13	1818	1	14	1818.040	NaN	1
14	1818	1	15	1818.042	NaN	1
15	1818	1	16	1818.045	NaN	1
16	1818	1	17	1818.048	46.0	1
17	1818	1	18	1818.051	59.0	1
18	1818	1	19	1818.053	63.0	1
19	1818	1	20	1818.056	NaN	1



Using na_values keyword (3)

```
In [15]: sunspots = pd.read_csv(filepath, header=None,  
    ...: names=col_names, na_values={'sunspots': ['-1']})
```

```
In [16]: sunspots.iloc[10:20, :]
```

```
Out[16]:
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	NaN	1
11	1818	1	12	1818.034	NaN	1
12	1818	1	13	1818.037	22.0	1
13	1818	1	14	1818.040	NaN	1
14	1818	1	15	1818.042	NaN	1
15	1818	1	16	1818.045	NaN	1
16	1818	1	17	1818.048	46.0	1
17	1818	1	18	1818.051	59.0	1
18	1818	1	19	1818.053	63.0	1
19	1818	1	20	1818.056	NaN	1



Using parse_dates keyword

```
In [17]: sunspots = pd.read_csv(filepath, header=None,  
...: names=col_names, na_values={'sunspots': ['-1']},  
...: parse_dates=[[0, 1, 2]])
```

```
In [18]: sunspots.iloc[10:20, :]
```

```
Out[18]:
```

	year_month_day	dec_date	sunspots	definite
10	1818-01-11	1818.031	NaN	1
11	1818-01-12	1818.034	NaN	1
12	1818-01-13	1818.037	22.0	1
13	1818-01-14	1818.040	NaN	1
14	1818-01-15	1818.042	NaN	1
15	1818-01-16	1818.045	NaN	1
16	1818-01-17	1818.048	46.0	1
17	1818-01-18	1818.051	59.0	1
18	1818-01-19	1818.053	63.0	1
19	1818-01-20	1818.056	NaN	1



Inspecting DataFrame

```
In [19]: sunspots.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 71922 entries, 0 to 71921
Data columns (total 4 columns):
year_month_day      71922 non-null datetime64[ns]
dec_date            71922 non-null float64
sunspots            68675 non-null float64
definite            71922 non-null int64
dtypes: datetime64[ns](1), float64(2), int64(1)
memory usage: 2.2 MB
```



Using dates as index

```
In [20]: sunspots.index = sunspots['year_month_day']
```

```
In [21]: sunspots.index.name = 'date'
```

```
In [22]: sunspots.info()
```

```
Out[22]:
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
DatetimeIndex: 71922 entries, 1818-01-01 to 2014-11-30
```

```
Data columns (total 4 columns):
```

```
year_month_day    71922 non-null datetime64[ns]
```

```
dec_date          71922 non-null float64
```

```
sunspots          68675 non-null float64
```

```
definite          71922 non-null int64
```

```
dtypes: datetime64[ns](1), float64(2), int64(1)
```

```
memory usage: 2.7 MB
```



Trimming redundant columns

```
In [23]: cols = ['sunspots', 'definite']
```

```
In [24]: sunspots = sunspots[cols]
```

```
In [25]: sunspots.iloc[10:20, :]
```

```
Out[25]:
```

	sunspots	definite
date		
1818-01-11	NaN	1
1818-01-12	NaN	1
1818-01-13	22.0	1
1818-01-14	NaN	1
1818-01-15	NaN	1
1818-01-16	NaN	1
1818-01-17	46.0	1
1818-01-18	59.0	1
1818-01-19	63.0	1
1818-01-20	NaN	1



Writing files

```
In [26]: out_csv = 'sunspots.csv'
```

```
In [27]: sunspots.to_csv(out_csv)
```

```
In [28]: out_tsv = 'sunspots.tsv'
```

```
In [29]: sunspots.to_csv(out_tsv, sep='\t')
```

```
In [30]: out_xlsx = 'sunspots.xlsx'
```

```
In [31]: sunspots.to_excel(out_xlsx)
```




PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Plotting with pandas



AAPL stock data

```
In [1]: import pandas as pd
```

```
In [2]: import matplotlib.pyplot as plt
```

```
In [3]: aapl = pd.read_csv('aapl.csv', index_col='date',  
....:                      parse_dates=True)
```

```
In [4]: aapl.head(6)
```

```
Out[4]:
```

	adj_close	close	high	low	open	volume
date						
2000-03-01	31.68	130.31	132.06	118.50	118.56	38478000
2000-03-02	29.66	122.00	127.94	120.69	127.00	11136800
2000-03-03	31.12	128.00	128.23	120.00	124.87	11565200
2000-03-06	30.56	125.69	129.13	125.00	126.00	7520000
2000-03-07	29.87	122.87	127.44	121.12	126.44	9767600
2000-03-08	29.66	122.00	123.94	118.56	122.87	9690800



Plotting arrays (matplotlib)

```
In [5]: close_arr = aapl['close'].values
```

```
In [6]: type(close_arr)
```

```
Out[6]: numpy.ndarray
```

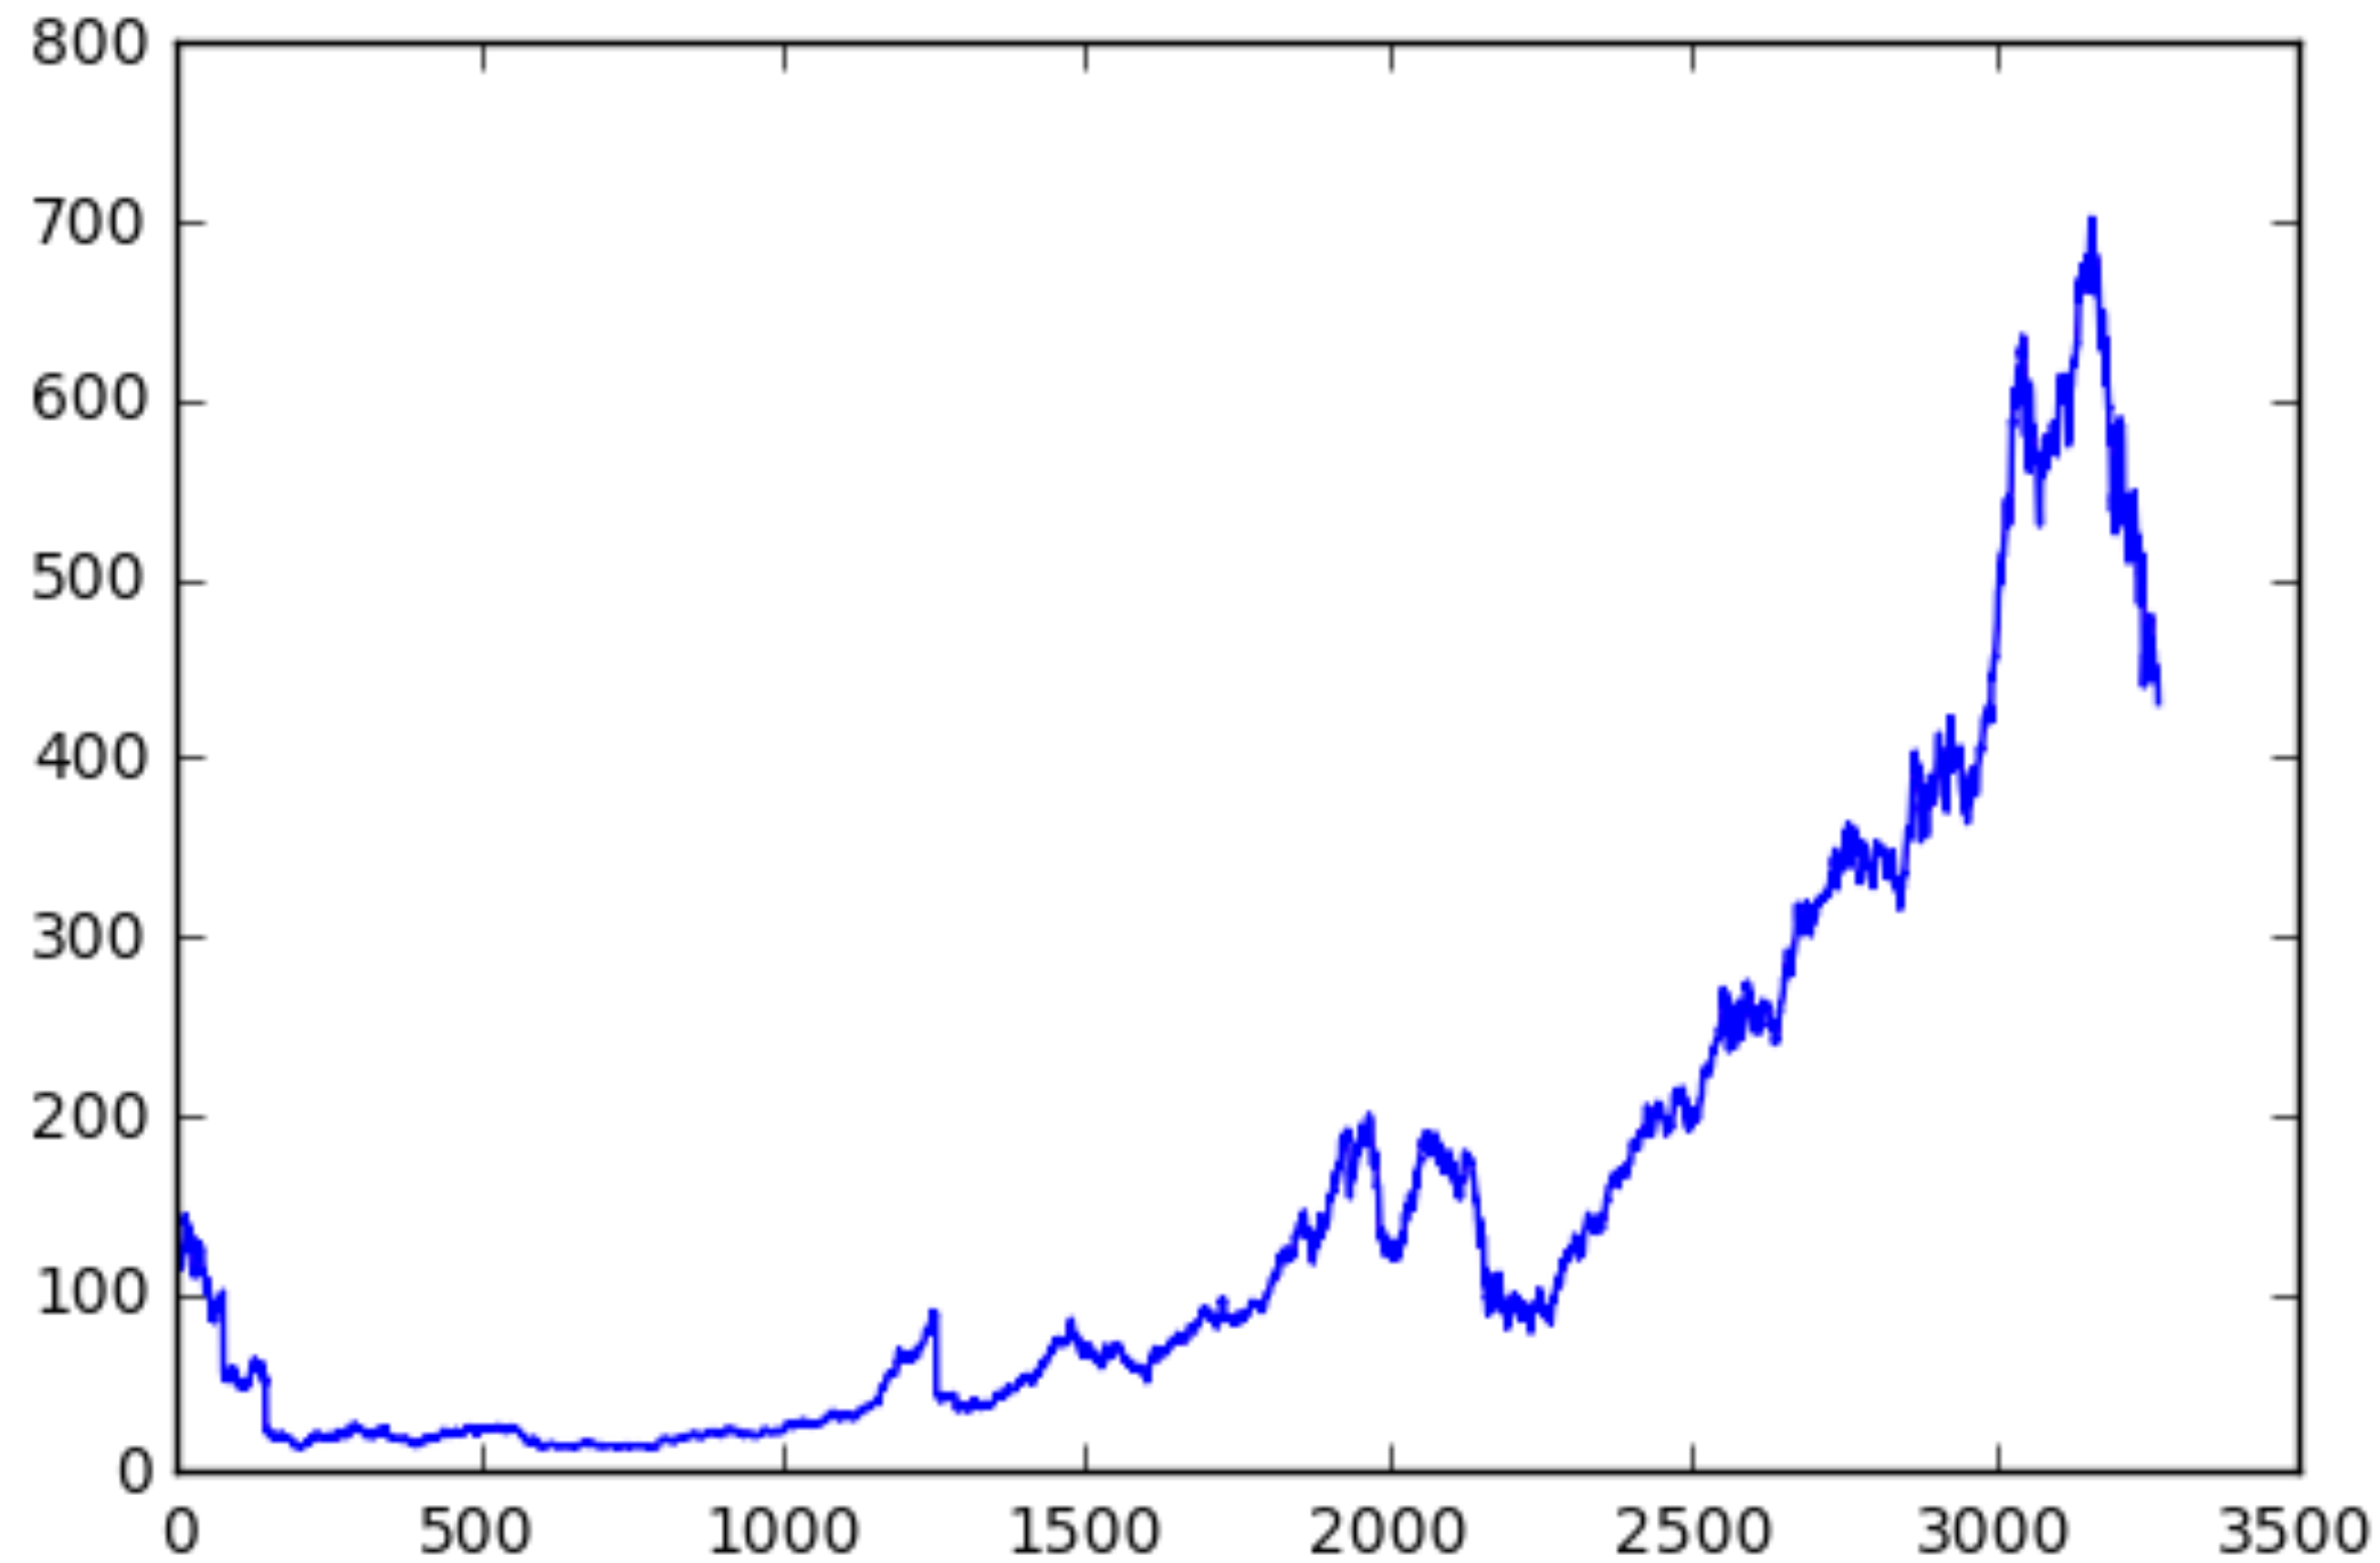
```
In [7]: plt.plot(close_arr)
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x115550358>]
```

```
In [8]: plt.show()
```



Plotting arrays (Matplotlib)





Plotting Series (matplotlib)

```
In [9]: close_series = aapl['close']
```


```
In [10]: type(close_series)
```

```
Out[10]: pandas.core.series.Series
```

```
In [11]: plt.plot(close_series)
```

```
Out[11]: [<matplotlib.lines.Line2D at 0x11801cd30>]
```

```
In [12]: plt.show()
```





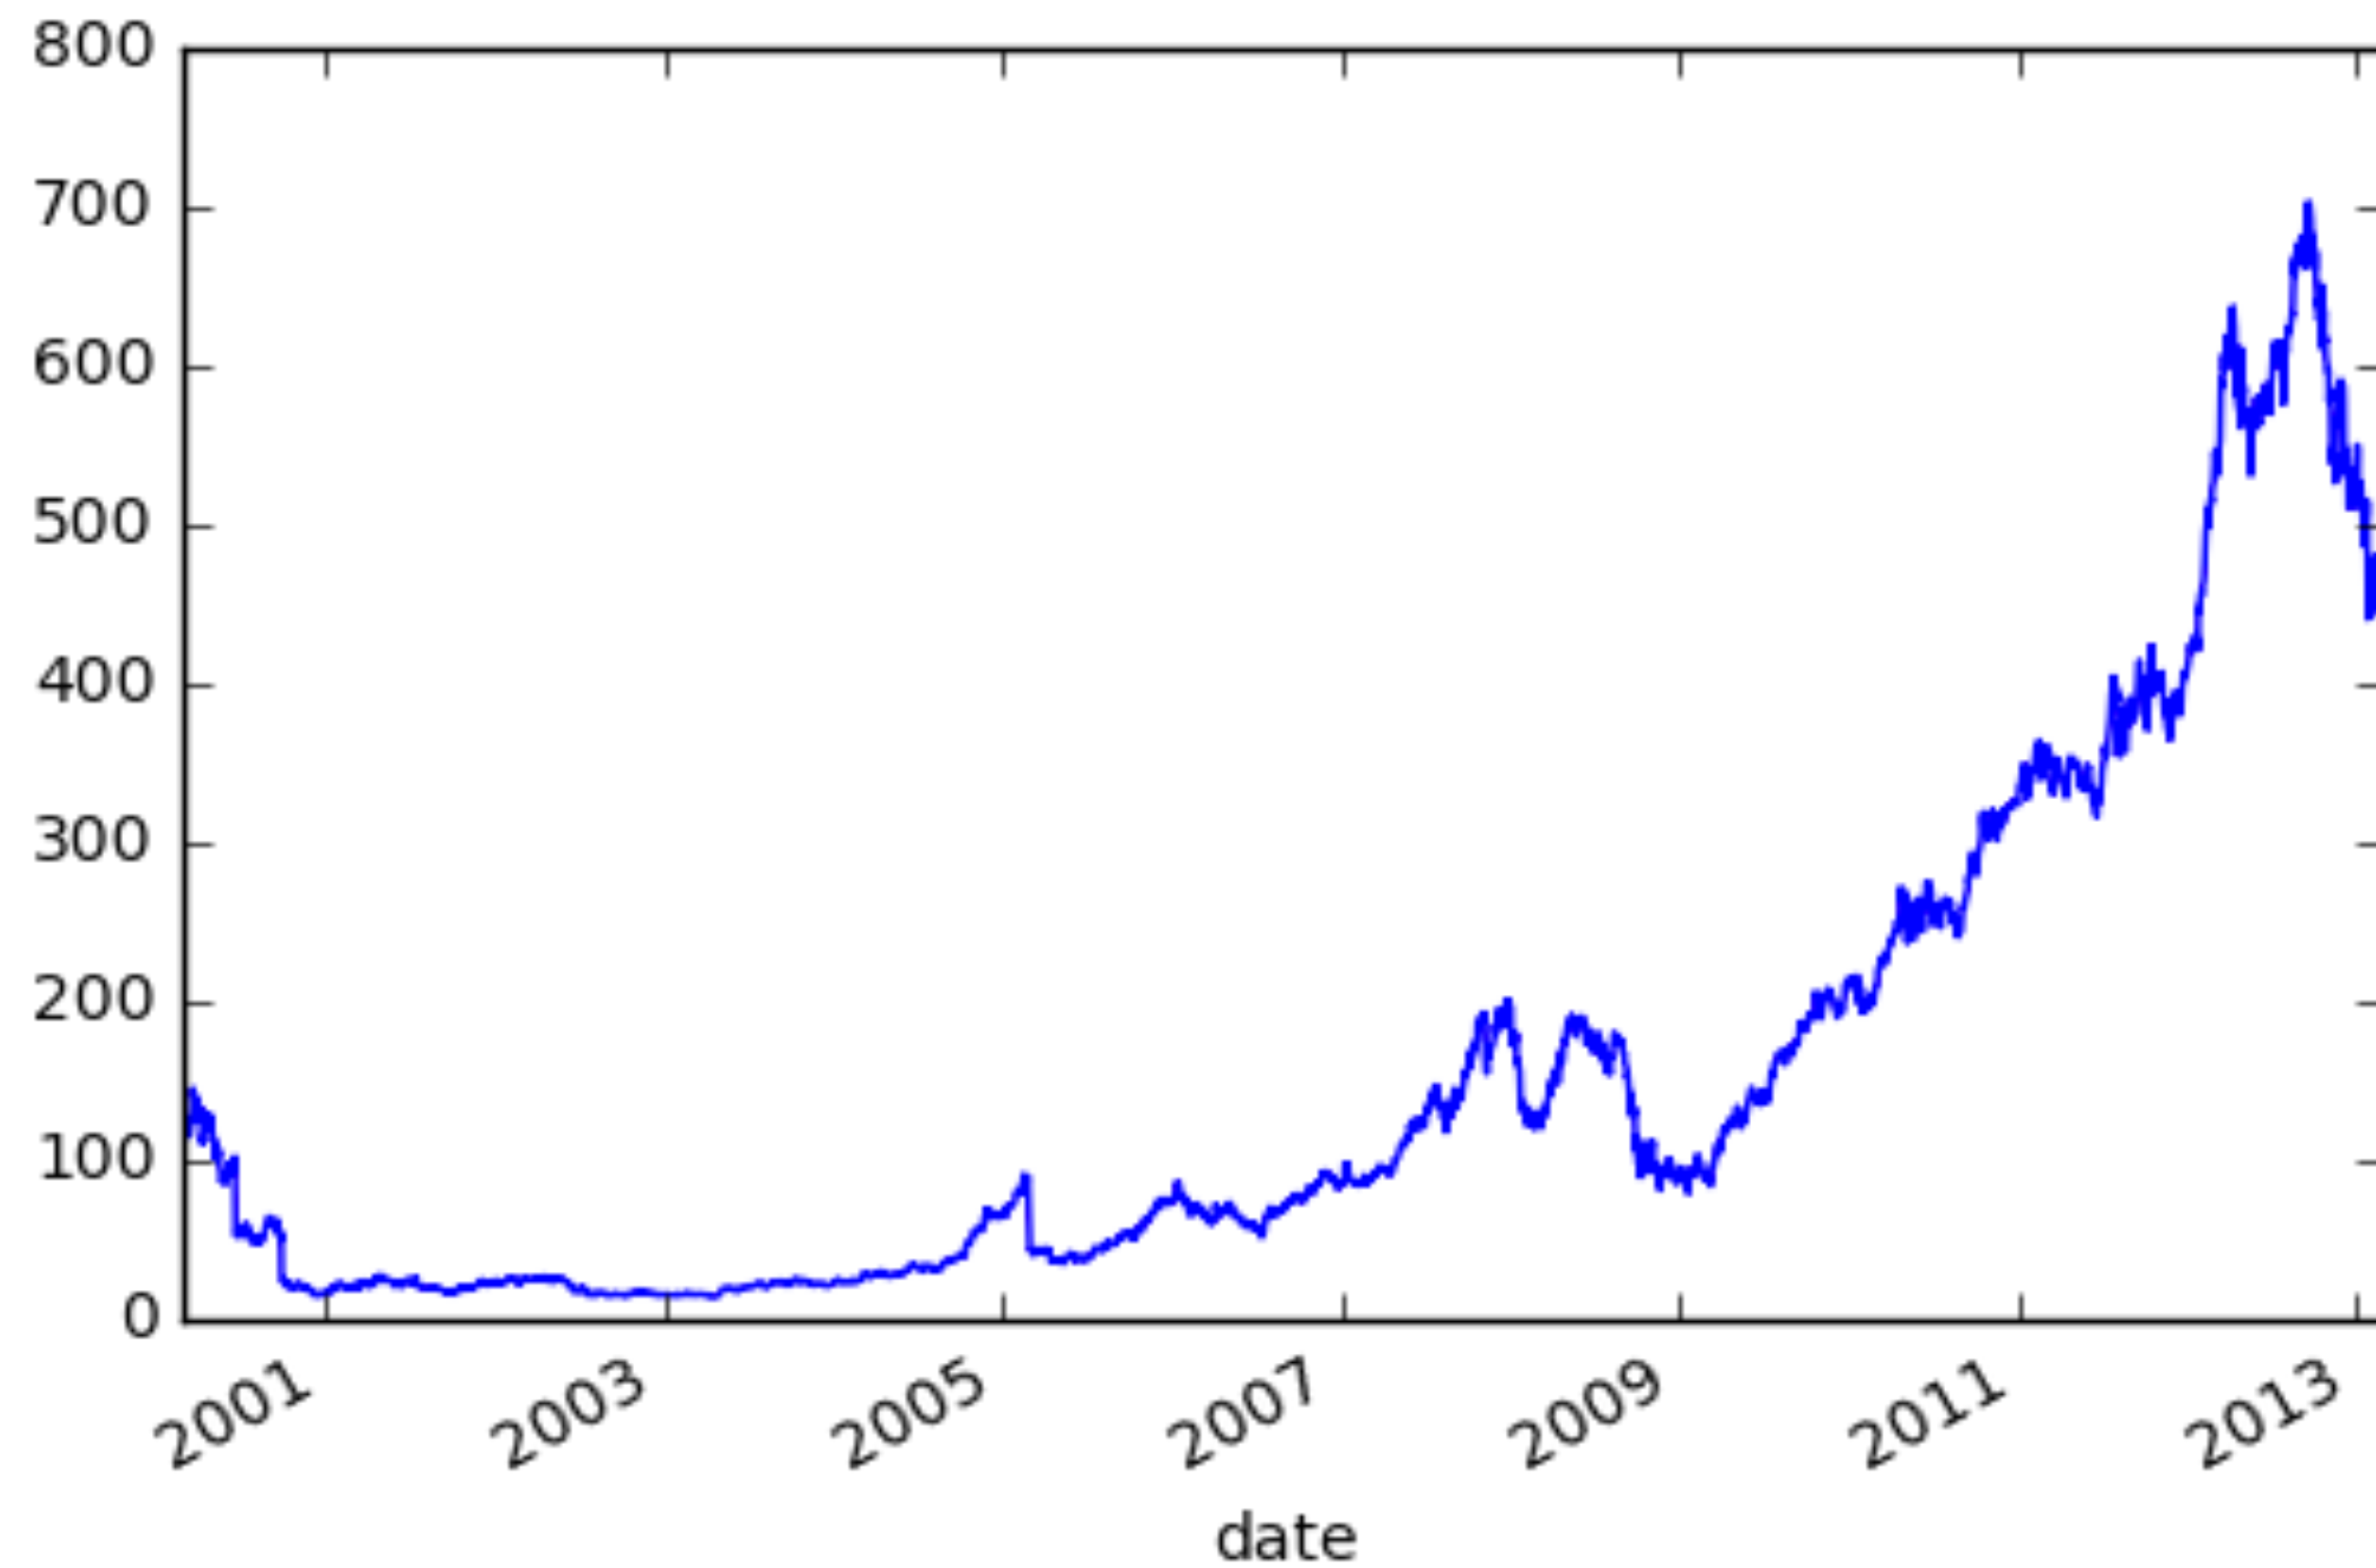
Plotting Series (pandas)

```
In [13]: close_series.plot() # plots Series directly
```

```
In [14]: plt.show()
```




Plotting Series (pandas)



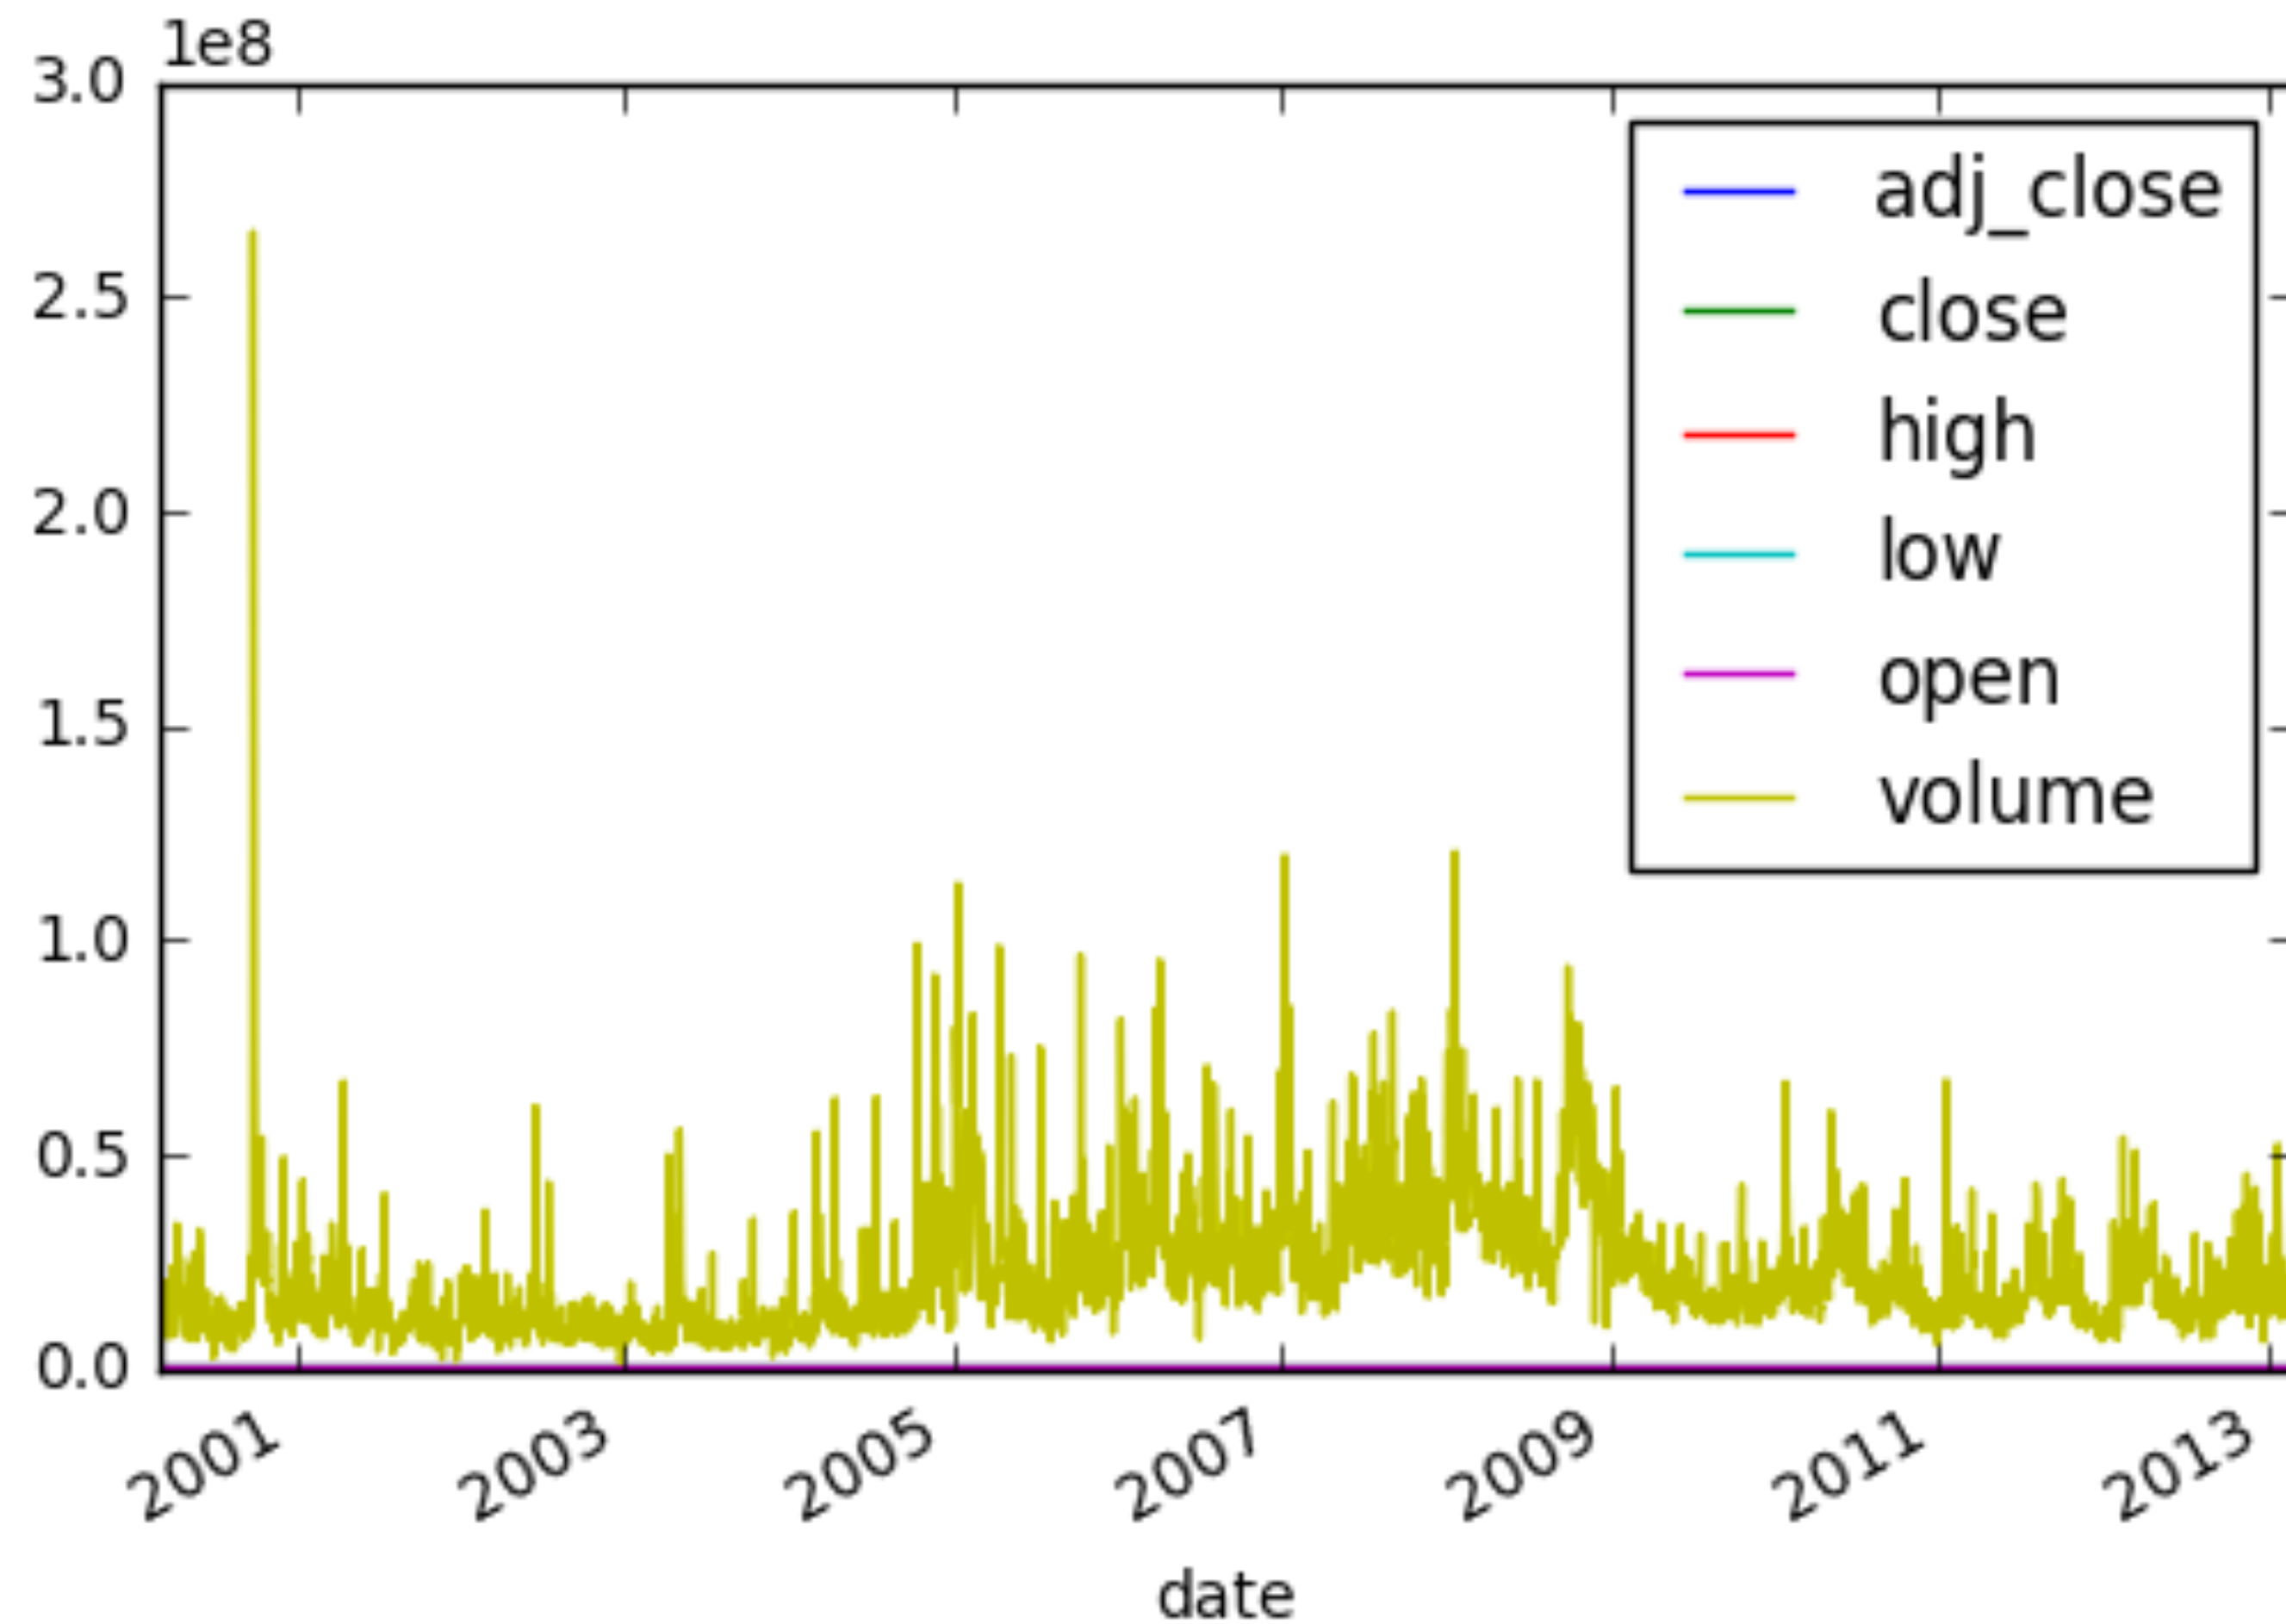


Plotting DataFrames (pandas)

```
In [15]: aapl.plot() # plots all Series at once  
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x118039b38>  
  
In [16]: plt.show()
```



Plotting DataFrames (pandas)





Plotting DataFrames (matplotlib)

```
In [17]: plt.plot(aapl) # plots all columns at once
```

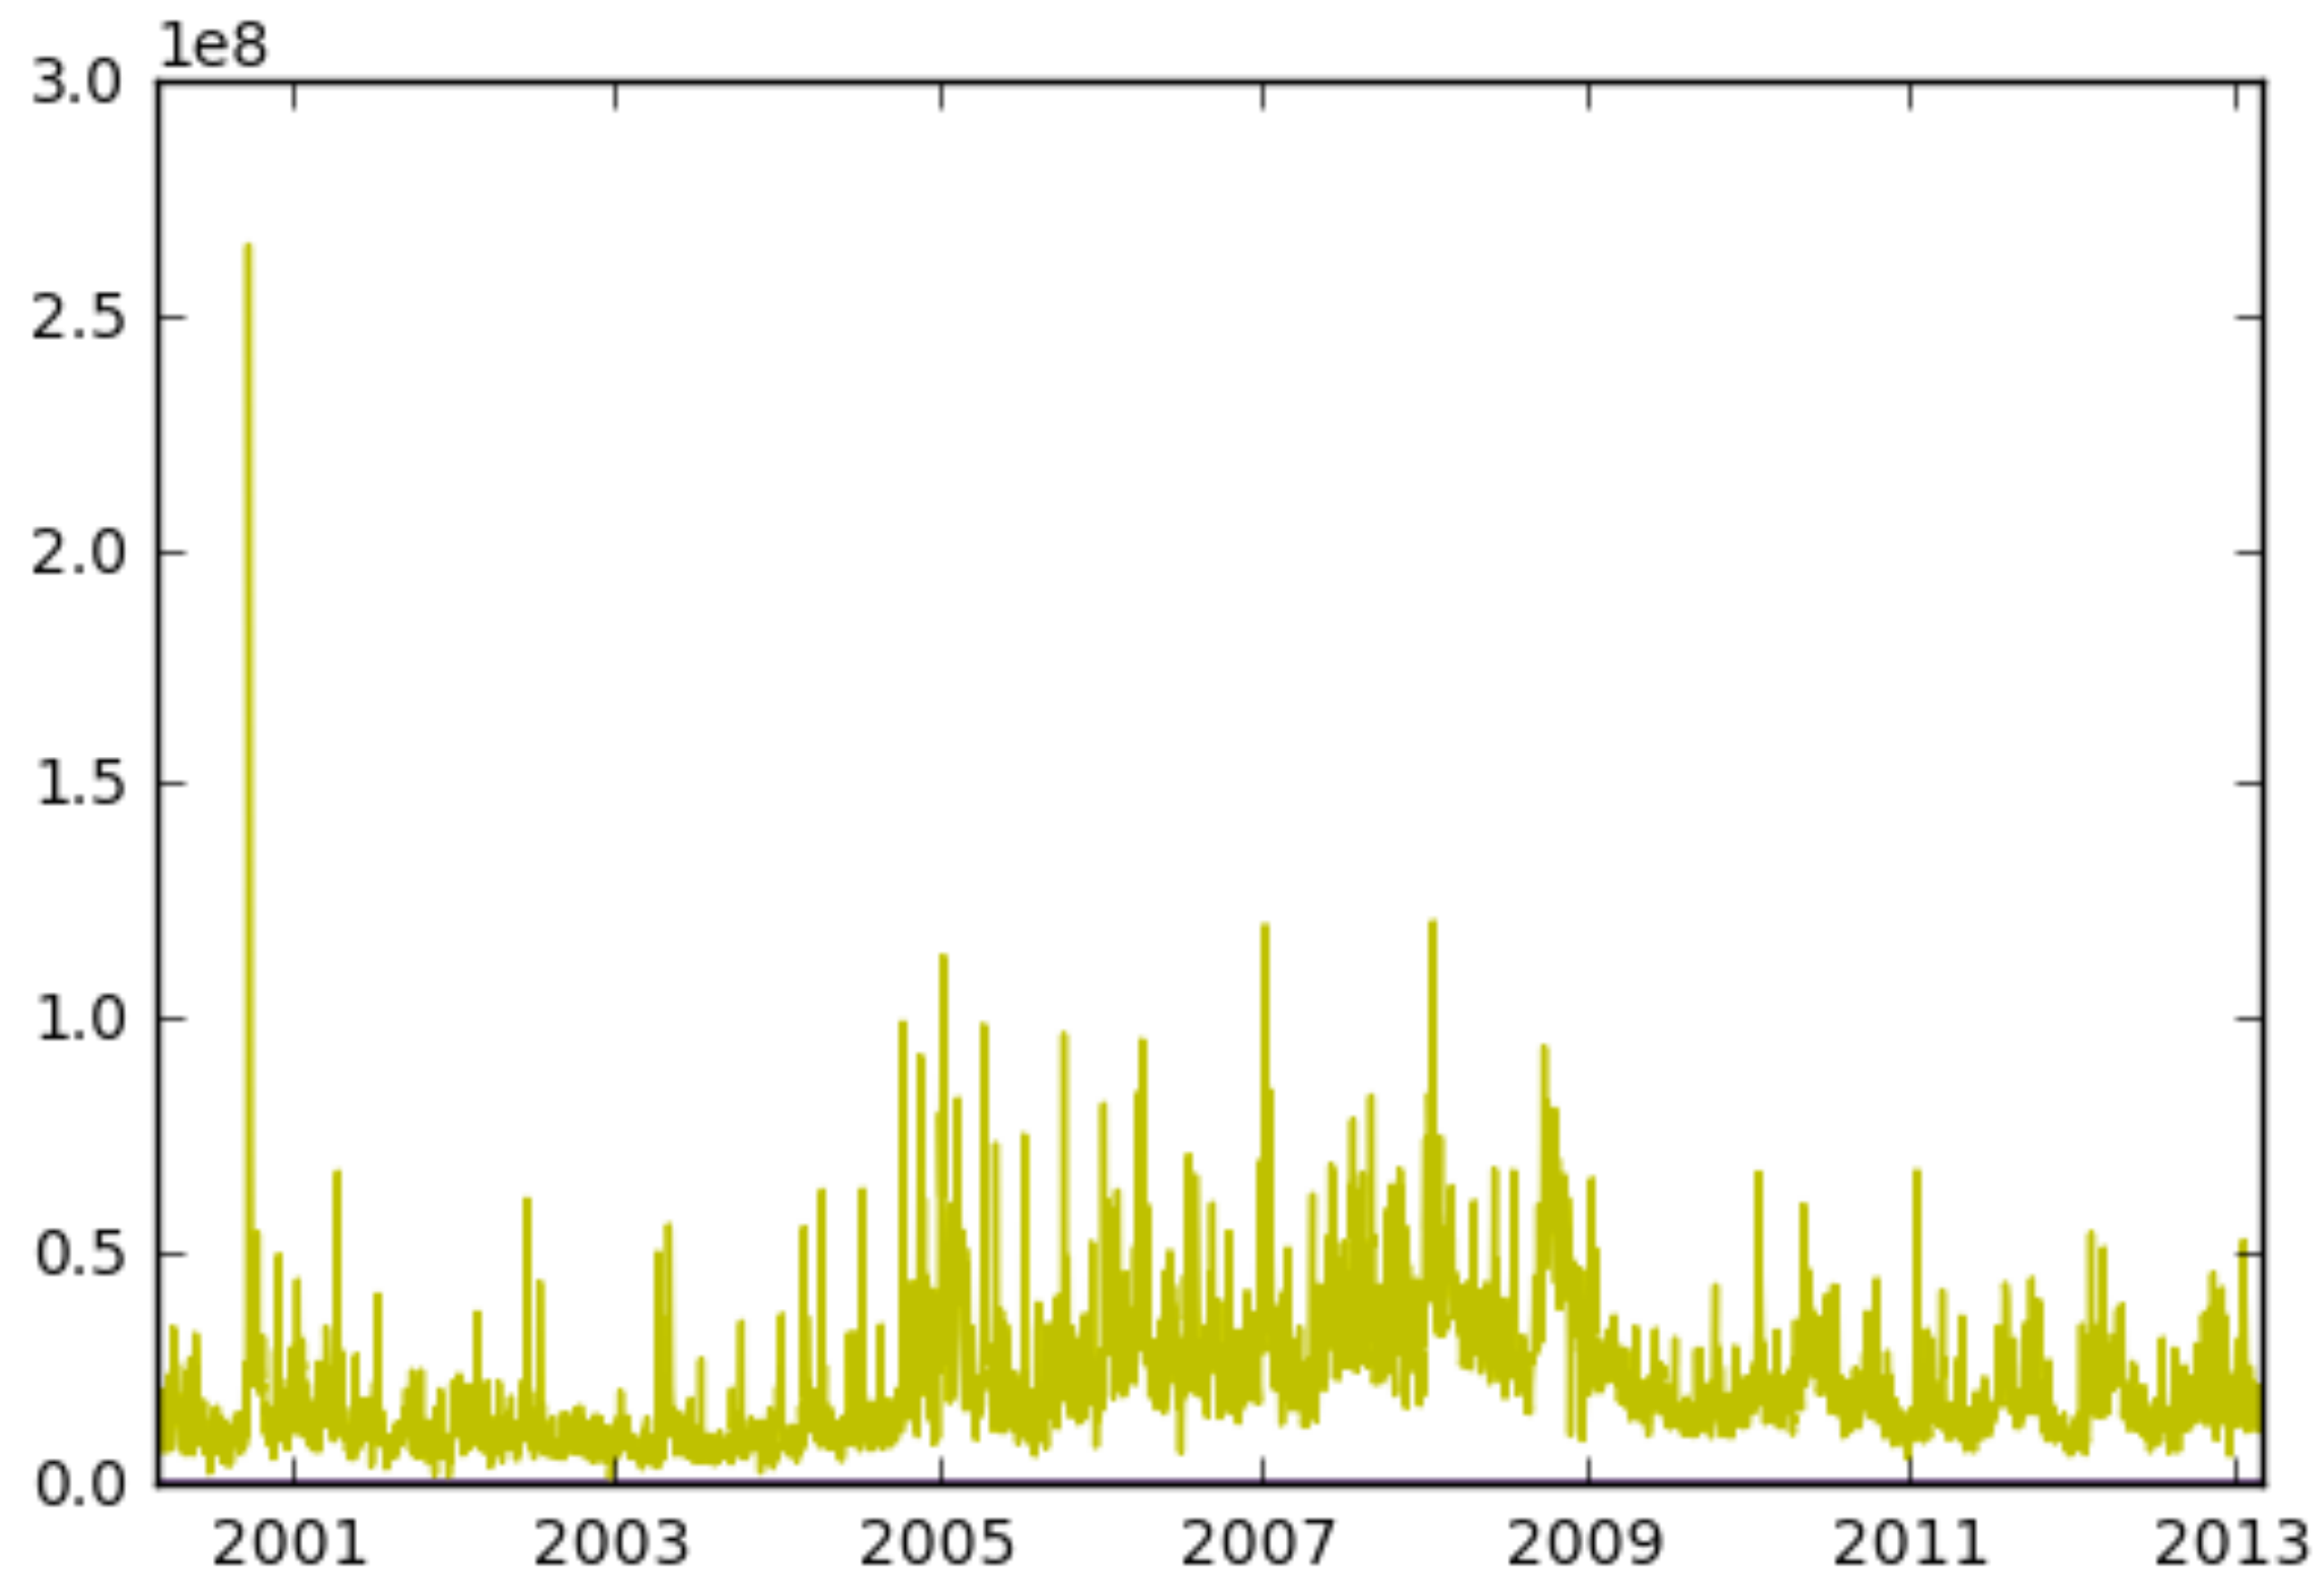
```
Out[17]:
```

```
<matplotlib.lines.Line2D at 0x1156290f0>,  
<matplotlib.lines.Line2D at 0x1156525f8>,  
<matplotlib.lines.Line2D at 0x1156527f0>,  
<matplotlib.lines.Line2D at 0x1156529e8>,  
<matplotlib.lines.Line2D at 0x115652be0>,  
<matplotlib.lines.Line2D at 0x115652dd8>
```

```
In [18]: plt.show()
```



Plotting DataFrames (matplotlib)





Fixing scales

```
In [19]: aapl.plot()
```

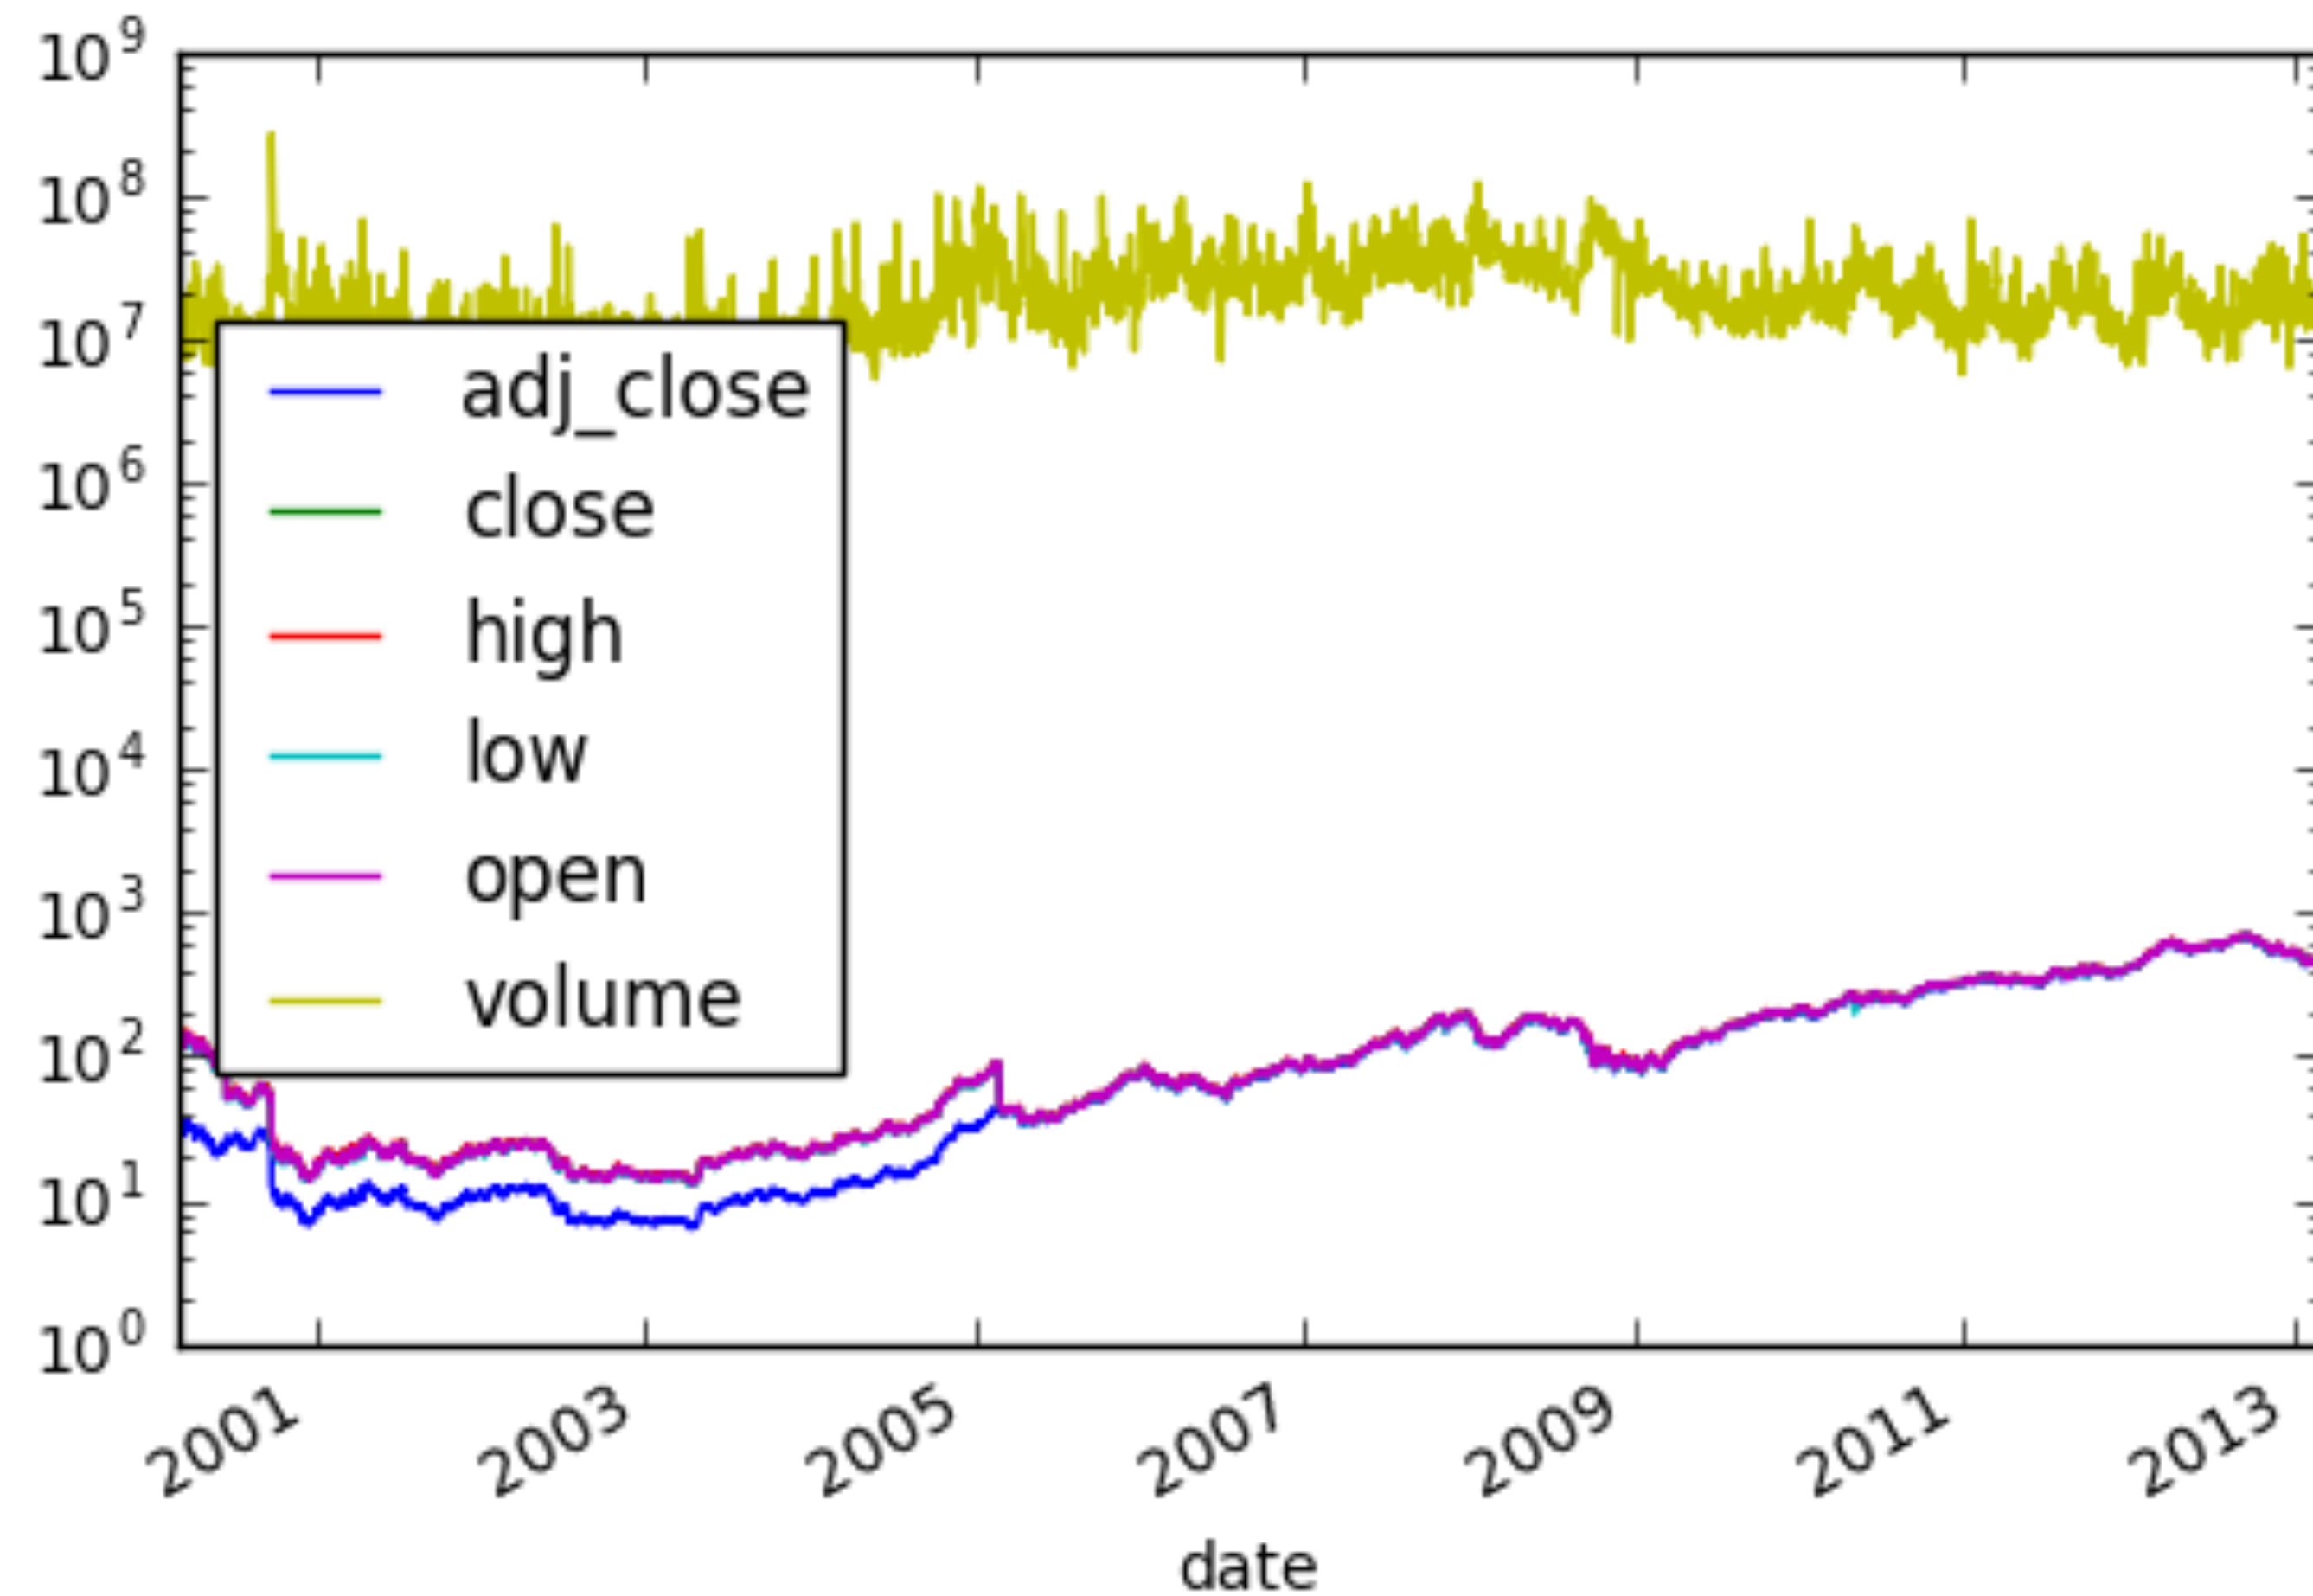
```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x118afe048>
```

```
In [20]: plt.yscale('log') # logarithmic scale on vertical axis
```

```
In [21]: plt.show()
```



Fixing scales





Customizing plots

```
In [22]: aapl['open'].plot(color='b', style='.-', legend=True)
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x11a17db38>
```

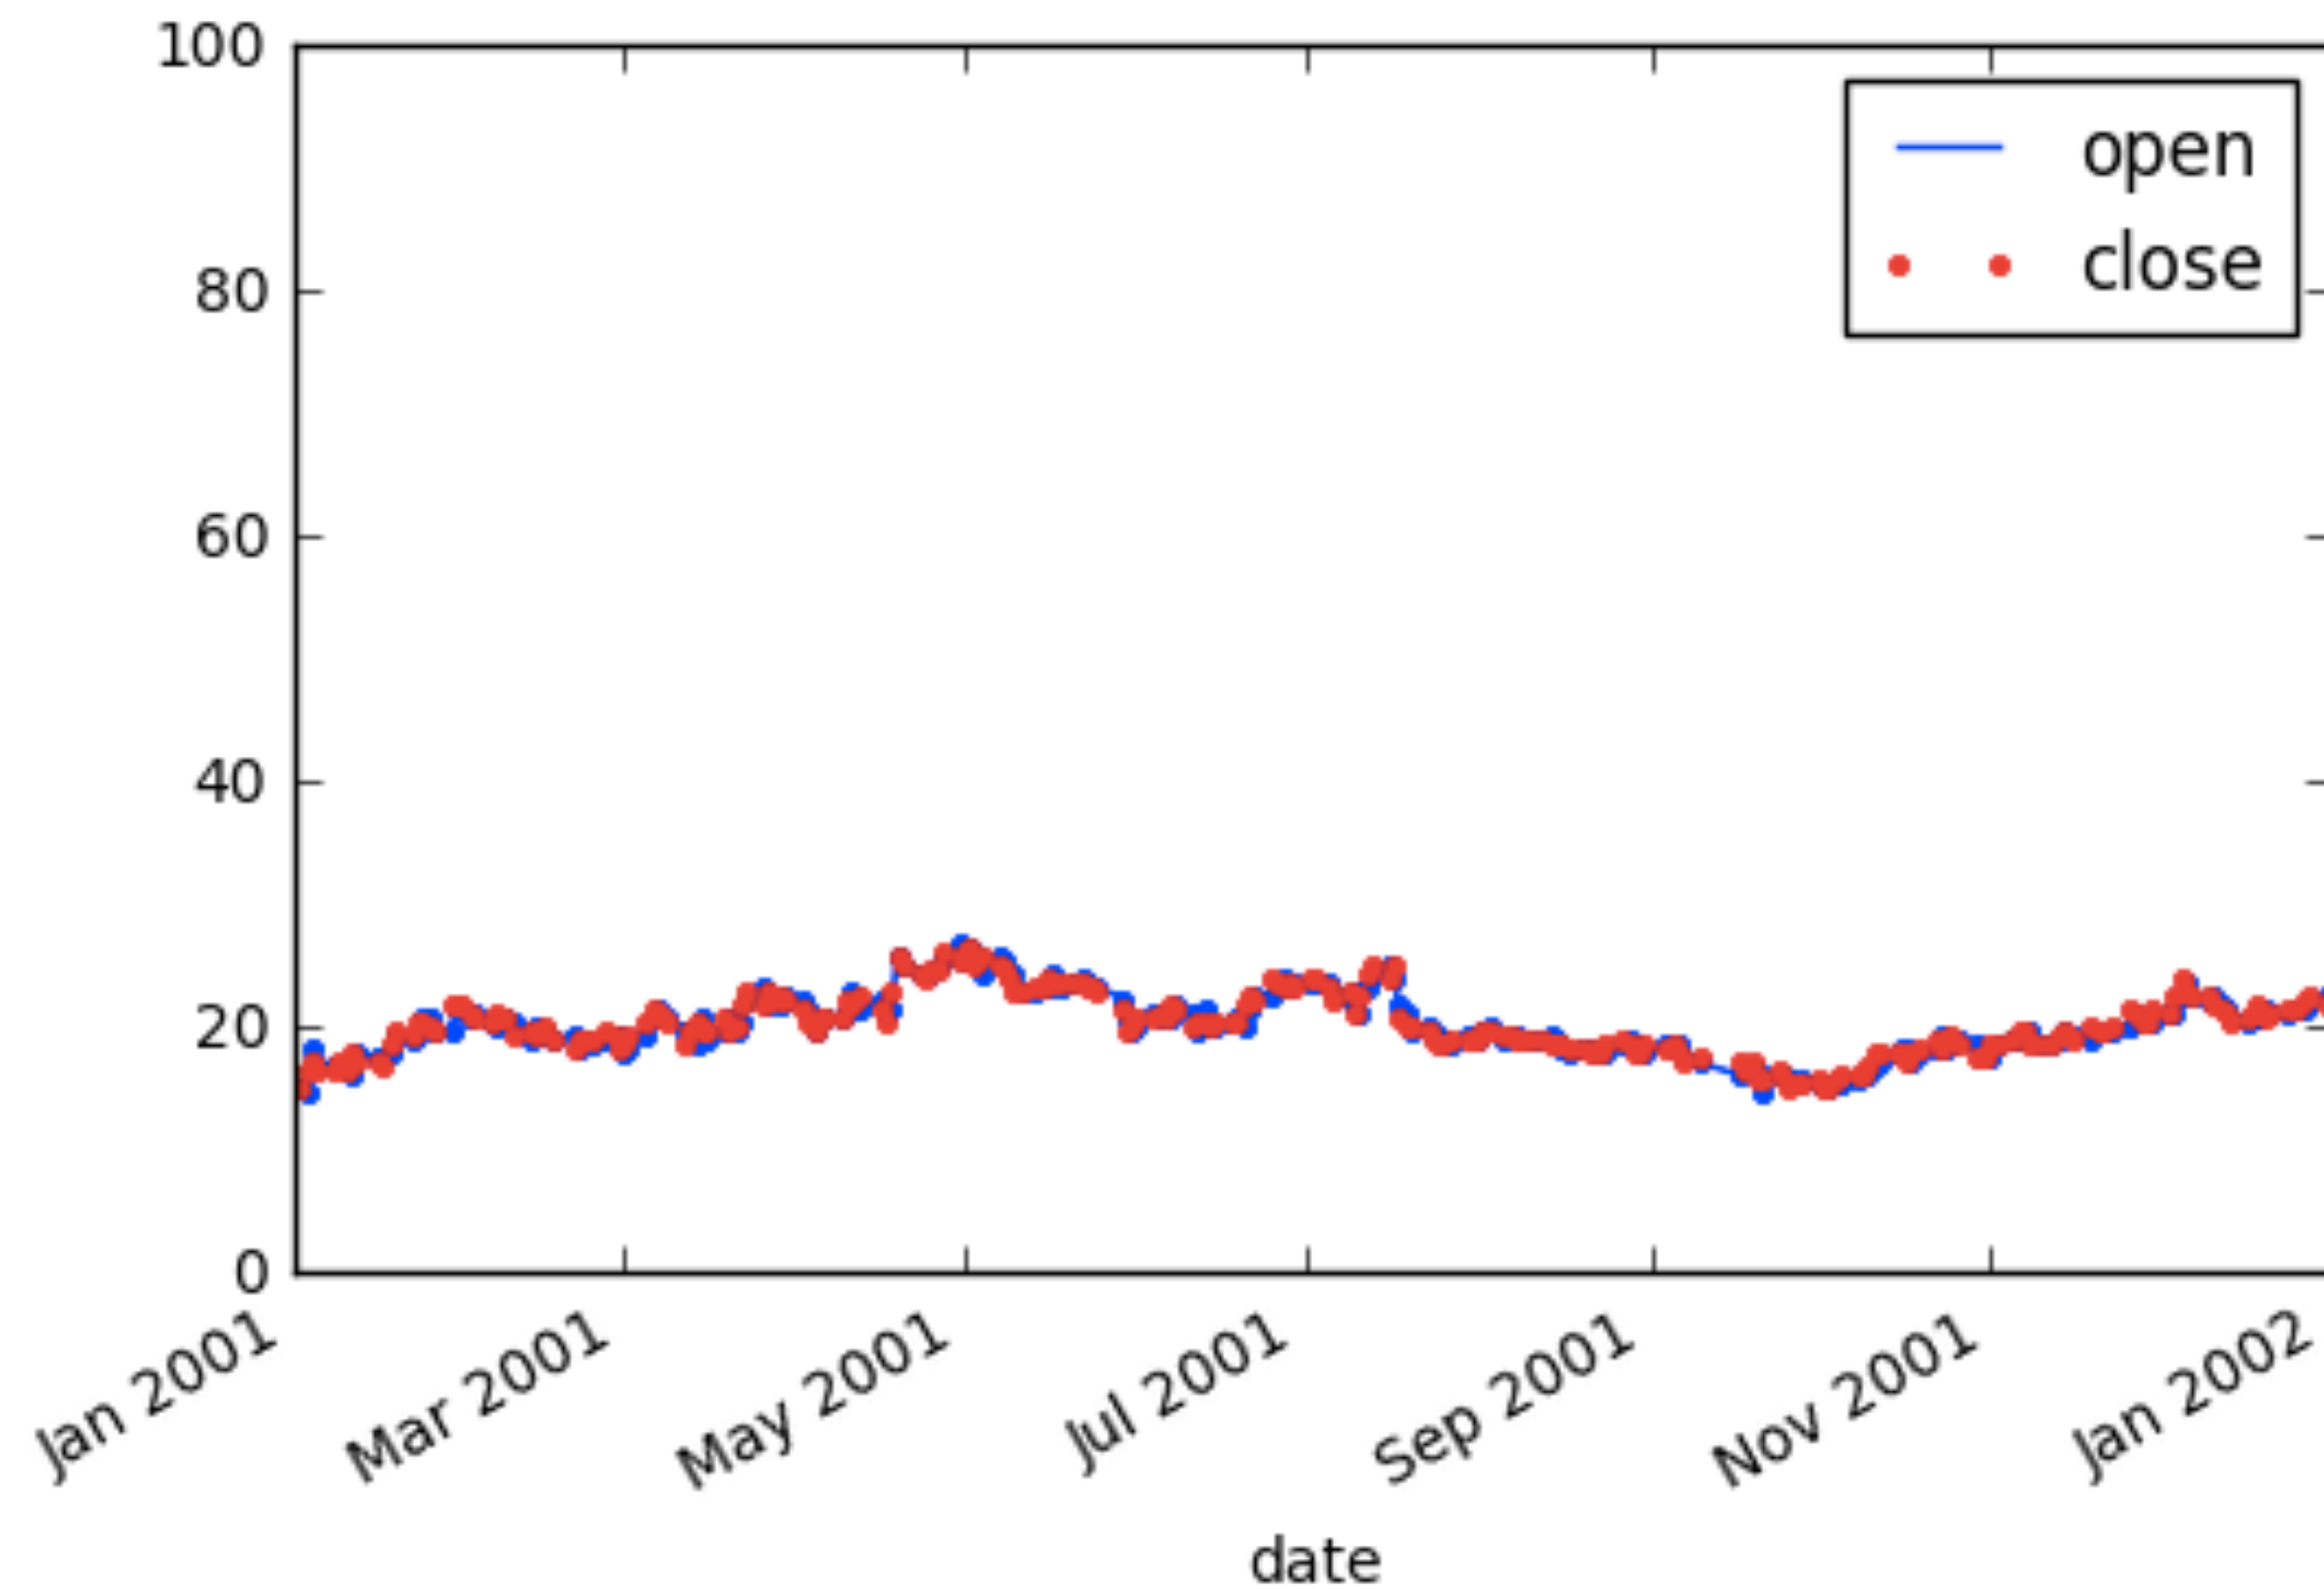
```
In [23]: aapl['close'].plot(color='r', style='.', legend=True)
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x11a17db38>
```

```
In [24]: plt.axis(('2001', '2002', 0, 100))
Out[24]: ('2001', '2002', 0, 100)
```

```
In [25]: plt.show()
```

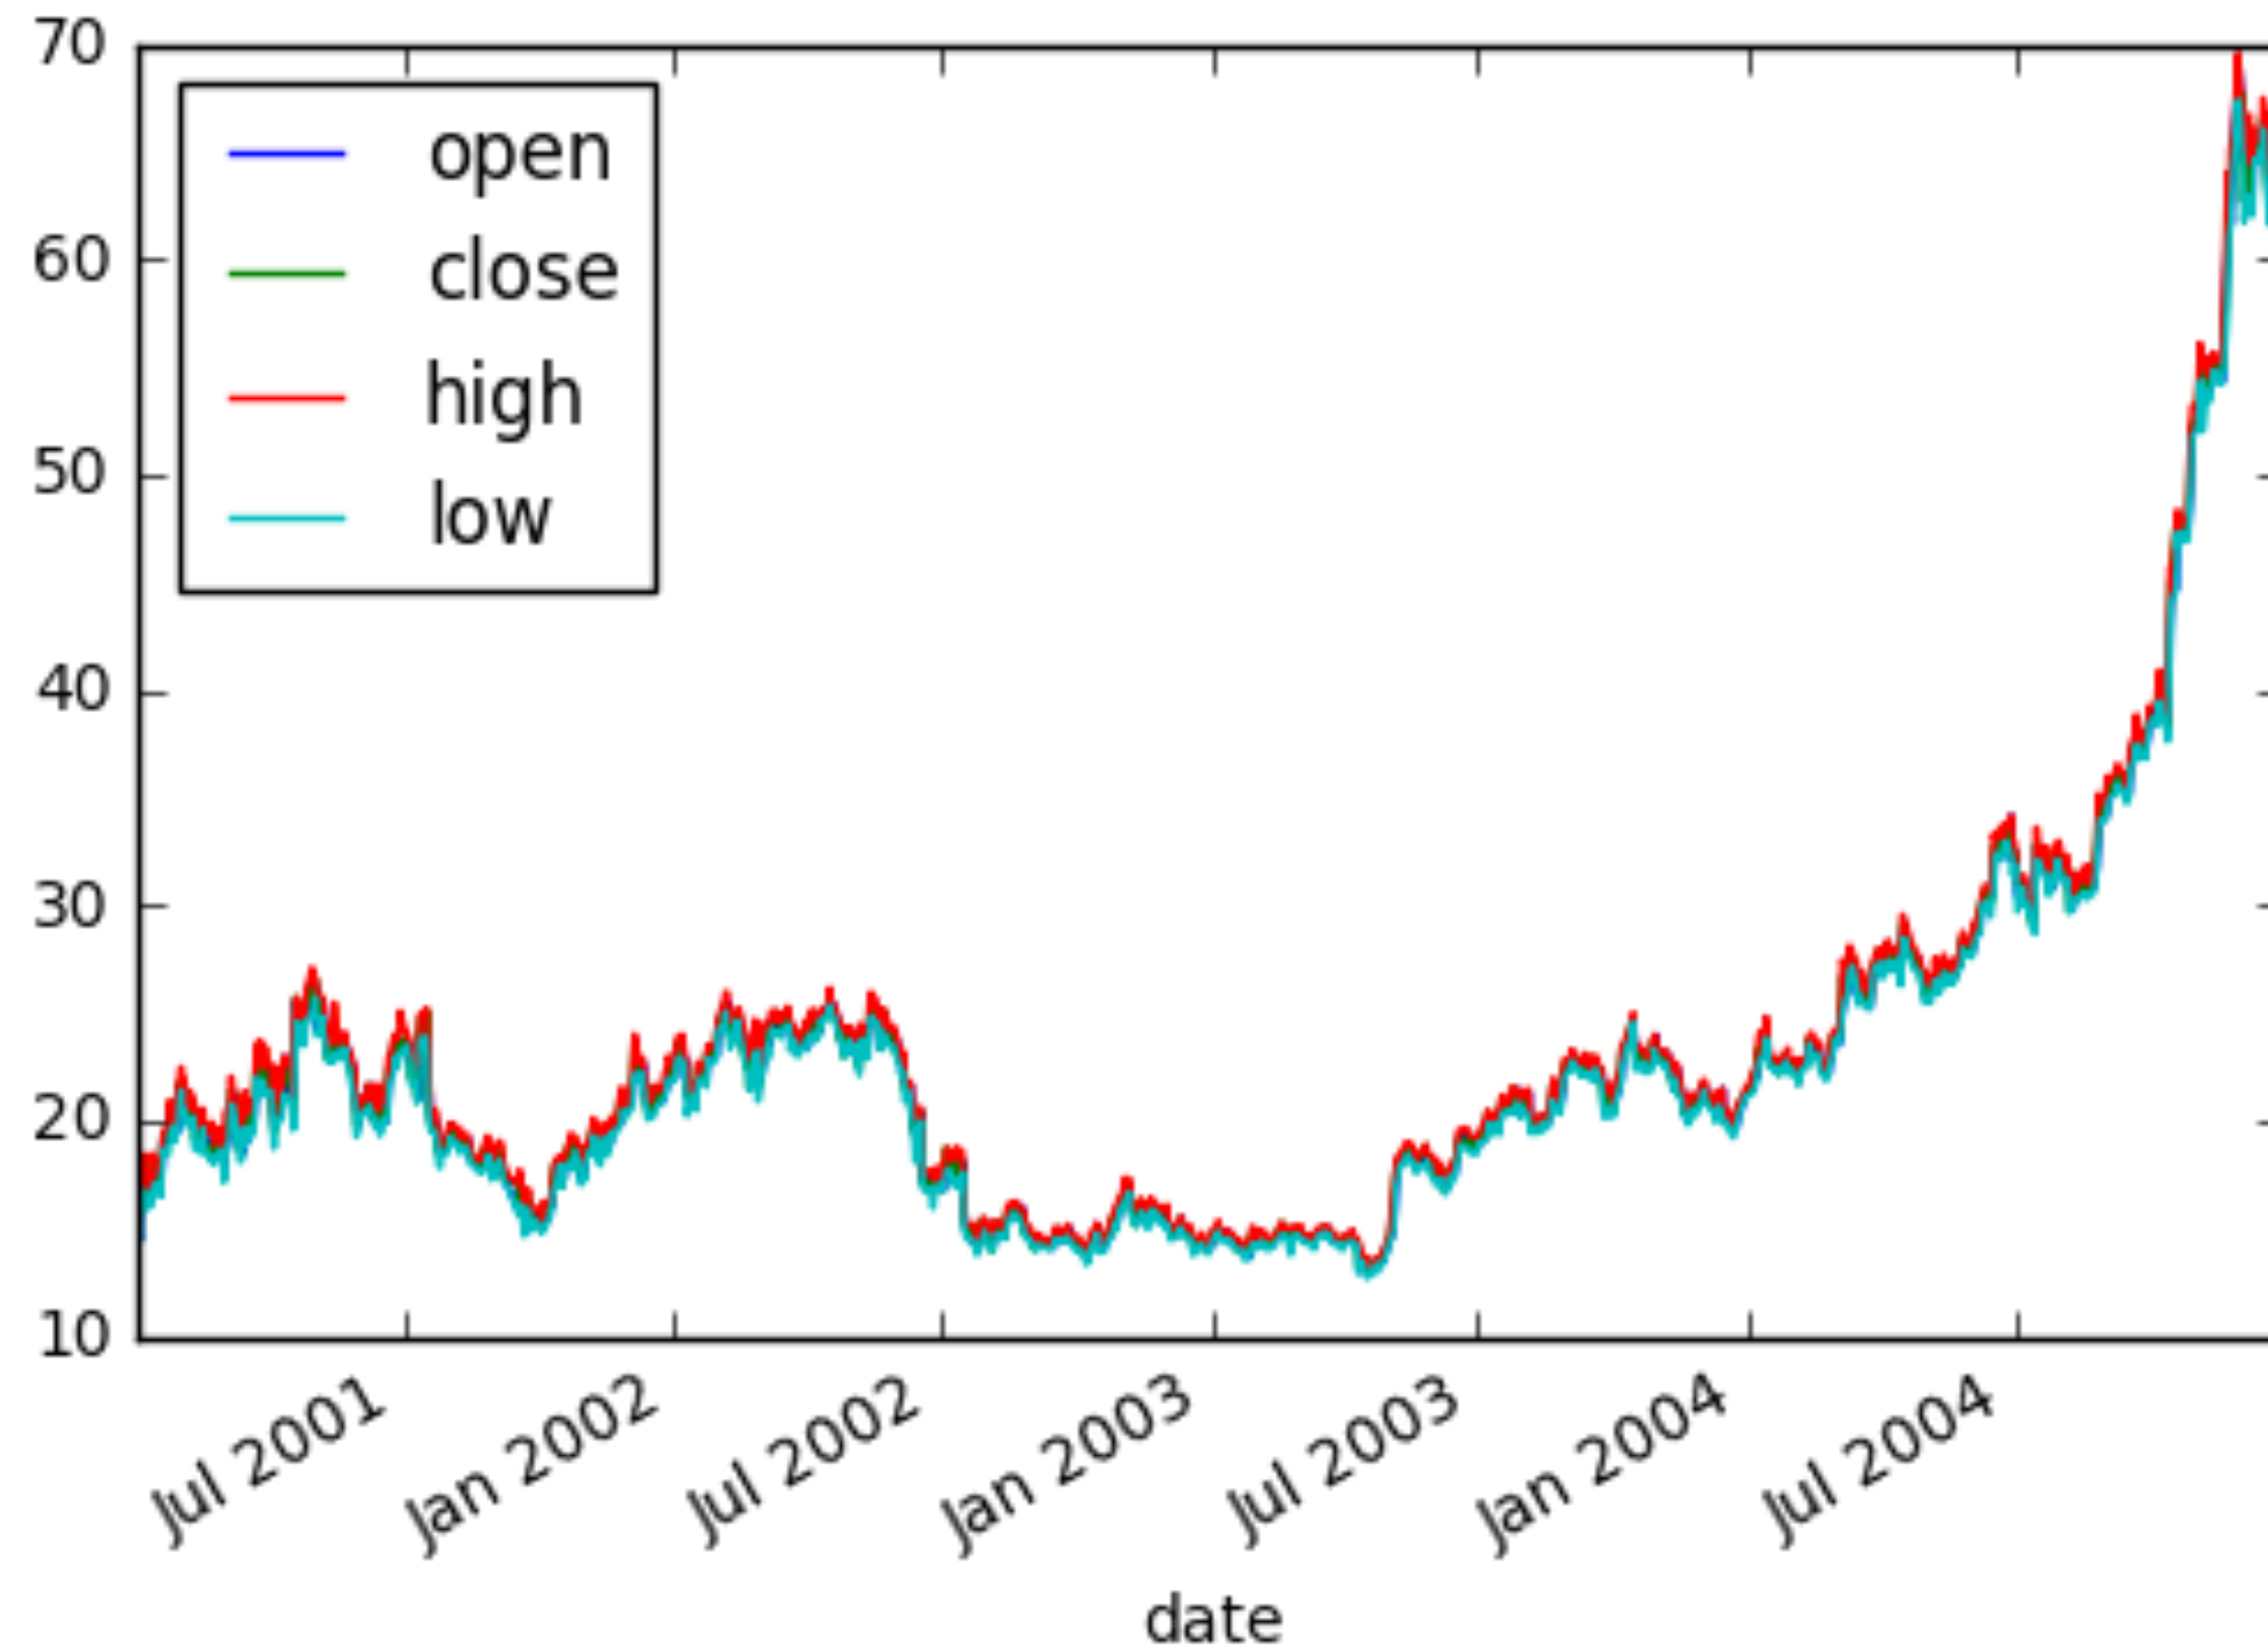



Customizing plots





Saving plots





Saving plots

```
In [26]: aapl.loc['2001':'2004',['open', 'close', 'high',  
....:      'low']].plot()  
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x11ab42978>  
  
In [27]: plt.savefig('aapl.png')  
  
In [28]: plt.savefig('aapl.jpg')  
  
In [29]: plt.savefig('aapl.pdf')  
  
In [30]: plt.show()
```



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Visual exploratory data analysis



The iris data set

- Famous data set in pattern recognition
- 150 observations, 4 features each
 - Sepal length
 - Sepal width
 - Petal length
 - Petal width
- 3 species: setosa, versicolor, virginica



Data import

```
In [1]: import pandas as pd
```

```
In [2]: import matplotlib.pyplot as plt
```

```
In [3]: iris = pd.read_csv('iris.csv', index_col=0)
```

```
In [4]: print(iris.shape)  
(150, 5)
```



Line plot

```
In [5]: iris.head()
```

```
Out[5]:
```

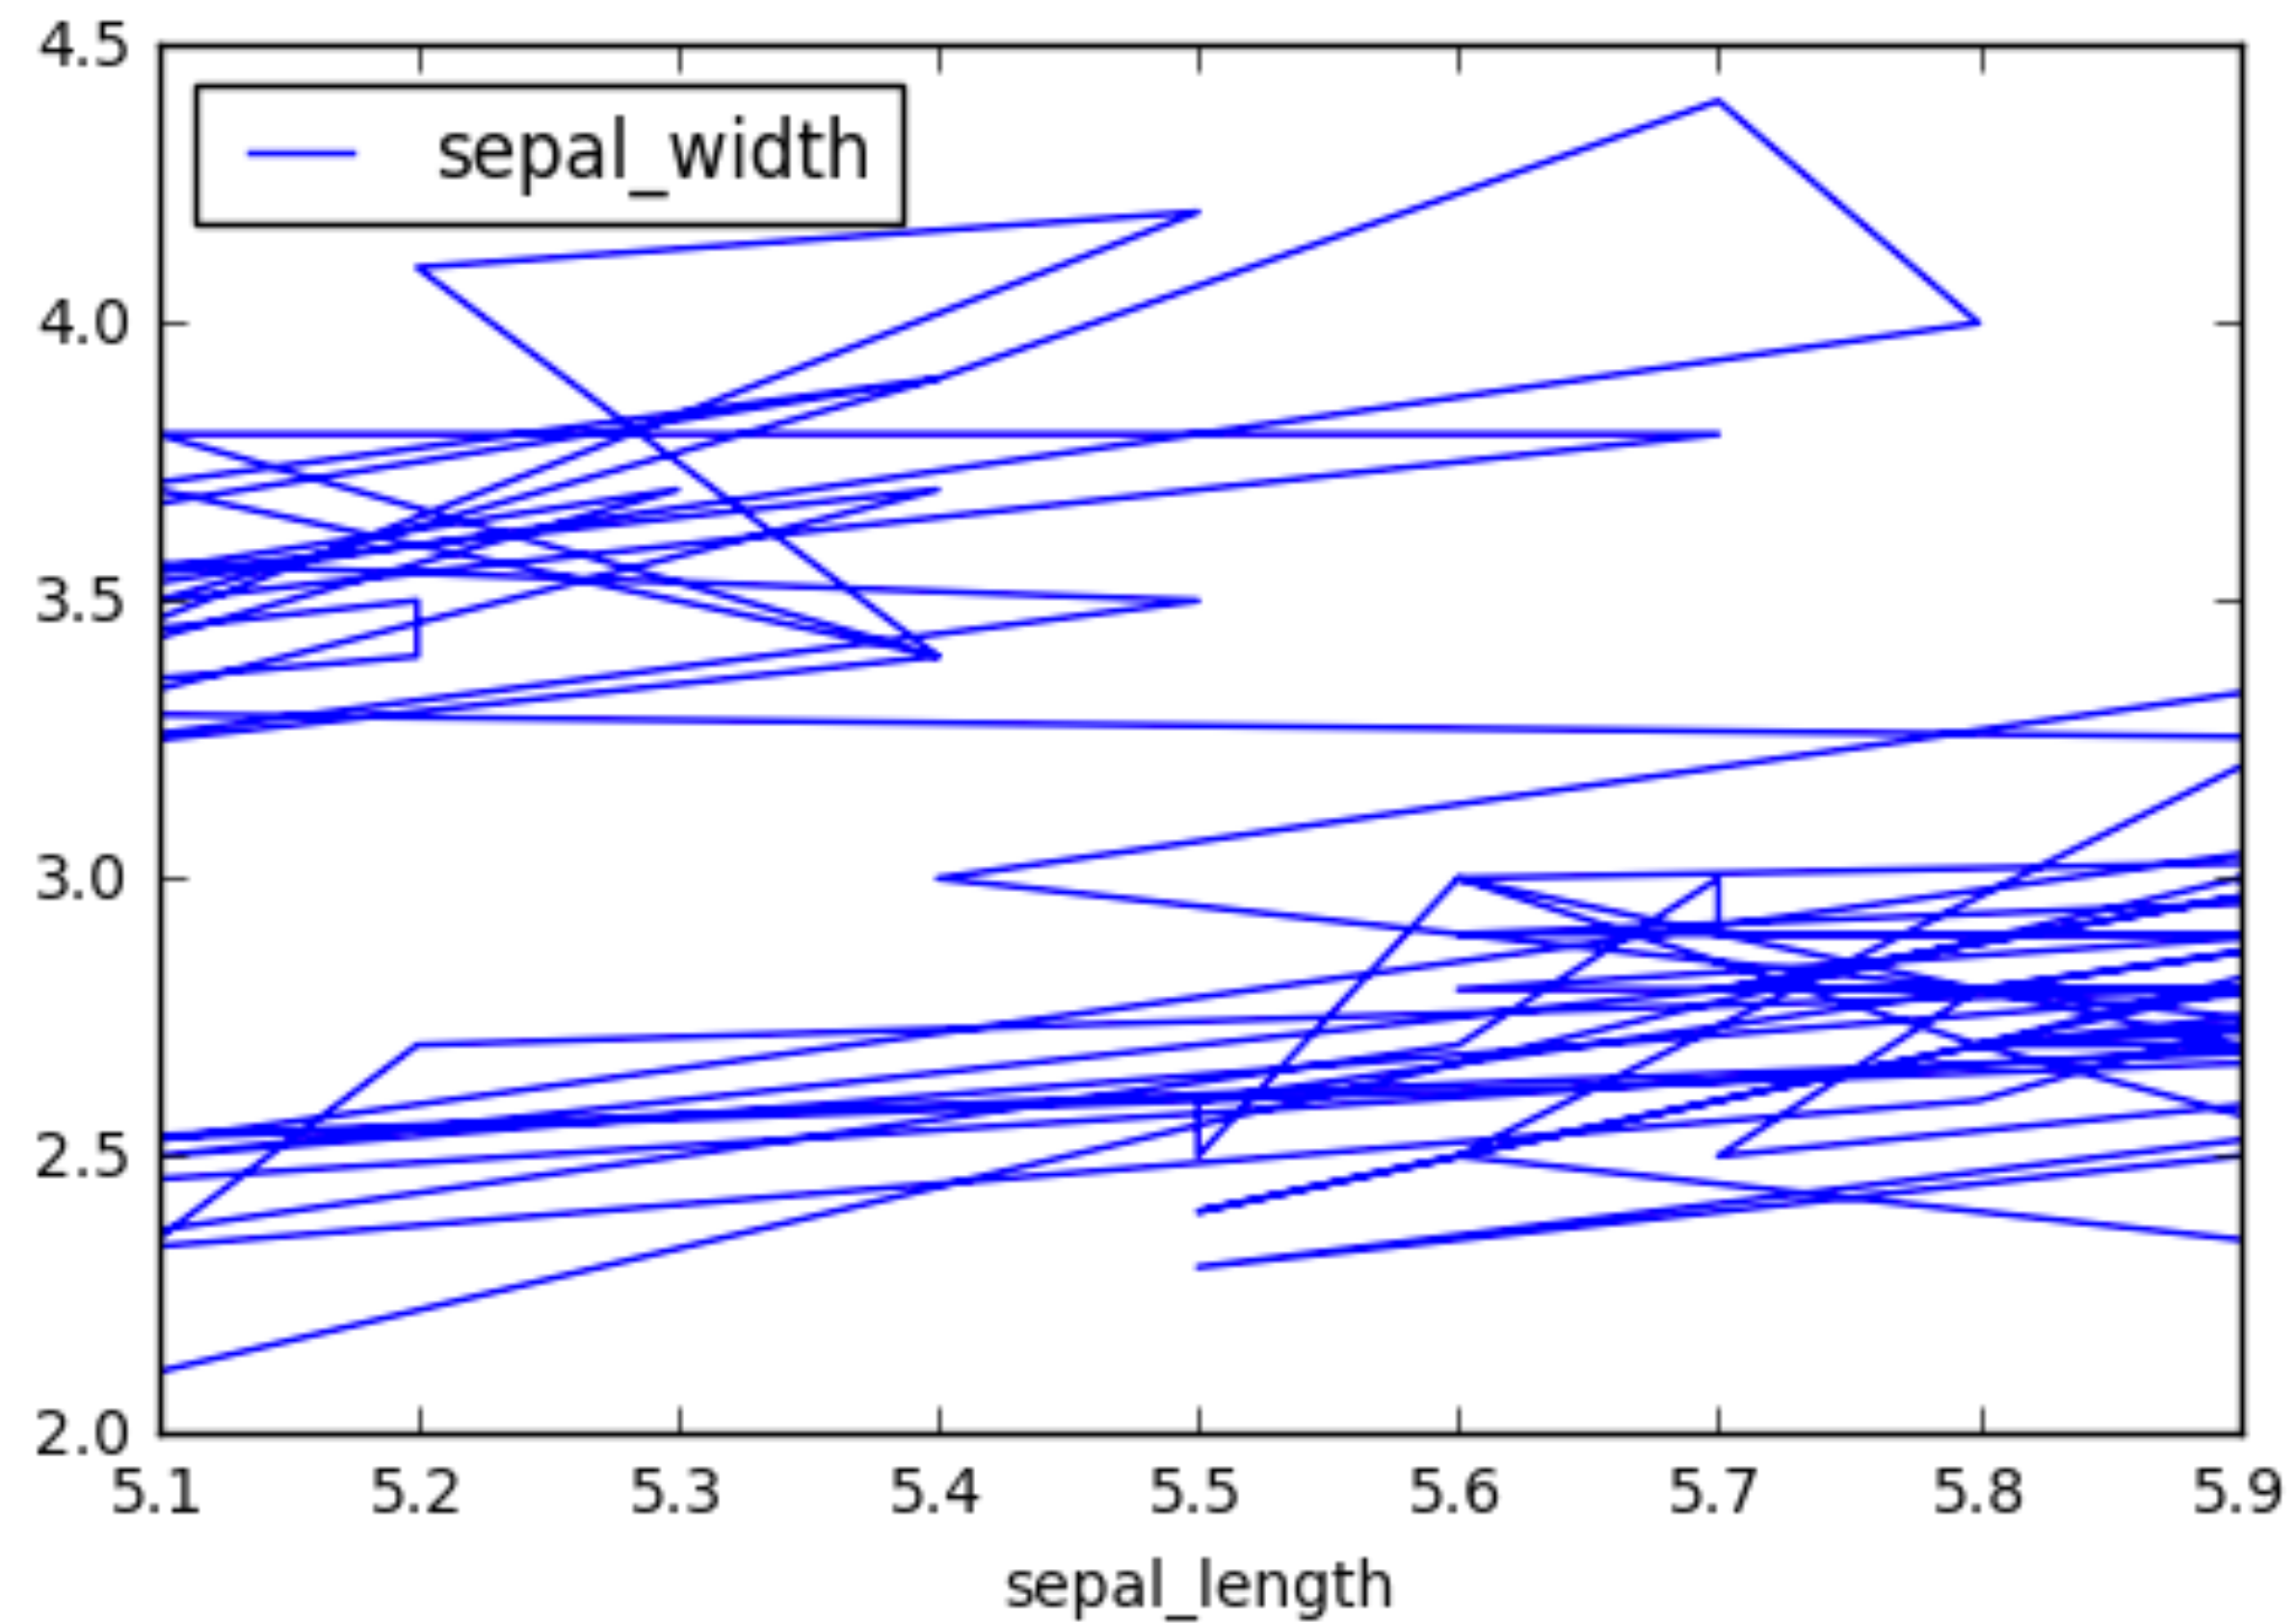
	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [6]: iris.plot(x='sepal_length', y='sepal_width')
```

```
In [7]: plt.show()
```




Line plot





Scatter plot

```
In [8]: iris.plot(x='sepal_length', y='sepal_width',  
....:             kind='scatter')
```

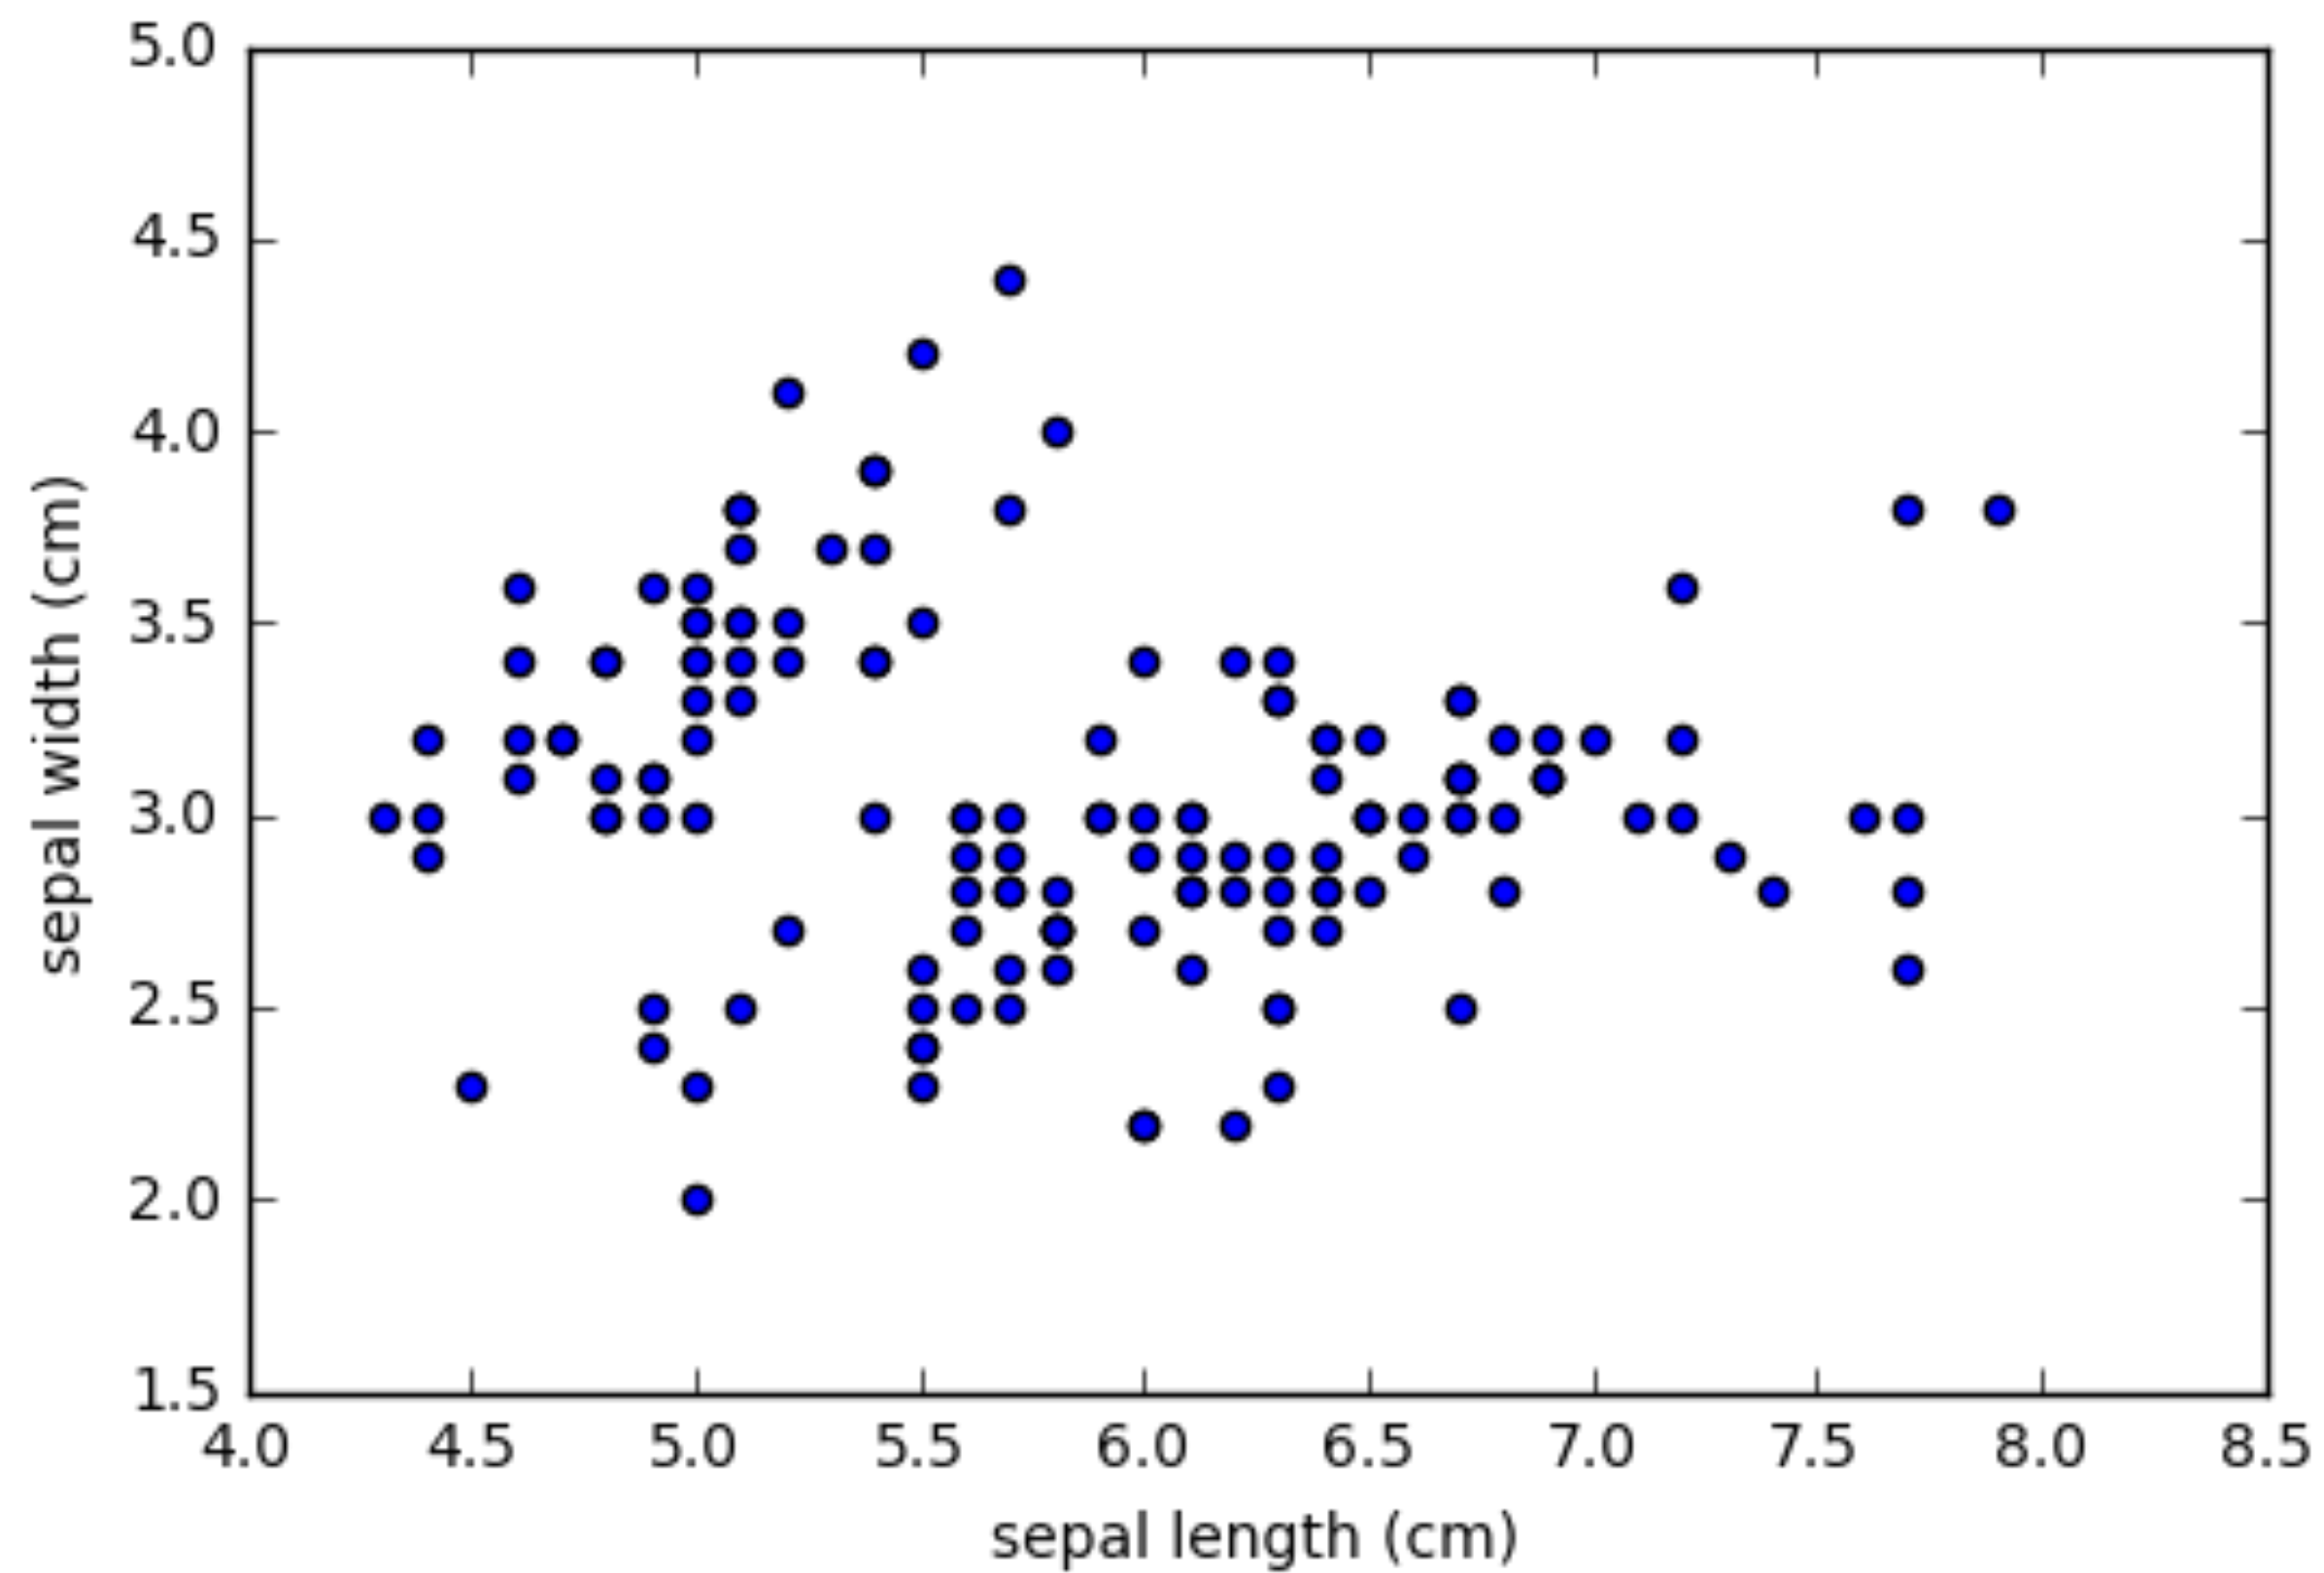
```
In [9]: plt.xlabel('sepal length (cm)')
```

```
In [10]: plt.ylabel('sepal width (cm)')
```

```
In [11]: plt.show()
```



Scatter plot





Box plot

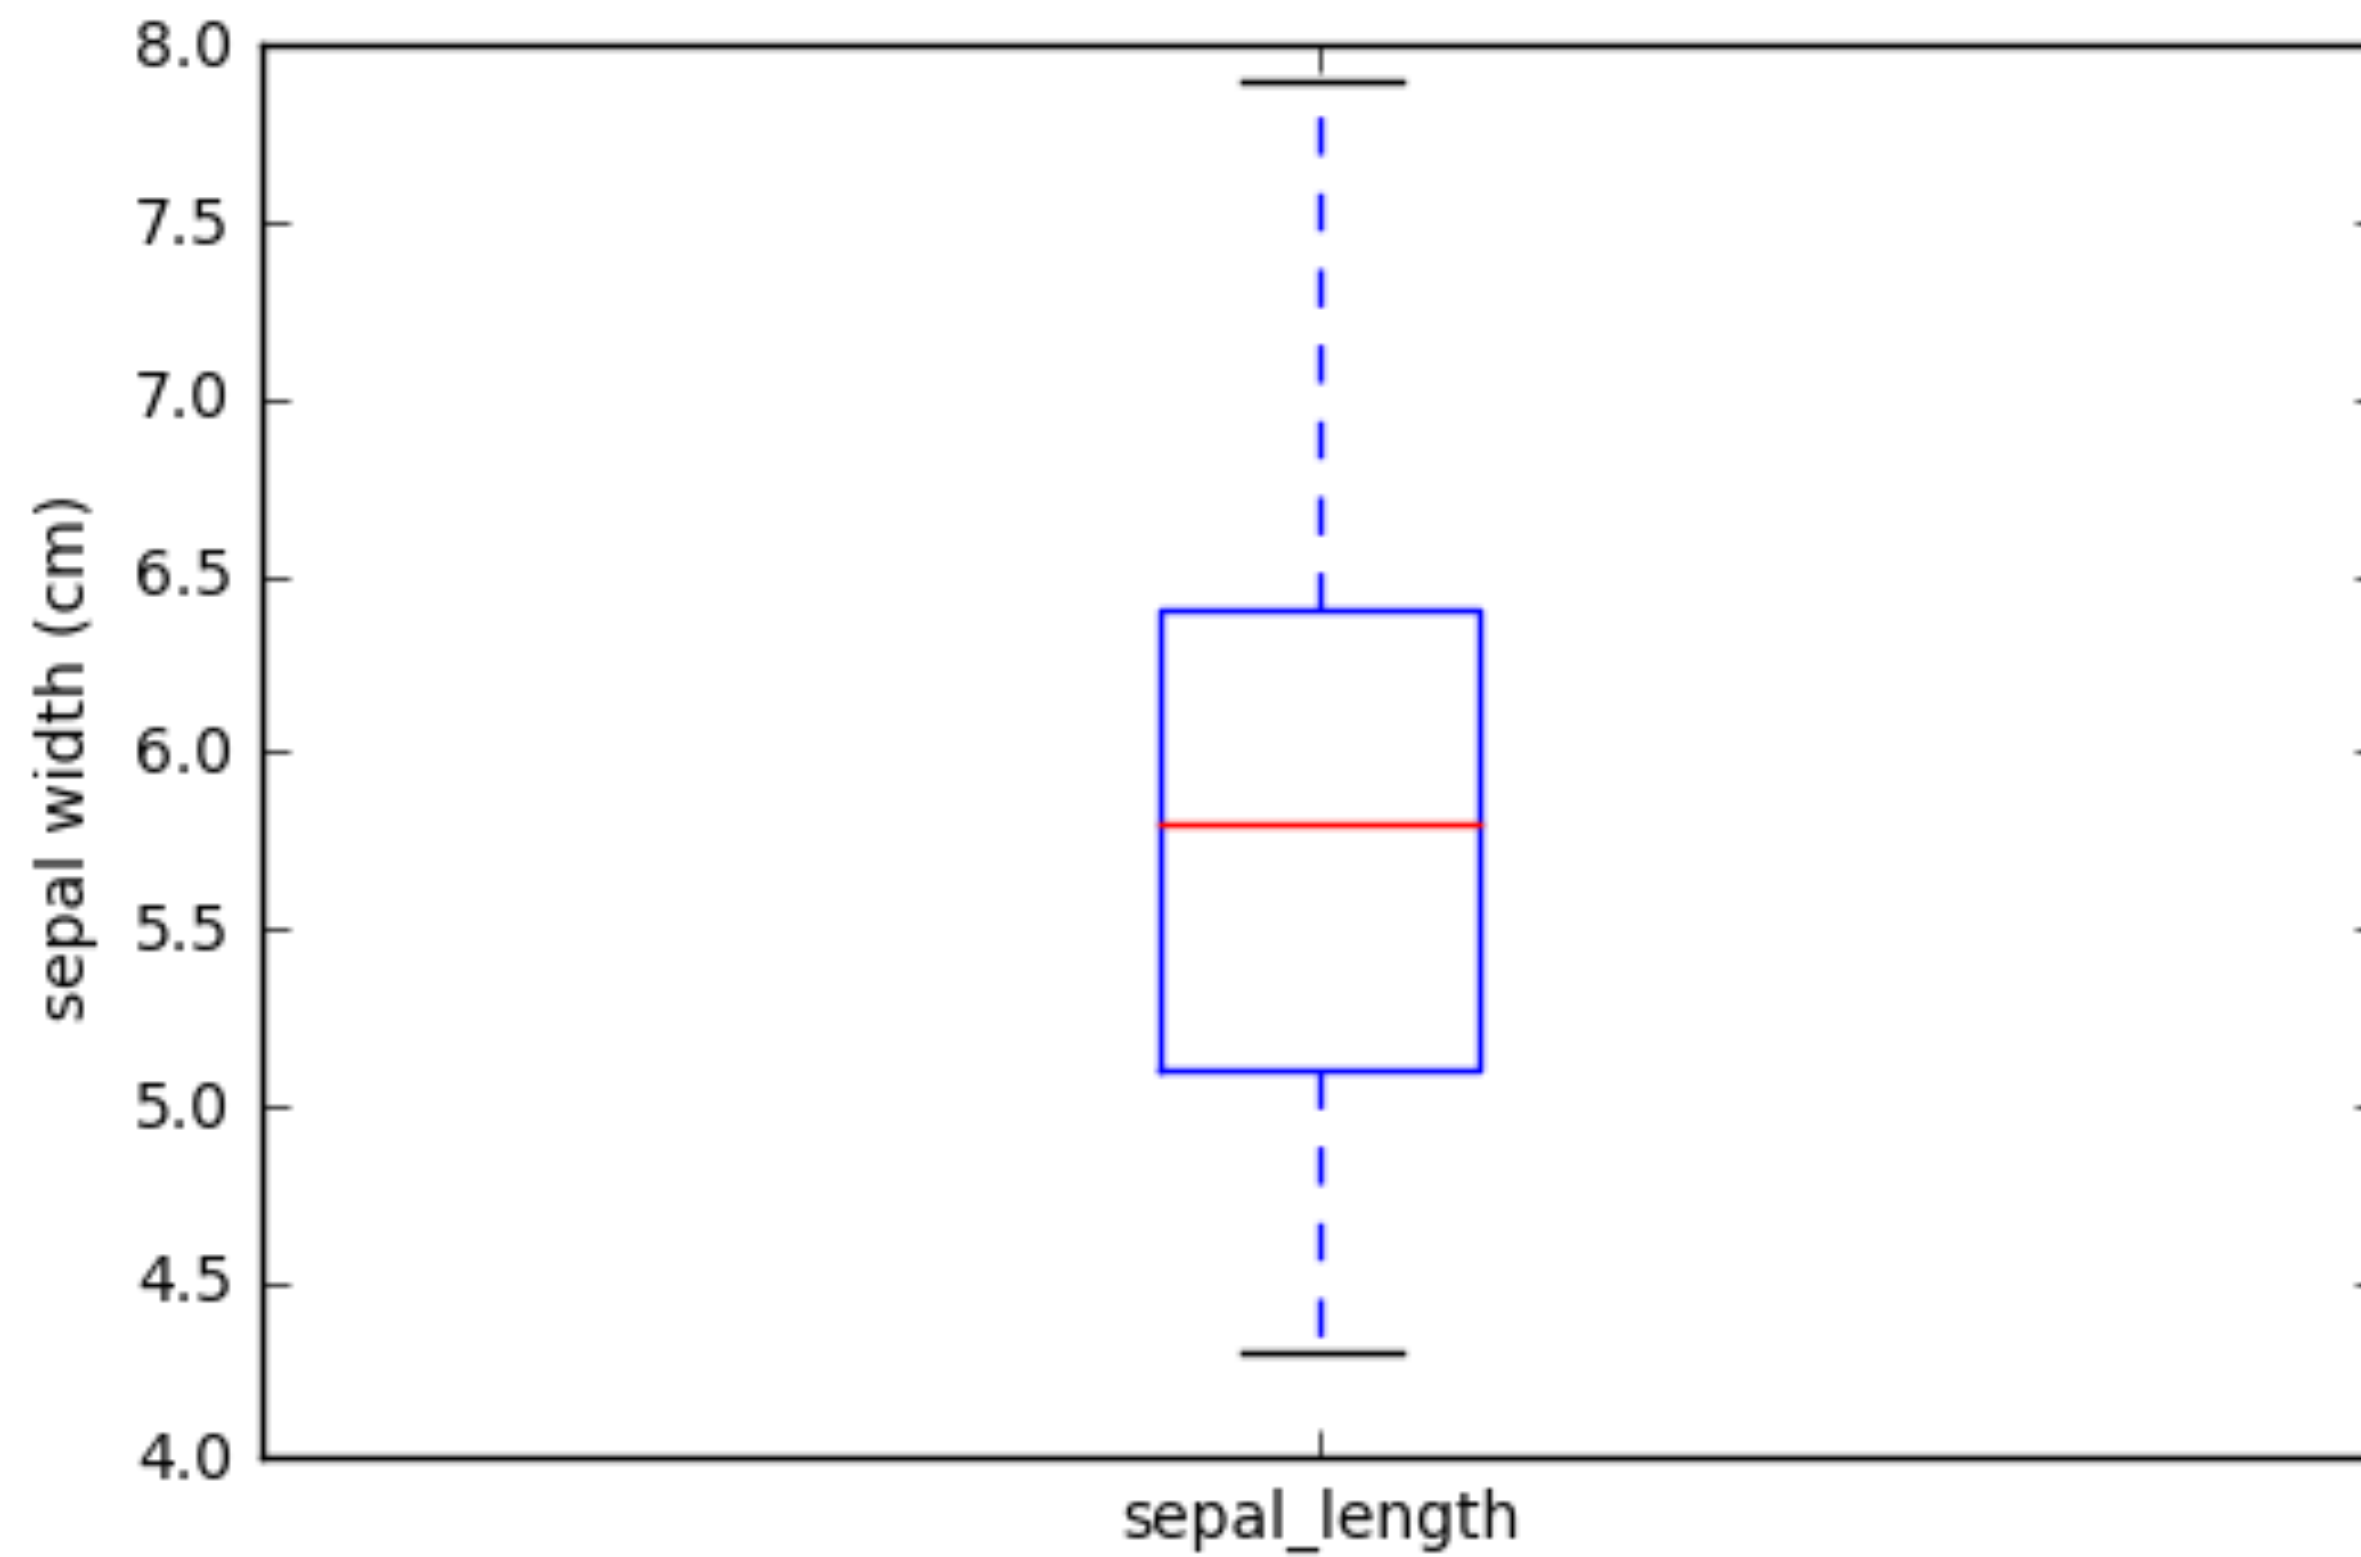
```
In [12]: iris.plot(y='sepal_length', kind='box')
```

```
In [13]: plt.ylabel('sepal width (cm)')
```

```
In [14]: plt.show()
```



Box plot





Histogram

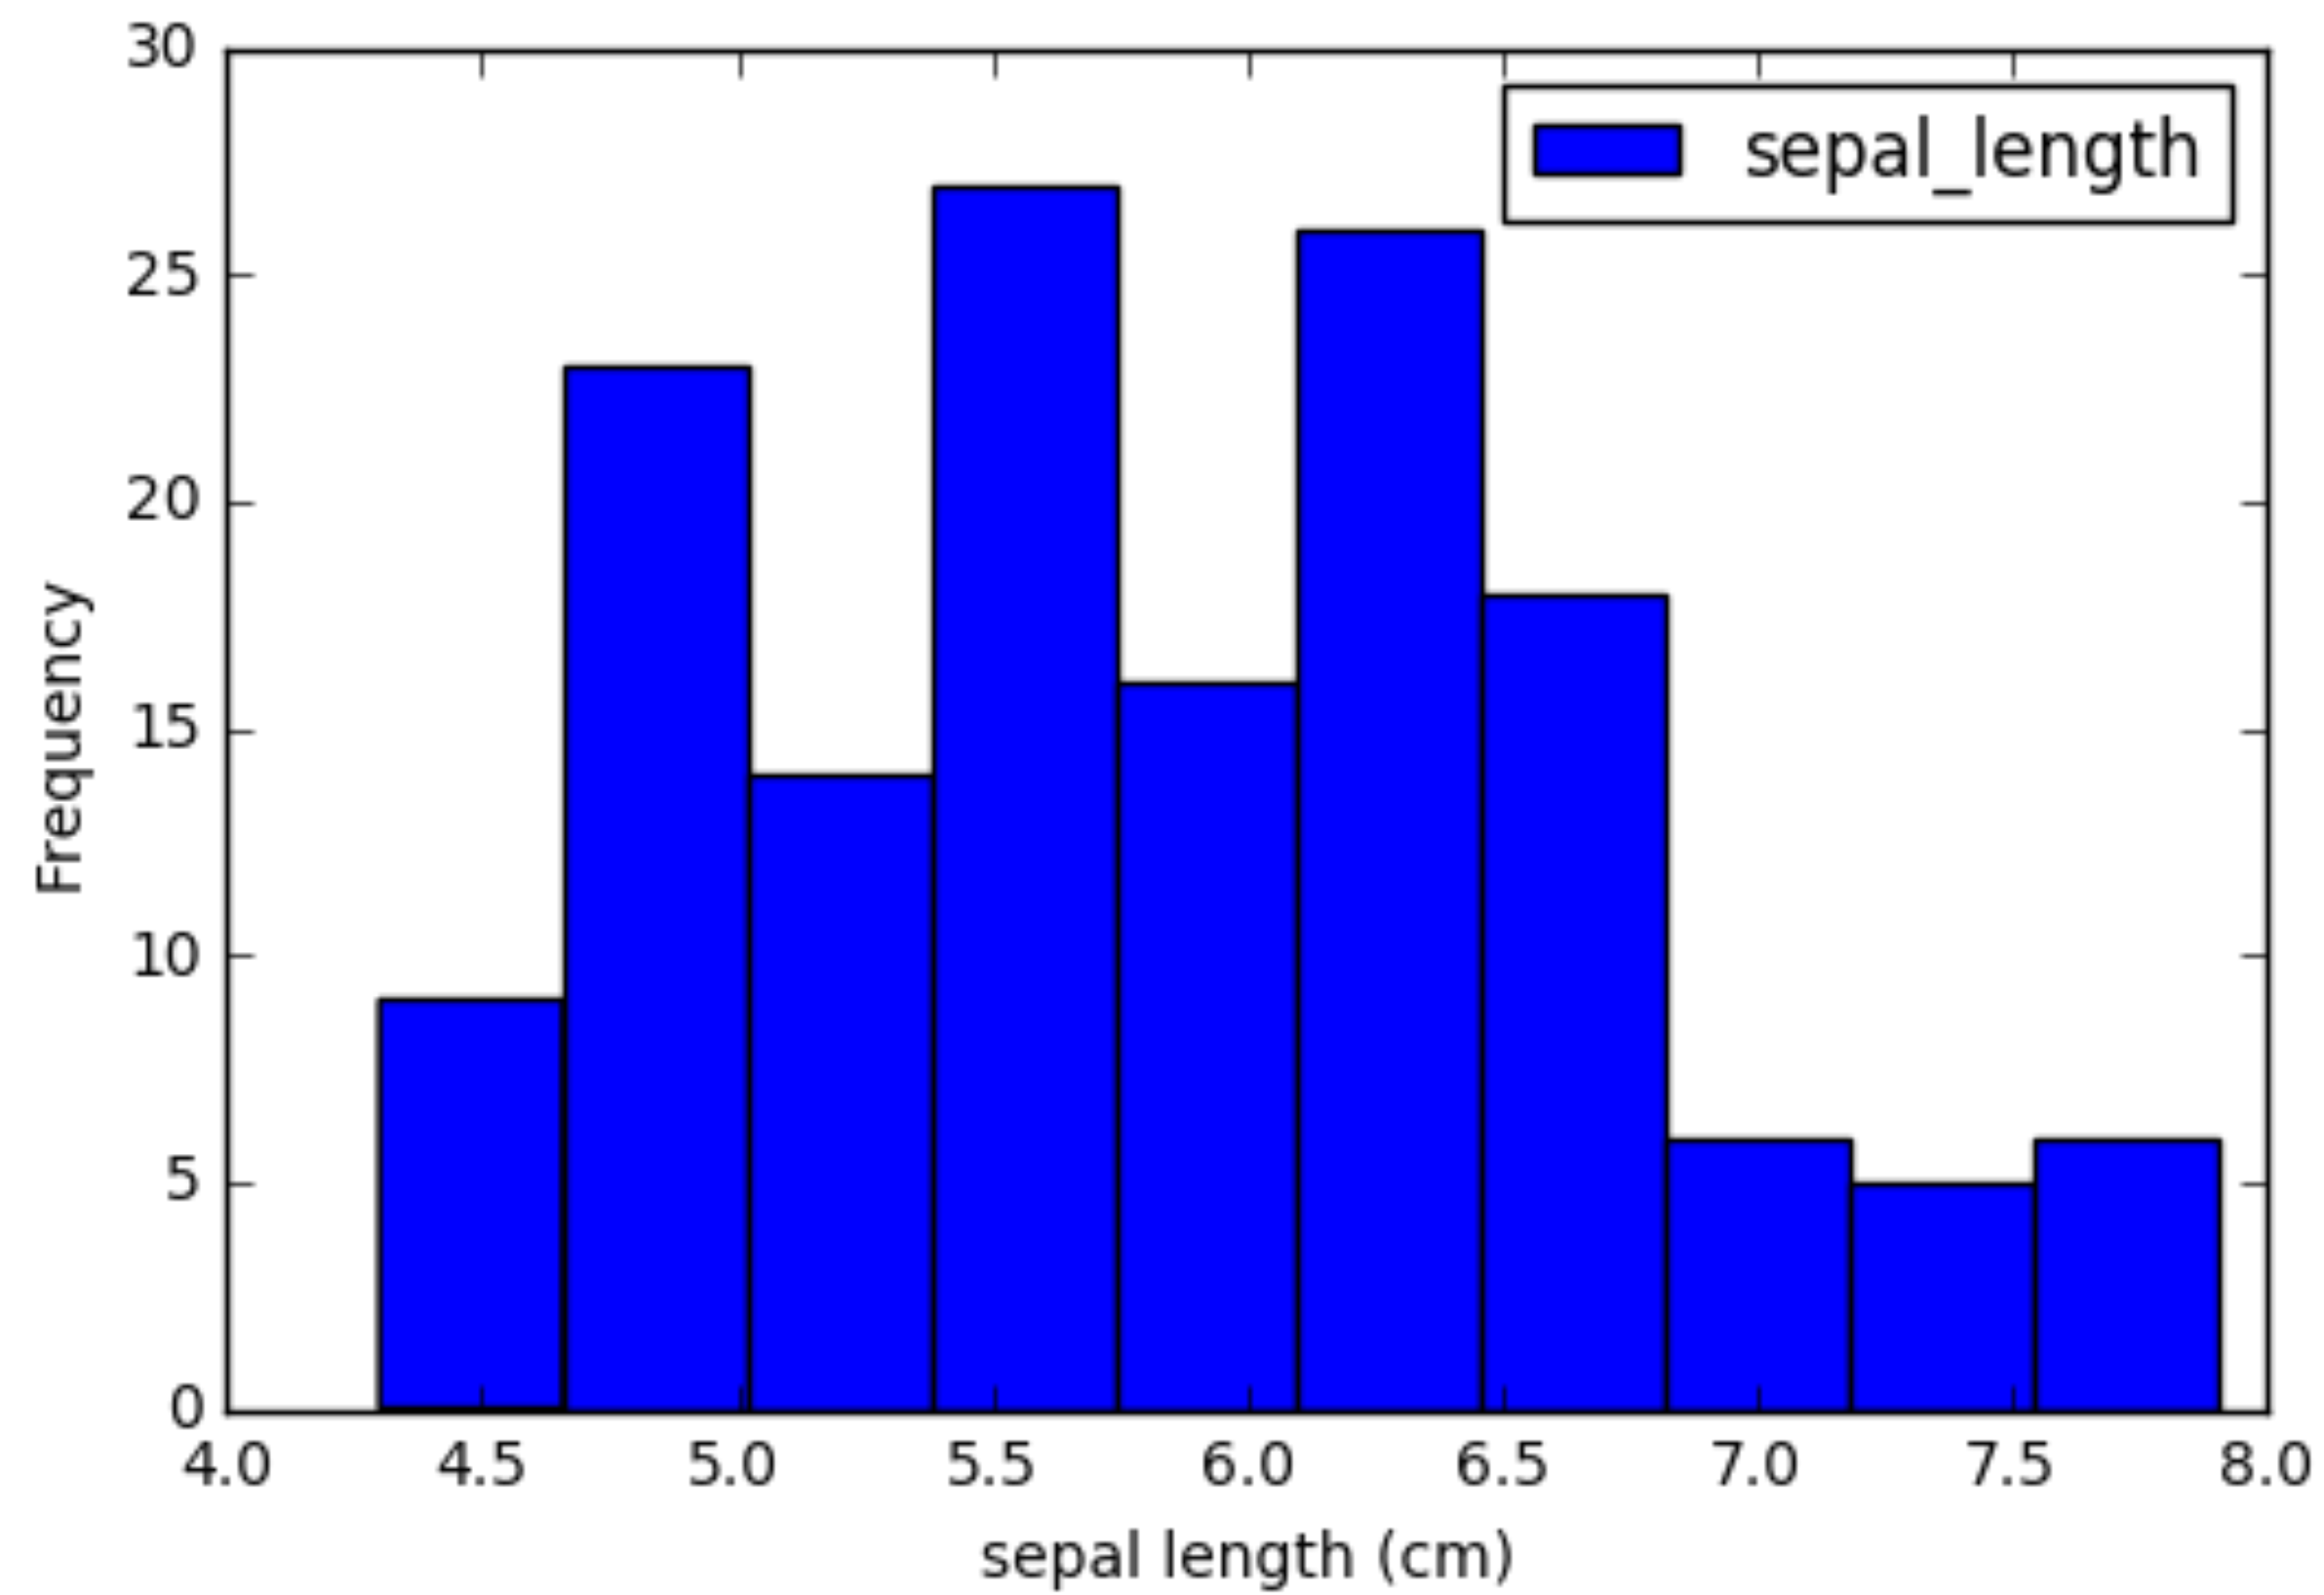
```
In [15]: iris.plot(y='sepal_length', kind='hist')
```

```
In [16]: plt.xlabel('sepal length (cm)')
```

```
In [17]: plt.show()
```



Histogram





Histogram options

- *bins* (integer): number of intervals or bins
- *range* (tuple): extrema of bins (minimum, maximum)
- *normed* (boolean): whether to normalize to one
- *cumulative* (boolean): compute Cumulative Distribution Function (CDF)
- ... more Matplotlib customizations



Customizing histogram

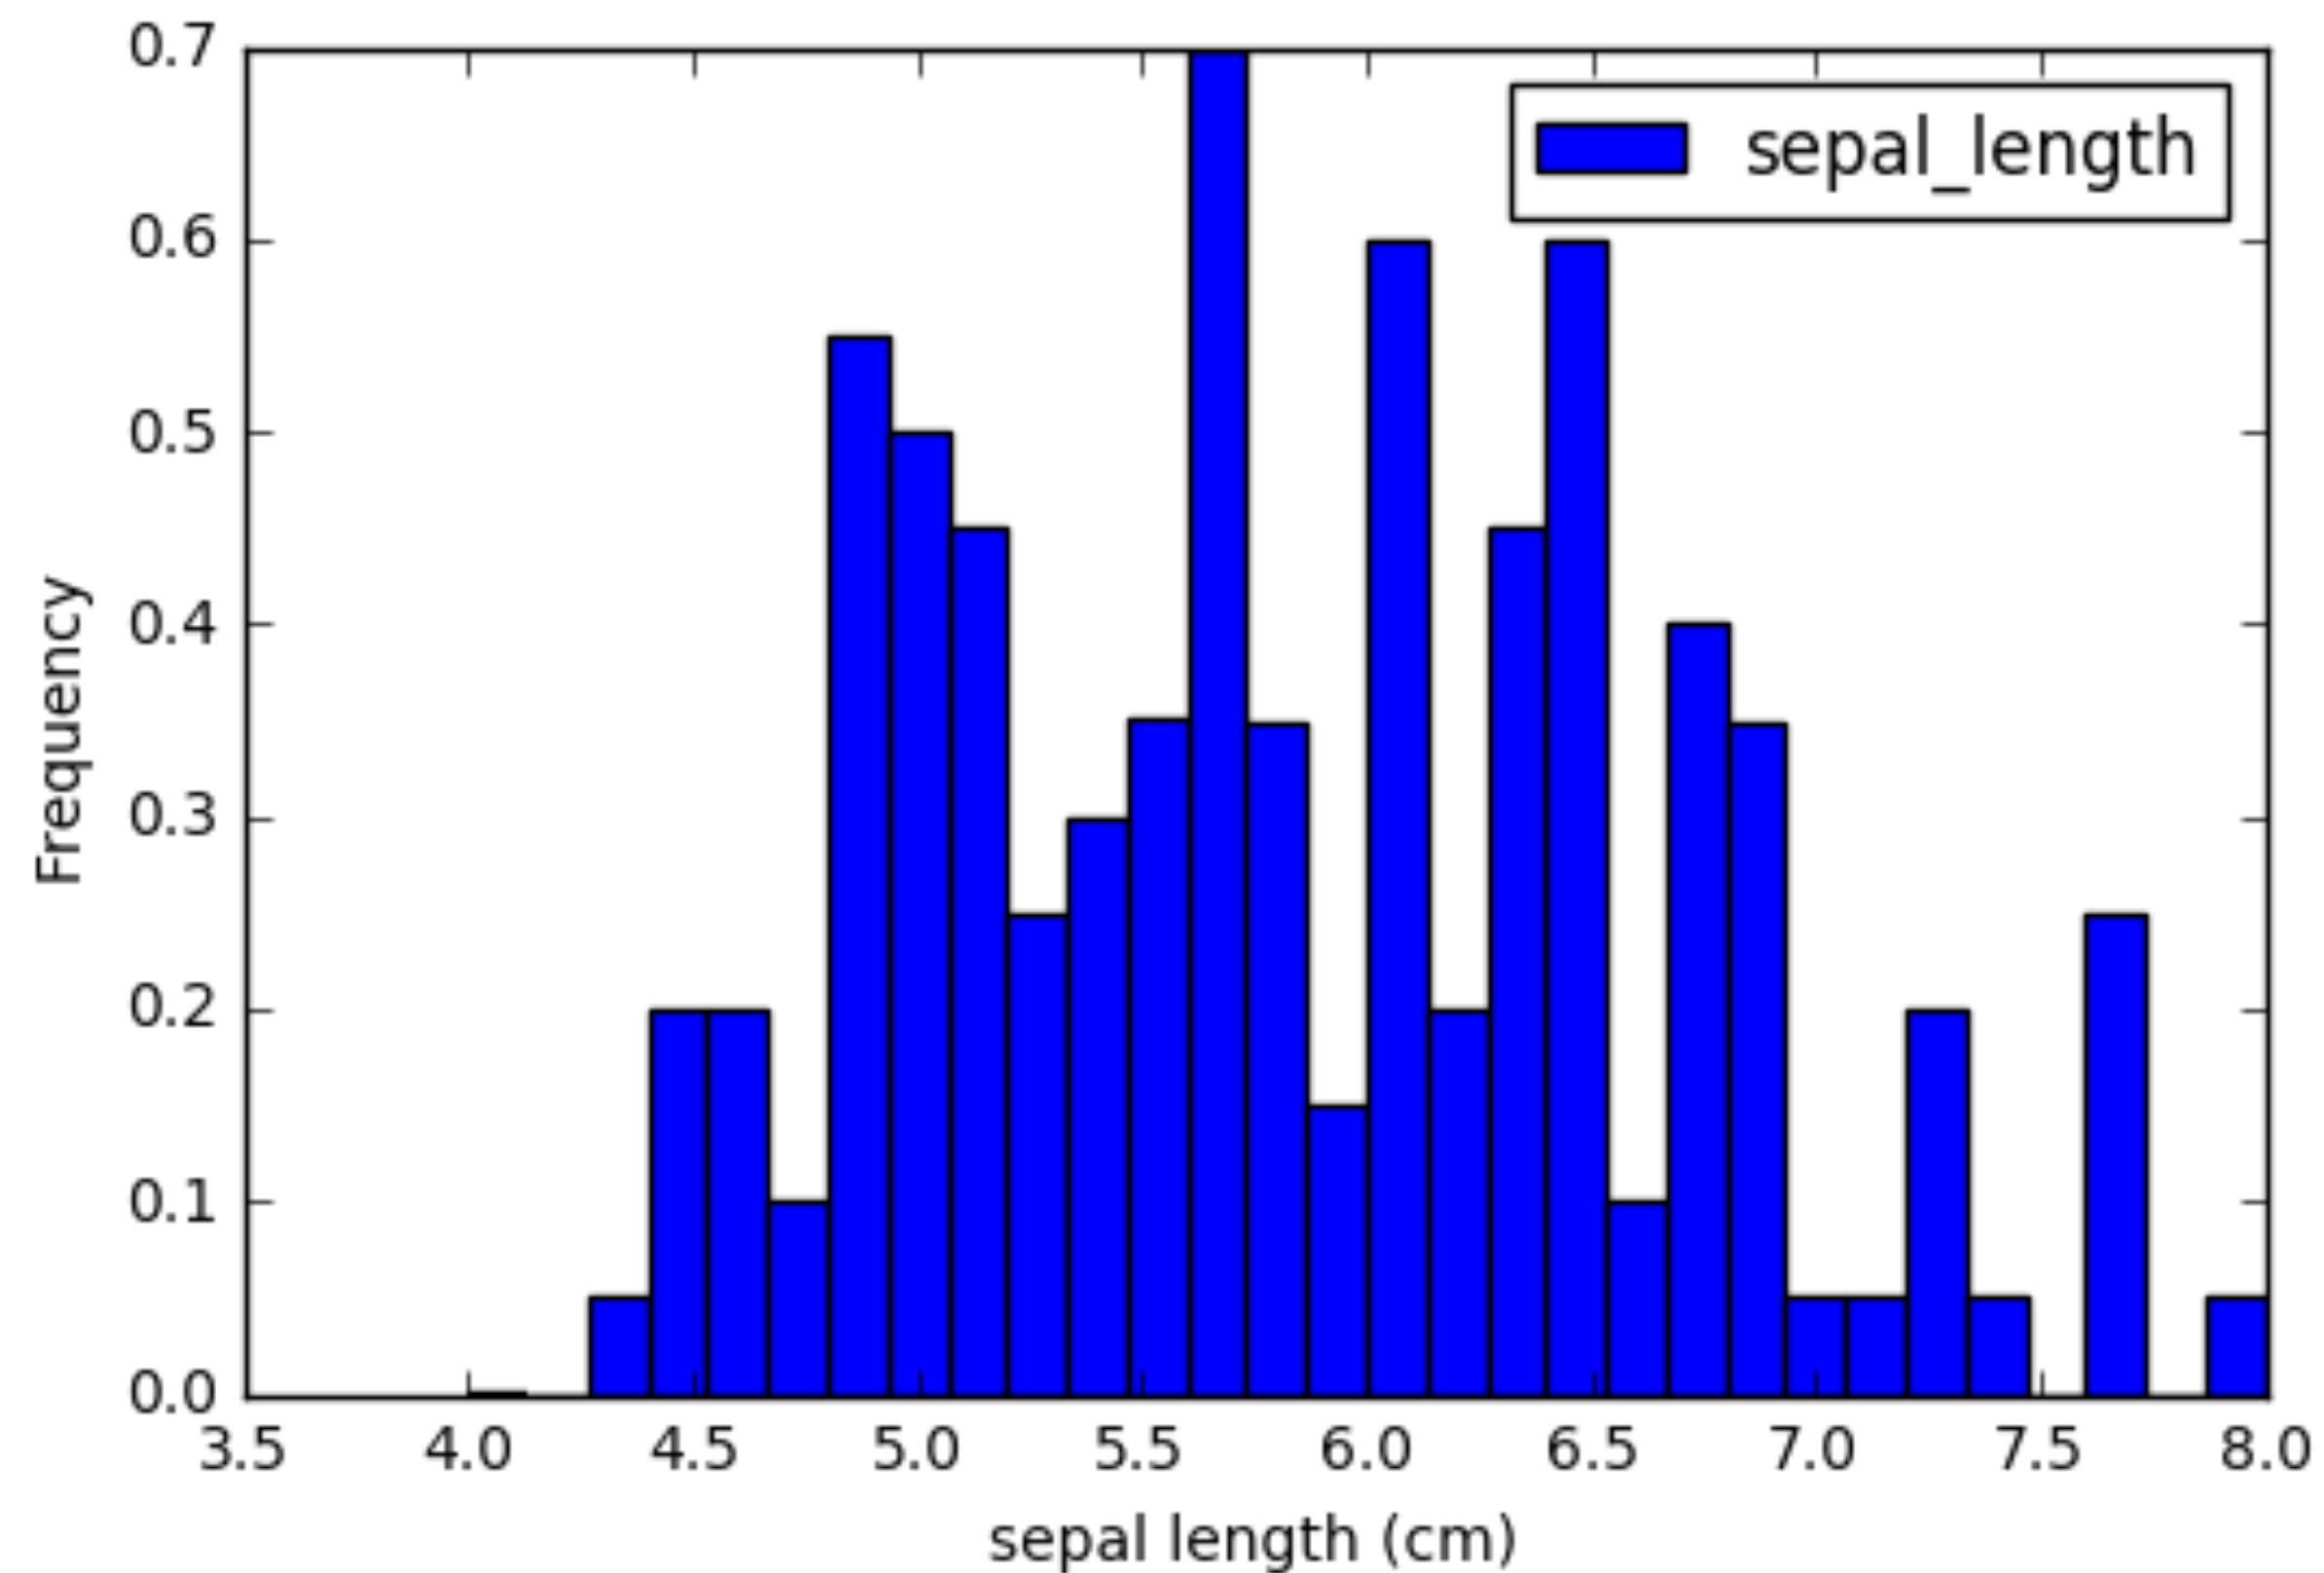
```
In [18]: iris.plot(y='sepal_length', kind='hist',  
....:             bins=30, range=(4,8), normed=True)
```

```
In [19]: plt.xlabel('sepal length (cm)')
```

```
In [20]: plt.show()
```



Customizing histogram






Cumulative distribution

```
In [21]: iris.plot(y='sepal_length', kind='hist', bins=30,  
....:             range=(4,8), cumulative=True, normed=True)  
  
In [22]: plt.xlabel('sepal length (cm)')  
  
In [23]: plt.title('Cumulative distribution function (CDF)')  
  
In [24]: plt.show()
```



Cumulative distribution





Word of warning

- Three different DataFrame plot idioms
 - *iris.plot(kind='hist')*
 - *iris.plt.hist()*
 - *iris.hist()*
- Syntax/results differ!
- Pandas API still evolving: check documentation!



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Statistical exploratory data analysis



Summarizing with describe()

```
In [1]: iris.describe() # summary statistics
```

```
Out[1]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000



Describe

- *count*: number of entries
- *mean*: average of entries
- *std*: standard deviation
- *min*: minimum entry
- *25%*: first quartile
- *50%*: median or second quartile
- *75%*: third quartile
- *max*: maximum entry



Counts

```
In [2]: iris['sepal_length'].count() # Applied to Series
Out[2]: 150
```

```
In [3]: iris['sepal_width'].count() # Applied to Series
Out[3]: 150
```

```
In [4]: iris[['petal_length', 'petal_width']].count() # Applied
....: to DataFrame
Out[4]:
petal_length    150
petal_width     150
dtype: int64
```

```
In [5]: type(iris[['petal_length', 'petal_width']].count()) #
....: returns Series
Out[5]: pandas.core.series.Series
```



Averages

```
In [6]: iris['sepal_length'].mean() # Applied to Series  
Out[6]: 5.8433333333333335
```

```
In [7]: iris.mean() # Applied to entire DataFrame  
Out[7]:  
sepal_length      5.843333  
sepal_width       3.057333  
petal_length      3.758000  
petal_width       1.199333  
dtype: float64
```




Standard deviations

```
In [8]: iris.std()
Out[8]:
sepal_length    0.828066
sepal_width     0.435866
petal_length    1.765298
petal_width     0.762238
dtype: float64
```



Mean and standard deviation on a bell curve





Medians

```
In [9]: iris.median()  
Out[9]:  
sepal_length    5.80  
sepal_width     3.00  
petal_length    4.35  
petal_width     1.30  
dtype: float64
```



Medians & 0.5 quantiles

```
In [10]: iris.median()
```

```
Out[10]:
```

```
sepal_length    5.80  
sepal_width     3.00  
petal_length    4.35  
petal_width     1.30  
dtype: float64
```

```
In [11]: q = 0.5
```

```
In [12]: iris.quantile(q)
```

```
Out[12]:
```

```
sepal_length    5.80  
sepal_width     3.00  
petal_length    4.35  
petal_width     1.30  
dtype: float64
```



Inter-quartile range (IQR)

```
In [13]: q = [0.25, 0.75]
```

```
In [14]: iris.quantile(q)
```

```
Out[14]:
```

	sepal_length	sepal_width	petal_length	petal_width
0.25	5.1	2.8	1.6	0.3
0.75	6.4	3.3	5.1	1.8



Ranges

```
In [15]: iris.min()
```

```
Out[15]:
```

```
sepal_length    4.3
sepal_width      2
petal_length     1
petal_width     0.1
species         setosa
dtype: object
```

```
In [16]: iris.max()
```

```
Out[16]:
```

```
sepal_length    7.9
sepal_width     4.4
petal_length     6.9
petal_width     2.5
species         virginica
dtype: object
```



Box plots

```
In [17]: iris.plot(kind= 'box')
```

```
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x118a3d5f8>
```


```
In [18]: plt.ylabel(' [cm] ')
```

```
Out[18]: <matplotlib.text.Text at 0x118a524e0>
```

```
In [19]: plt.show()
```



Box plots





Percentiles as quantiles

```
In [20]: iris.describe() # summary statistics
```

```
Out[20]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Separating populations



```
In [1]: iris.head()
```

```
Out[1]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa



Describe species column

```
In [2]: iris['species'].describe()
```

```
Out[2]:
```

```
count          150  
unique           3  
top          setosa  
freq           50
```

```
Name: species, dtype: object
```

count: # non-null entries

unique: # distinct values

top: most frequent category

freq: # occurrences of top



Unique & factors

```
In [3]: iris['species'].unique()  
Out[3]: array(['setosa', 'versicolor', 'virginica'], dtype=object)
```



Filtering by species

```
In [4]: indices = iris['species'] == 'setosa'
```

```
In [5]: setosa = iris.loc[indices,:] # extract new DataFrame
```

```
In [6]: indices = iris['species'] == 'versicolor'
```

```
In [7]: versicolor = iris.loc[indices,:] # extract new DataFrame
```

```
In [8]: indices = iris['species'] == 'virginica'
```

```
In [9]: virginica = iris.loc[indices,:] # extract new DataFrame
```



Checking species

```
In [10]: setosa['species'].unique()  
Out[10]: array(['setosa'], dtype=object)
```

```
In [11]: versicolor['species'].unique()  
Out[11]: array(['versicolor'], dtype=object)
```

```
In [12]: virginica['species'].unique()  
Out[12]: array(['virginica'], dtype=object)
```

```
In [13]: del setosa['species'], versicolor['species'],  
....:     virginica['species']
```



Checking indexes

```
In [14]: setosa.head(2)
```

```
Out[14]:
```

	sepal_length	sepal_width	petal_length	petal_width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2

```
In [15]: versicolor.head(2)
```

```
Out[15]:
```

	sepal_length	sepal_width	petal_length	petal_width
50	7.0	3.2	4.7	1.4
51	6.4	3.2	4.5	1.5

```
In [16]: virginica.head(2)
```

```
Out[16]:
```

	sepal_length	sepal_width	petal_length	petal_width
100	6.3	3.3	6.0	2.5
101	5.8	2.7	5.1	1.9



Visual EDA: all data

```
In [17]: iris.plot(kind= 'hist', bins=50, range=(0,8), alpha=0.3)
```

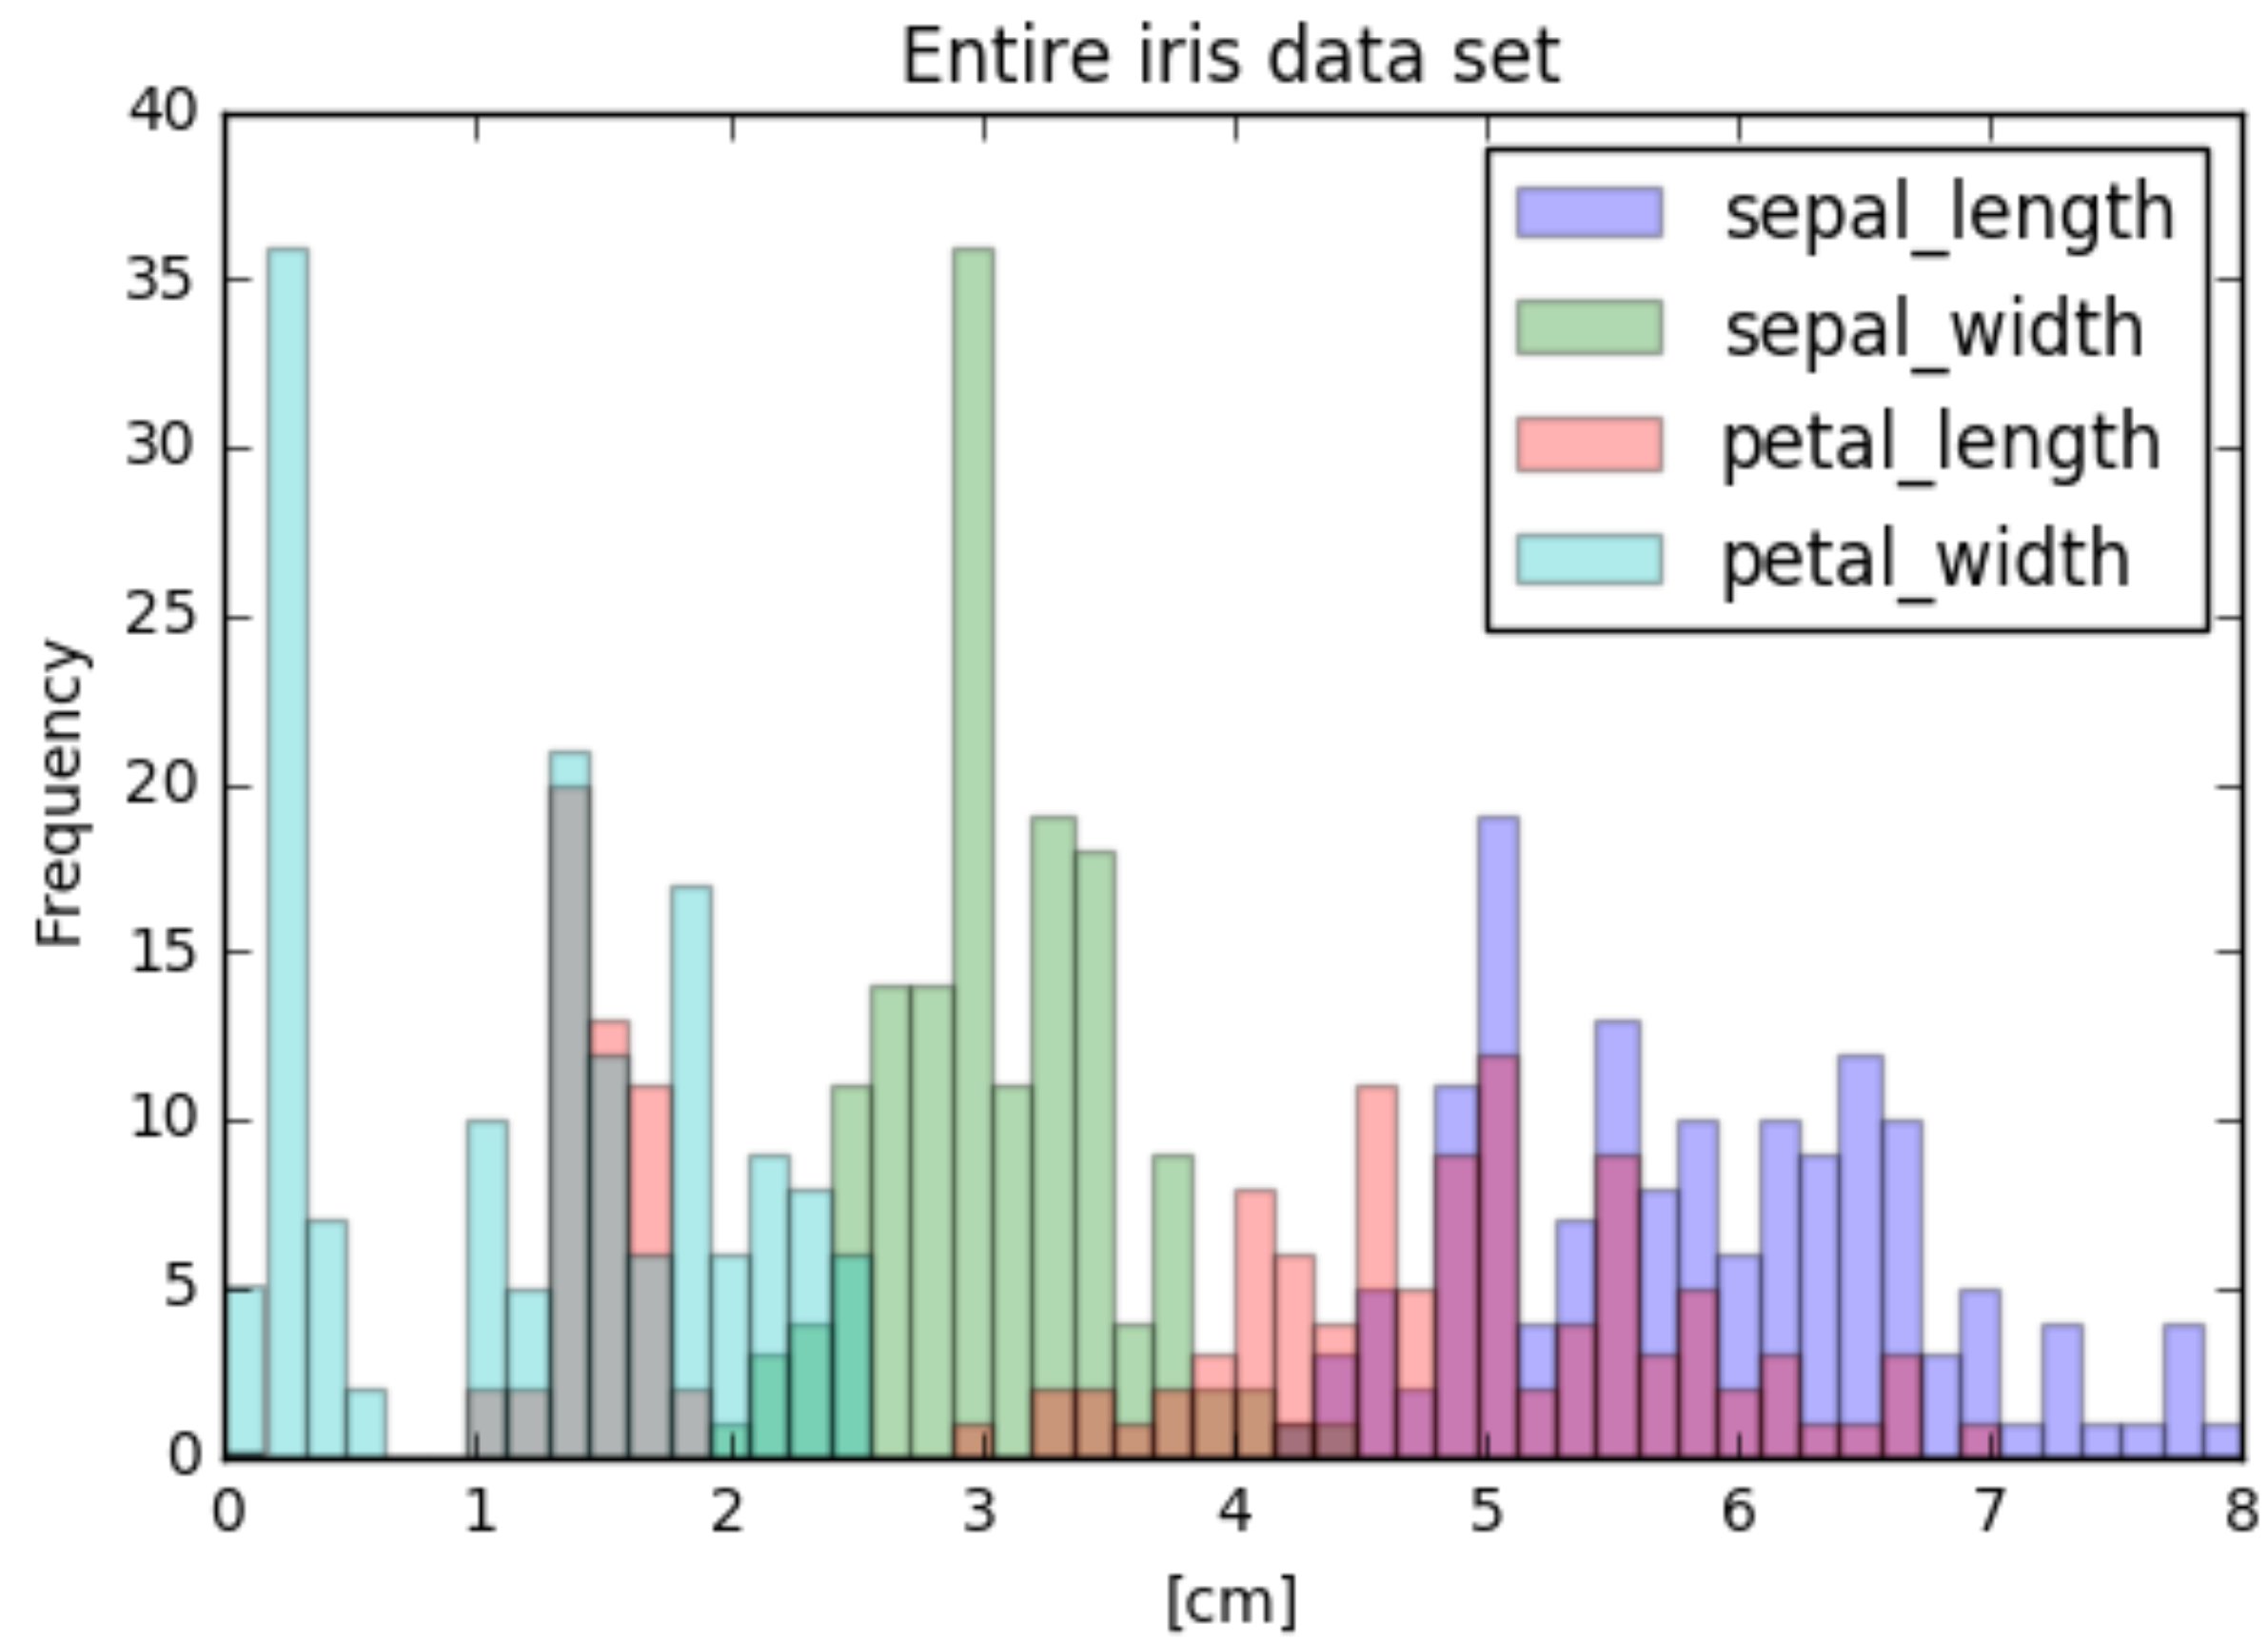
```
In [18]: plt.title('Entire iris data set')
```

```
In [19]: plt.xlabel(' [cm] ')
```

```
In [20]: plt.show()
```



Visual EDA: all data





Visual EDA: individual factors

```
In [21]: setosa.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)
```

```
In [22]: plt.title('Setosa data set')
```

```
In [23]: plt.xlabel('[cm]')
```

```
In [24]: versicolor.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)
```

```
In [25]: plt.title('Versicolor data set')
```

```
In [26]: plt.xlabel('[cm]')
```

```
In [27]: virginica.plot(kind='hist', bins=50, range=(0,8), alpha=0.3)
```

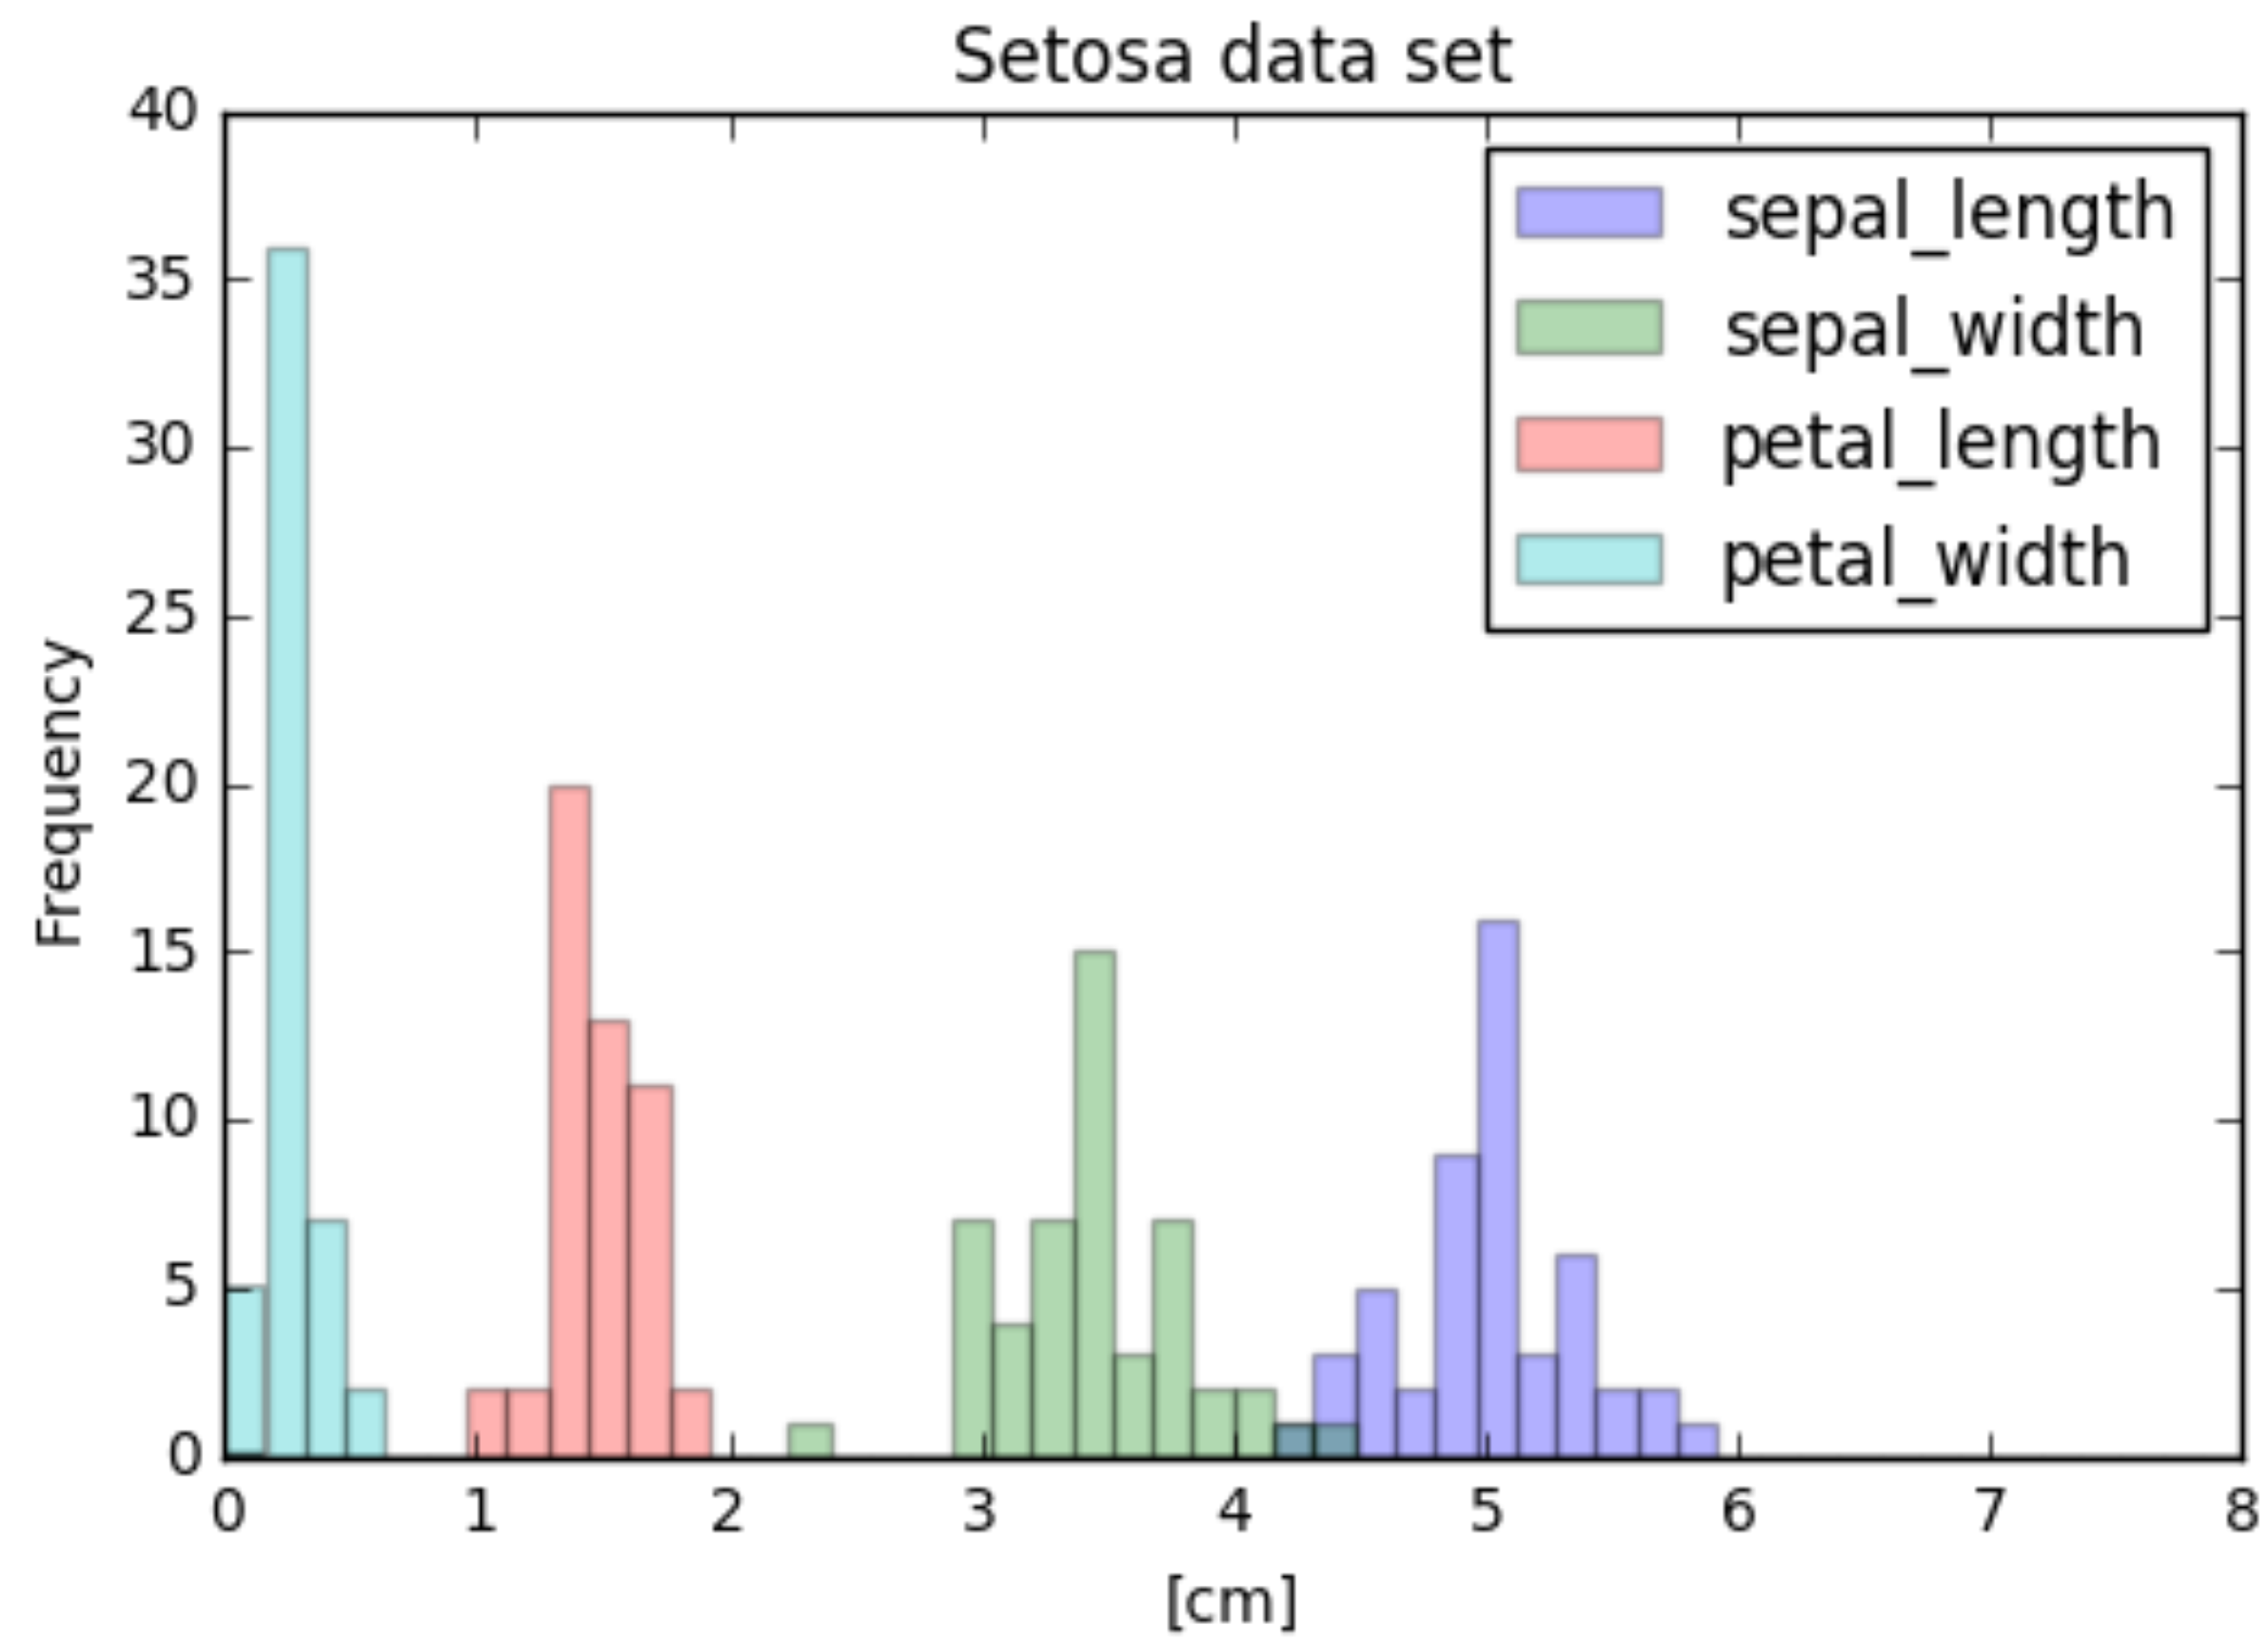
```
In [28]: plt.title('Virginica data set')
```

```
In [29]: plt.xlabel('[cm]')
```

```
In [30]: plt.show()
```

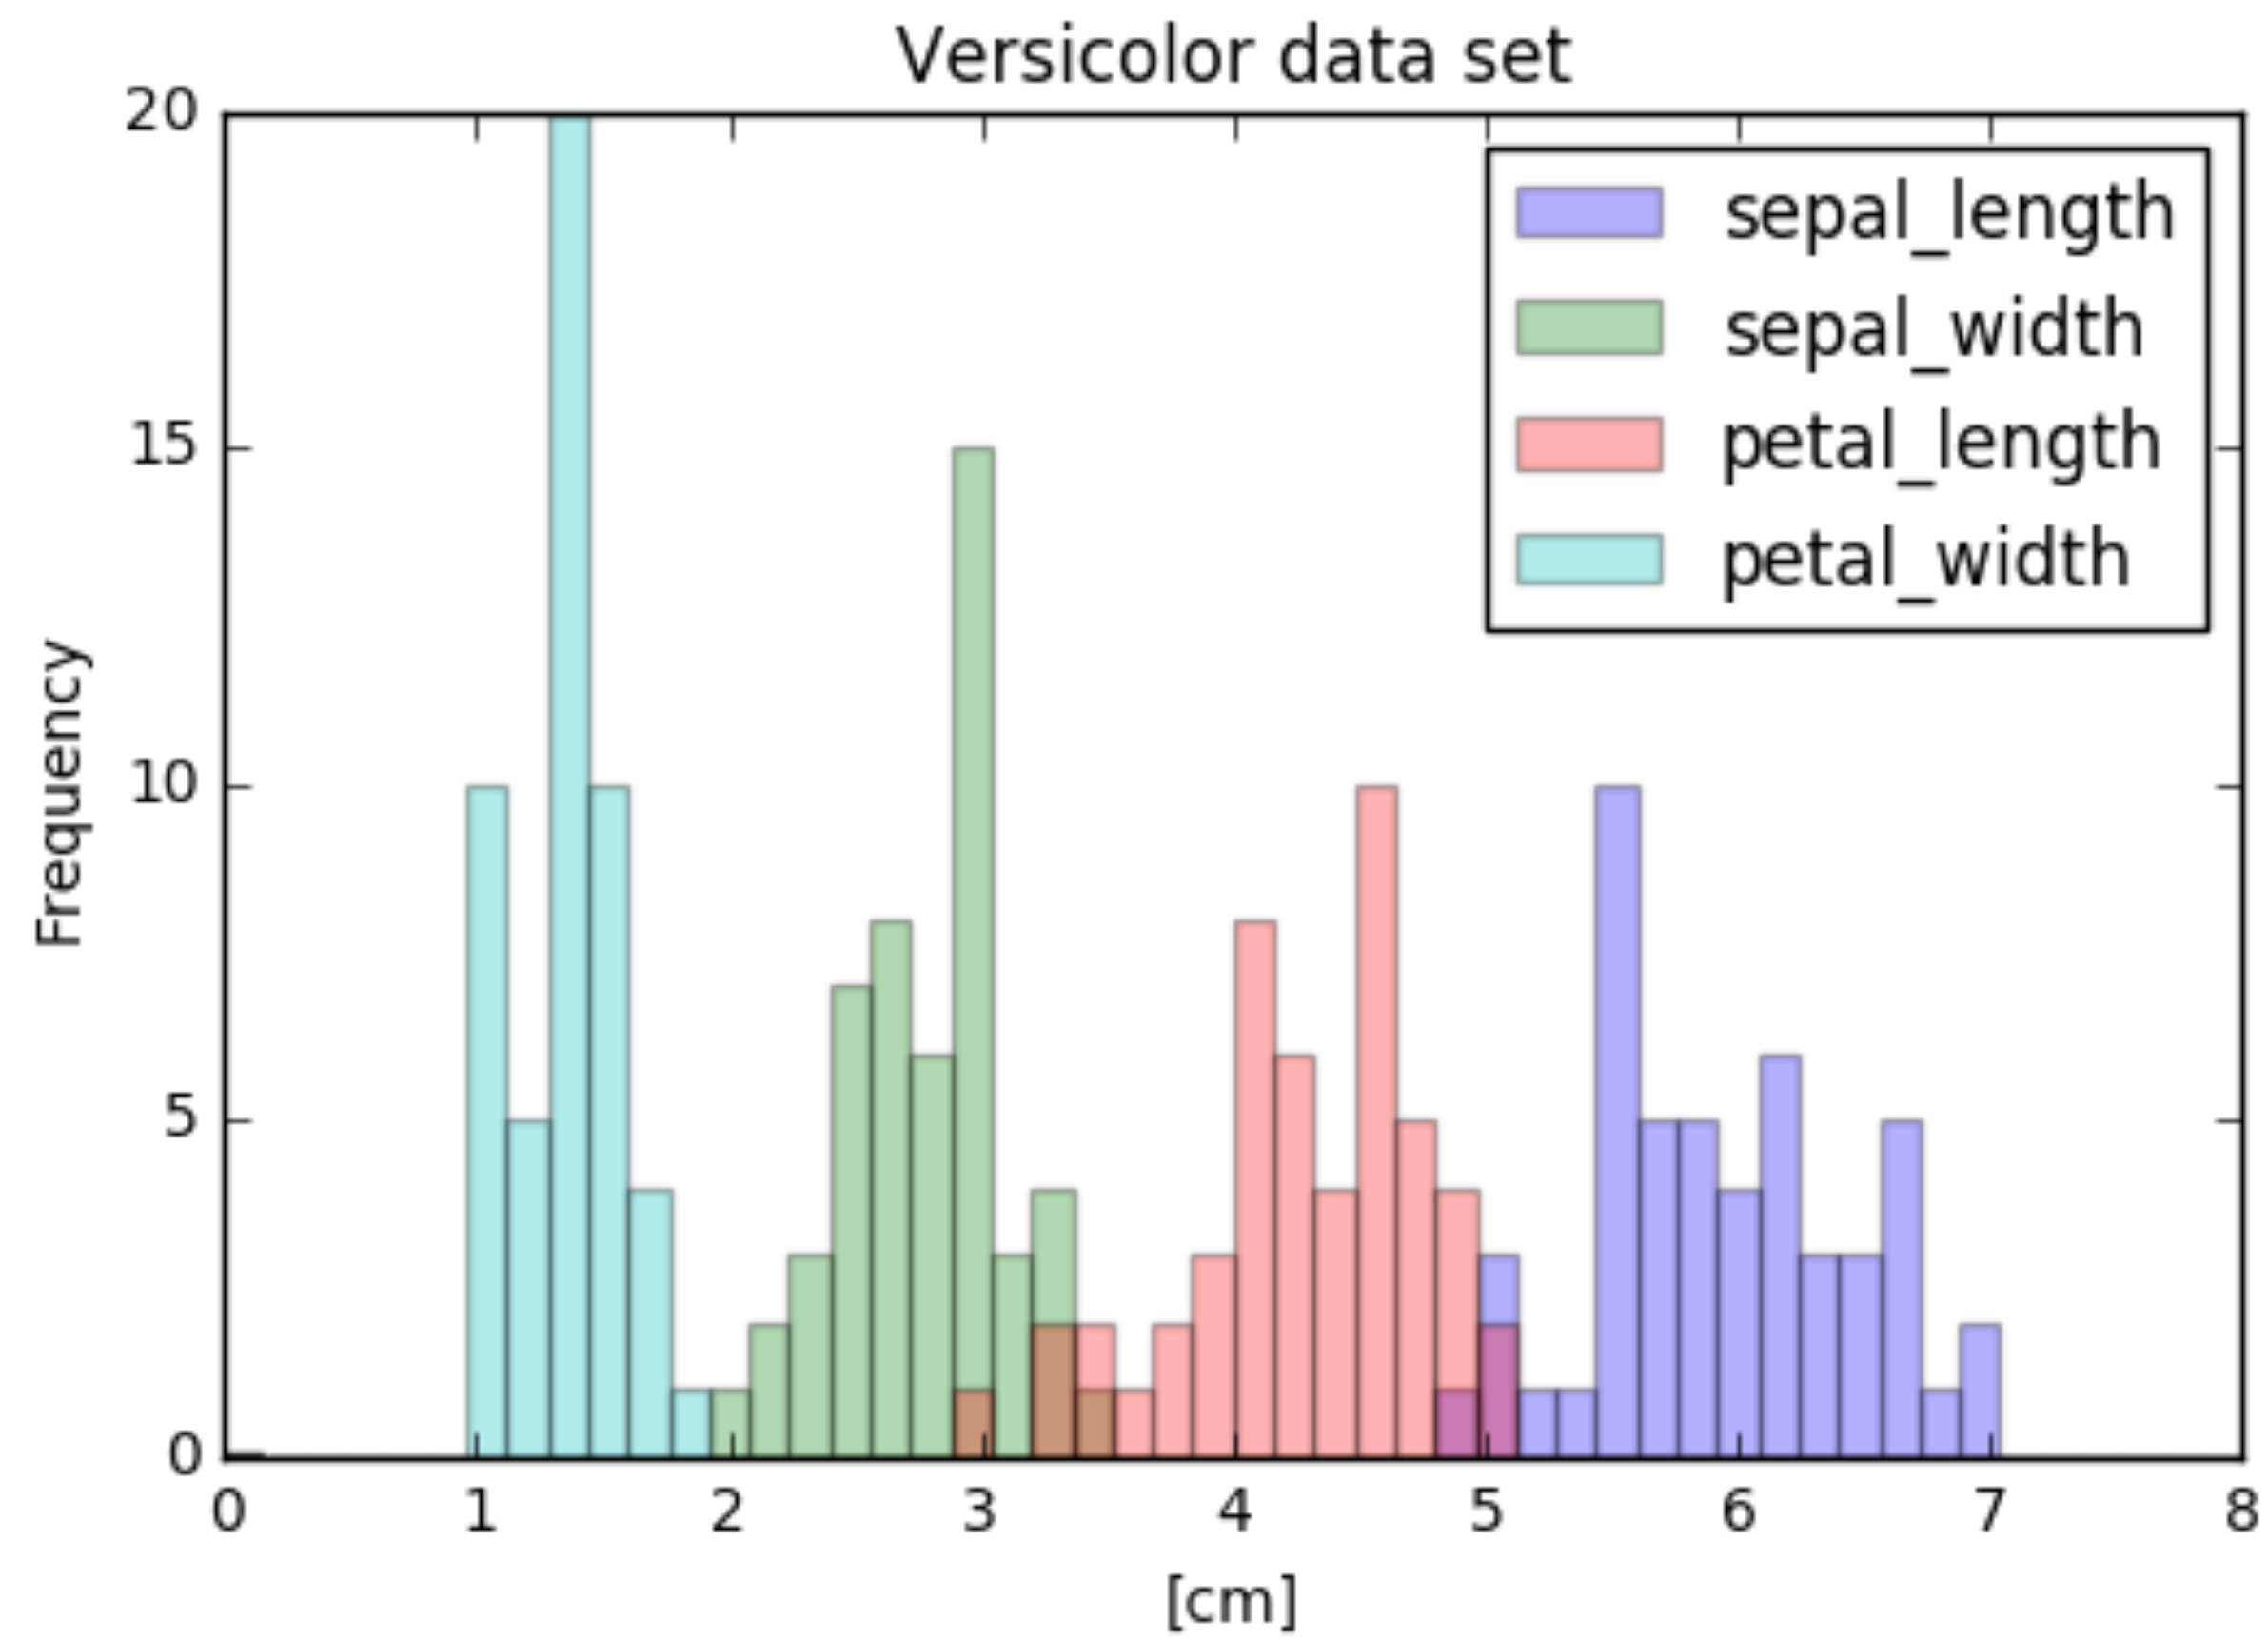


Visual EDA: Setosa data



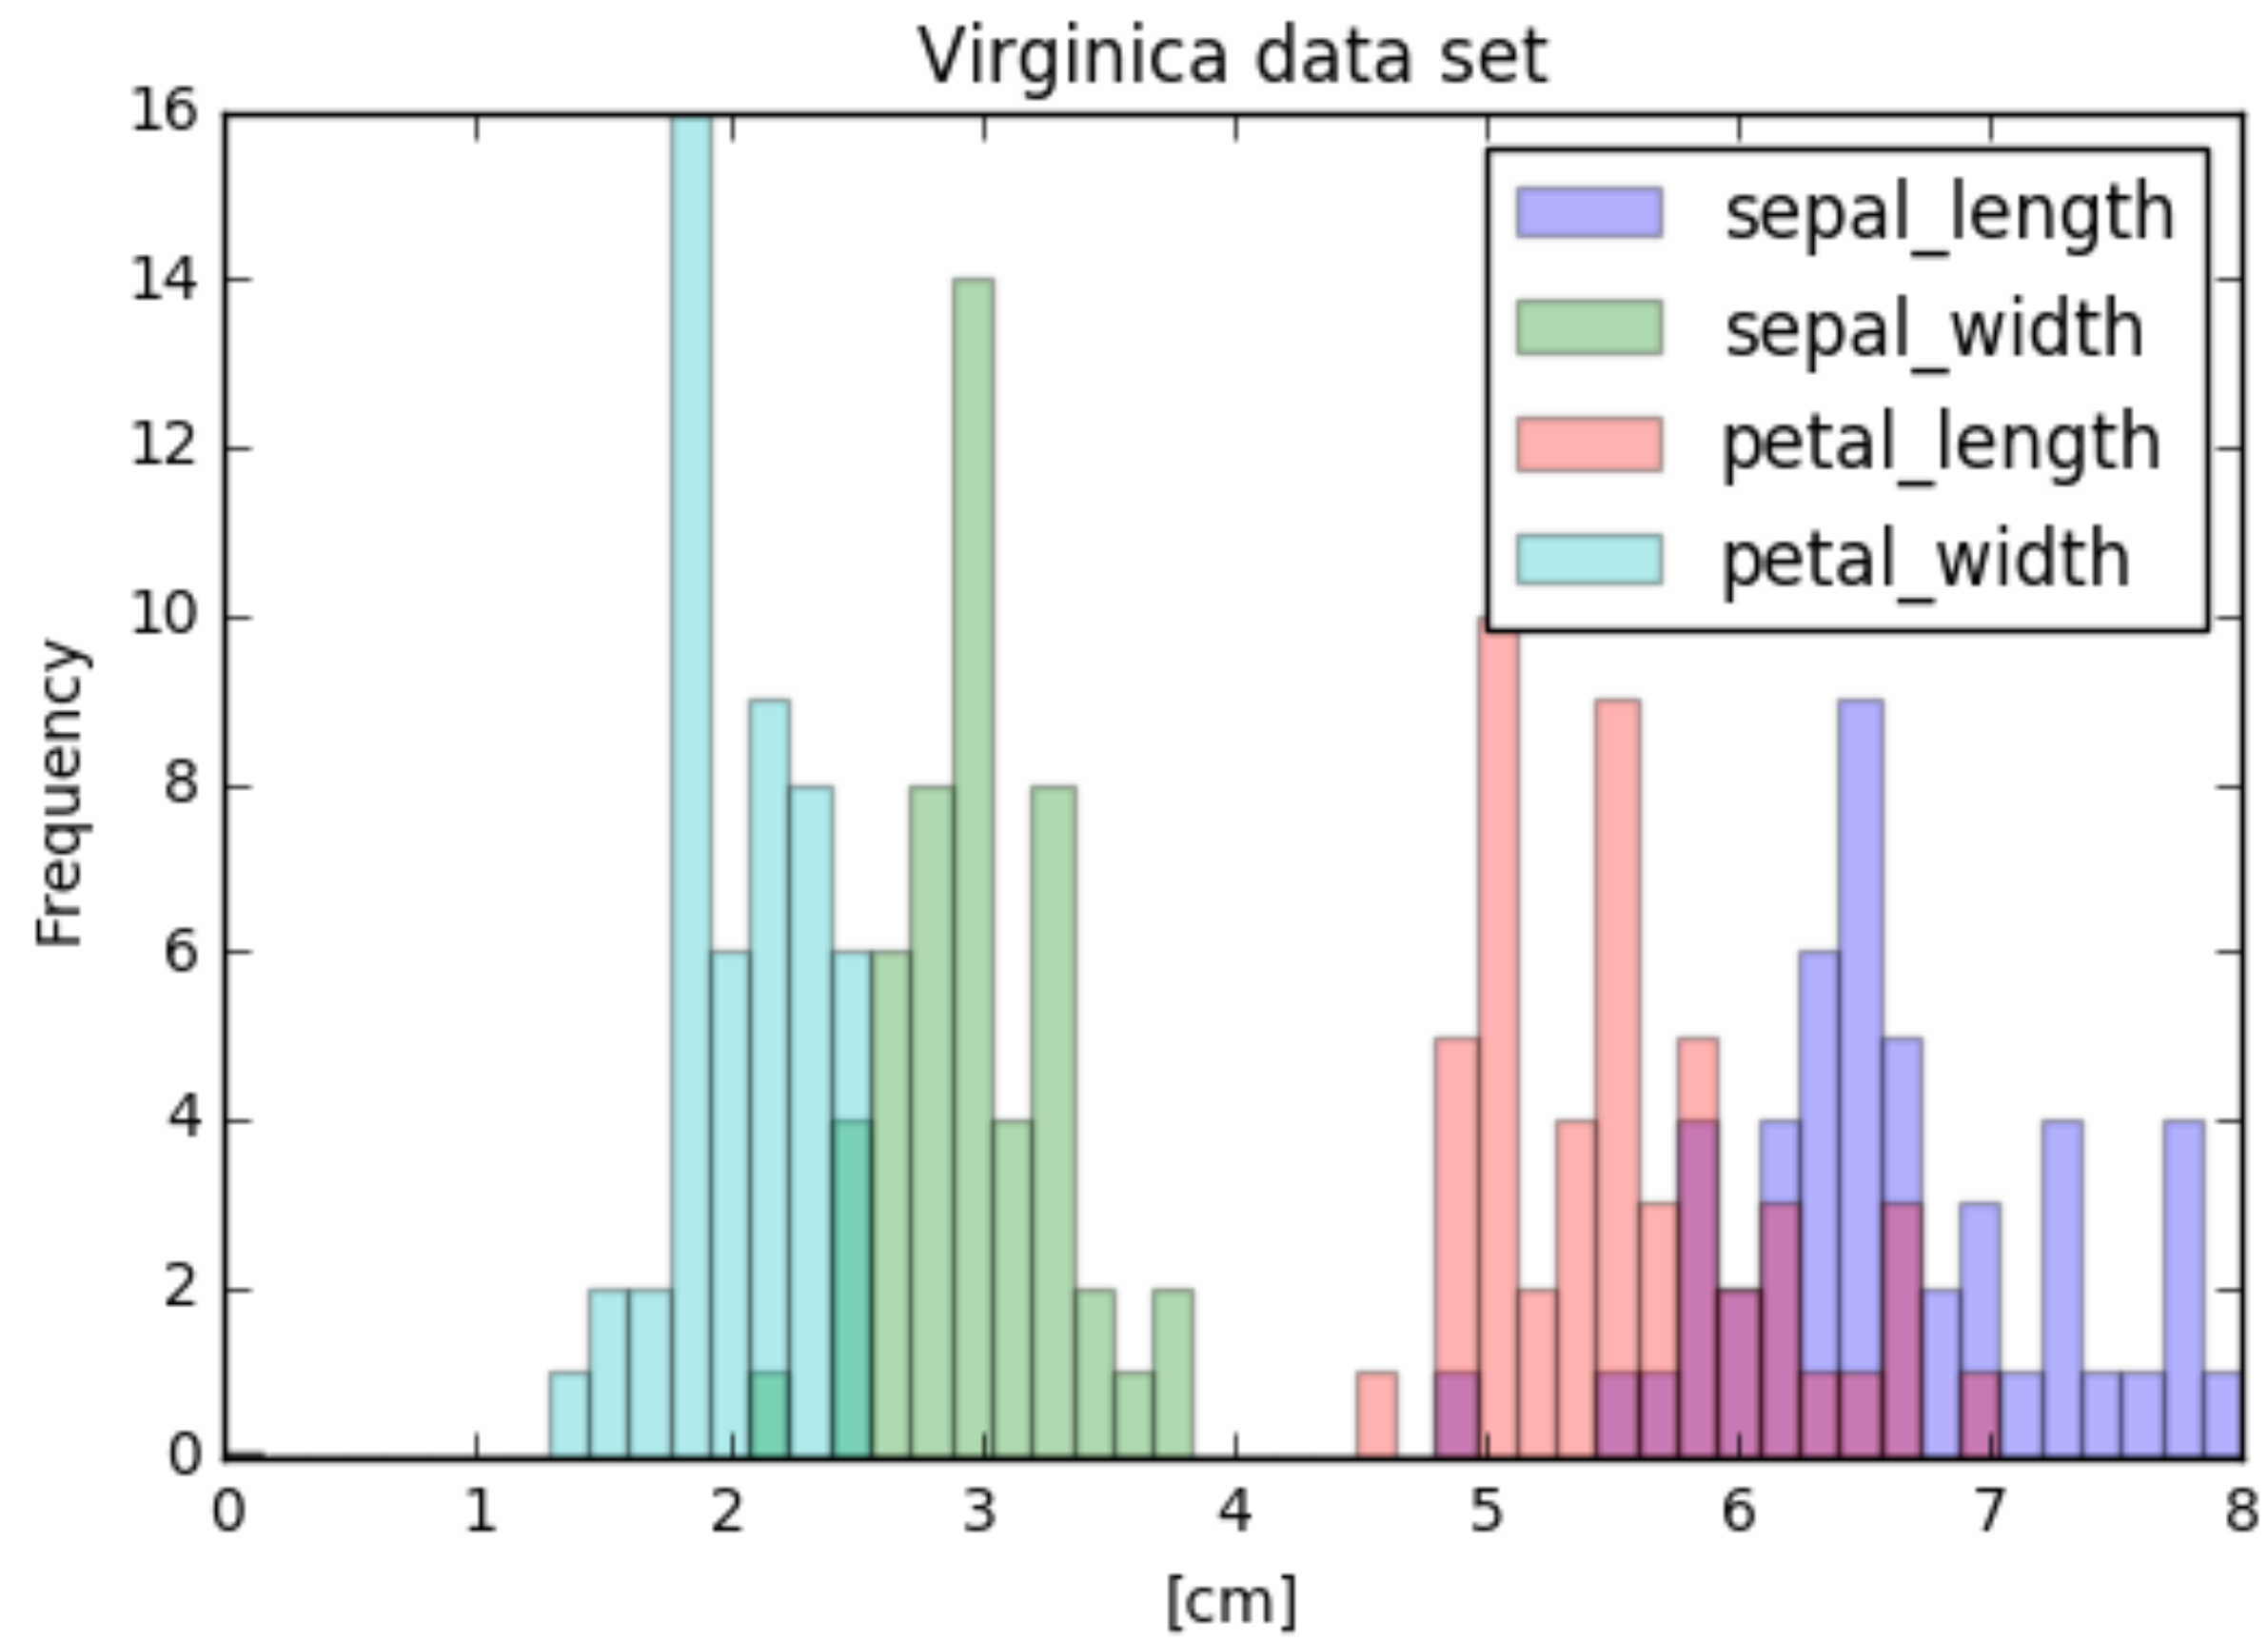


Visual EDA: Versicolor data





Visual EDA: Virginica data





Statistical EDA: describe()

```
In [31]: describe_all = iris.describe()
```

```
In [32]: print(describe_all)
```

```
Out[32]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
In [33]: describe_setosa = setosa.describe()
```

```
In [34]: describe_versicolor = versicolor.describe()
```

```
In [35]: describe_virginica = virginica.describe()
```



Computing errors

```
In [36]: error_setosa = 100 * np.abs(describe_setosa -  
....: describe_all)
```

```
In [37]: error_setosa = error_setosa/describe_setosa
```

```
In [38]: error_versicolor = 100 * np.abs(describe_versicolor -  
....: describe_all)
```

```
In [39]: error_versicolor = error_versicolor/describe_versicolor
```

```
In [40]: error_virginica = 100 * np.abs(describe_virginica -  
....: describe_all)
```

```
In [41]: error_virginica = error_virginica/describe_virginica
```



Viewing errors

```
In [42]: print(error_setosa)
```

	sepal_length	sepal_width
count	200.000000	200.000000
mean	16.726595	10.812913
std	134.919250	14.984768
min	0.000000	13.043478
25%	6.250000	12.500000
50%	16.000000	11.764706
75%	23.076923	10.204082
max	36.206897	0.000000

petal_length
200.000000
157.045144
916.502136
0.000000
14.285714
190.000000
223.809524
263.157895

petal_width
200.000000
387.533875
623.284534
0.000000
50.000000
550.000000
500.000000
316.666667



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Indexing time series

Using pandas to read datetime objects

- `read_csv()` function
 - Can read strings into datetime objects
 - Need to specify `'parse_dates=True'`
- ISO 8601 format
 - `yyyy-mm-dd hh:mm:ss`



Product sales CSV

	Date	Company	Product	Units
0	2015-02-02 08:30:00	Hooli	Software	3
1	2015-02-02 21:00:00	Mediacore	Hardware	9
2	2015-02-03 14:00:00	Initech	Software	13
3	2015-02-04 15:30:00	Streeplex	Software	13
4	2015-02-04 22:00:00	Acme Coporation	Hardware	14



Parse dates

```
In [3]: sales.head()
```

```
Out[3]:
```

Date	Company	Product	Units
2015-02-02 08:30:00	Hooli	Software	3
2015-02-02 21:00:00	Mediacore	Hardware	9
2015-02-03 14:00:00	Initech	Software	13
2015-02-04 15:30:00	Streeplex	Software	13
2015-02-04 22:00:00	Acme Coporation	Hardware	14



Parse dates

```
In [4]: sales.info()
DatetimeIndex: 19 entries, 2015-02-02 08:30:00 to 2015-02-26
09:00:00
Data columns (total 3 columns):
Company      19 non-null object
Product      19 non-null object
Units        19 non-null int64
dtypes: int64(1), object(2)
memory usage: 608.0+ bytes
```



Selecting single datetime

```
In [5]: sales.loc['2015-02-19 11:00:00', 'Company']  
Out[5]: 'Mediacore'
```



Selecting whole day

```
In [6]: sales.loc['2015-2-5']
```

```
Out[6]:
```

Date		Company	Product	Units
2015-02-05 02:00:00	Acme	Coporation	Software	19
2015-02-05 22:00:00		Hooli	Service	10

Partial datetime string selection

- Alternative formats:
 - `sales.loc['February 5, 2015']`
 - `sales.loc['2015-Feb-5']`
- Whole month: `sales.loc['2015-2']`
- Whole year: `sales.loc['2015']`



Selecting whole month

```
In [7]: sales.loc['2015-2']  
Out[7]:
```

Date	Company	Product	Units
2015-02-02 08:30:00	Hooli	Software	3
2015-02-02 21:00:00	Mediacore	Hardware	9
2015-02-03 14:00:00	Initech	Software	13
2015-02-04 15:30:00	Streeplex	Software	13
2015-02-04 22:00:00	Acme Coporation	Hardware	14
2015-02-05 02:00:00	Acme Coporation	Software	19
2015-02-05 22:00:00	Hooli	Service	10
2015-02-07 23:00:00	Acme Coporation	Hardware	1
2015-02-09 09:00:00	Streeplex	Service	19
2015-02-09 13:00:00	Mediacore	Software	7
2015-02-11 20:00:00	Initech	Software	7
2015-02-11 23:00:00	Hooli	Software	4
2015-02-16 12:00:00	Hooli	Software	10
2015-02-19 11:00:00	Mediacore	Hardware	16
...			



Slicing using dates/times

```
In [8]: sales.loc['2015-2-16':'2015-2-20']
```

```
Out[8]:
```

Date	Company	Product	Units
2015-02-16 12:00:00	Hooli	Software	10
2015-02-19 11:00:00	Mediacore	Hardware	16
2015-02-19 16:00:00	Mediacore	Service	10



Convert strings to datetime

```
In [9]: evening_2_11 = pd.to_datetime(['2015-2-11 20:00',  
....: '2015-2-11 21:00', '2015-2-11 22:00', '2015-2-11 23:00'])
```

```
In [10]: evening_2_11
```

```
Out[10]:
```

```
DatetimeIndex(['2015-02-11 20:00:00', '2015-02-11 21:00:00',  
               '2015-02-11 22:00:00', '2015-02-11 23:00:00'],  
              dtype='datetime64[ns]', freq=None)
```



Reindexing DataFrame

```
In [11]: sales.reindex(evening_2_11)
```

```
Out[11]:
```

		Company	Product	Units
2015-02-11	20:00:00	Initech	Software	7.0
2015-02-11	21:00:00	NaN	NaN	NaN
2015-02-11	22:00:00	NaN	NaN	NaN
2015-02-11	23:00:00	Hooli	Software	4.0



Filling missing values

```
In [12]: sales.reindex(evening_2_11, method='ffill')  
Out[12]:
```

		Company	Product	Units
2015-02-11	20:00:00	Initech	Software	7
2015-02-11	21:00:00	Initech	Software	7
2015-02-11	22:00:00	Initech	Software	7
2015-02-11	23:00:00	Hooli	Software	4

```
In [13]: sales.reindex(evening_2_11, method='bfill')  
Out[13]:
```

		Company	Product	Units
2015-02-11	20:00:00	Initech	Software	7
2015-02-11	21:00:00	Hooli	Software	4
2015-02-11	22:00:00	Hooli	Software	4
2015-02-11	23:00:00	Hooli	Software	4



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Resampling time series data



Sales data

```
In [1]: import pandas as pd
```

```
In [2]: sales = pd.read_csv('sales-feb-2015.csv',  
    ....:                   parse_dates=True, index_col= 'Date')
```

```
In [3]: sales.head()
```

```
Out[3]:
```

		Company	Product	Units
Date				
2015-02-02 08:30:00		Hooli	Software	3
2015-02-02 21:00:00		Mediacore	Hardware	9
2015-02-03 14:00:00		Initech	Software	13
2015-02-04 15:30:00		Streeplex	Software	13
2015-02-04 22:00:00	Acme	Coporation	Hardware	14

Resampling

- Statistical methods over different time intervals
 - `mean()`, `sum()`, `count()`, etc.
- Down-sampling
 - reduce datetime rows to slower frequency
- Up-sampling
 - increase datetime rows to faster frequency



Aggregating means

```
In [4]: daily_mean = sales.resample('D').mean()
```

```
In [5]: daily_mean
```

```
Out[5]:
```

	Units
Date	
2015-02-02	6.0
2015-02-03	13.0
2015-02-04	13.5
2015-02-05	14.5
2015-02-06	NaN
2015-02-07	1.0
2015-02-08	NaN
2015-02-09	13.0
2015-02-10	NaN
2015-02-11	5.5
2015-02-12	NaN
2015-02-13	NaN
2015-02-14	NaN



Verifying

```
In [6]: print(daily_mean.loc['2015-2-2'])  
Units      6.0  
Name: 2015-02-02 00:00:00, dtype: float64
```

```
In [7]: print(sales.loc['2015-2-2', 'Units'])  
Date  
2015-02-02 08:30:00      3  
2015-02-02 21:00:00      9  
Name: Units, dtype: int64
```

```
In [8]: sales.loc['2015-2-2', 'Units'].mean()  
Out[8]: 6.0
```



Method chaining

```
In [9]: sales.resample('D').sum()
```

```
Out[9]:
```

	Units
Date	
2015-02-02	6.0
2015-02-03	13.0
2015-02-04	13.5
2015-02-05	14.5
2015-02-06	NaN
2015-02-07	1.0
2015-02-08	NaN
2015-02-09	13.0
2015-02-10	NaN
2015-02-11	5.5
2015-02-12	NaN
2015-02-13	NaN

Method chaining

```
In [10]: sales.resample('D').sum().max()  
Out[10]:  
Units      29.0  
dtype: float64
```



Resampling strings

```
In [11]: sales.resample('W').count()
```

```
Out[11]:
```

	Company	Product	Units
Date			
2015-02-08	8	8	8
2015-02-15	4	4	4
2015-02-22	5	5	5
2015-03-01	2	2	2



Resampling frequencies

Input	Description
'min', 'T'	minute
'H'	hour
'D'	day
'B'	business day
'W'	week
'M'	month
'Q'	quarter
'A'	year



Multiplying frequencies

```
In [12]: sales.loc[:, 'Units'].resample('2W').sum()  
Out[12]:  
Date  
2015-02-08      82  
2015-02-22      79  
2015-03-08      14  
Freq: 2W-SUN, Name: Units, dtype: int64
```



Upsampling

```
In [13]: two_days = sales.loc['2015-2-4': '2015-2-5', 'Units']
```

```
In [13]: two_days
```

```
Out[13]:
```

```
Date
```

```
2015-02-04 15:30:00    13
```

```
2015-02-04 22:00:00    14
```

```
2015-02-05 02:00:00    19
```

```
2015-02-05 22:00:00    10
```

```
Name: Units, dtype: int64
```




Upsampling and filling

```
In [14]: two_days.resample('4H').ffill()
```

```
Out[14]:
```

```
Date
```

```
Date
```

```
2015-02-04 12:00:00      NaN
```

```
2015-02-04 16:00:00    13.0
```

```
2015-02-04 20:00:00    13.0
```

```
2015-02-05 00:00:00    14.0
```

```
2015-02-05 04:00:00    19.0
```

```
2015-02-05 08:00:00    19.0
```

```
2015-02-05 12:00:00    19.0
```

```
2015-02-05 16:00:00    19.0
```

```
2015-02-05 20:00:00    19.0
```

```
Freq: 4H, Name: Units, dtype: float64
```



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Manipulating time series data



Sales data

```
In [1]: import pandas as pd
```

```
In [2]: sales = pd.read_csv('sales-feb-2015.csv',  
    ....:                  parse_dates=['Date'])
```

```
In [3]: sales.head()
```

```
Out[3]:
```

	Date	Company	Product	Units
0	2015-02-02 08:30:00	Hooli	Software	3
1	2015-02-02 21:00:00	Mediacore	Hardware	9
2	2015-02-03 14:00:00	Initech	Software	13
3	2015-02-04 15:30:00	Streeplex	Software	13
4	2015-02-04 22:00:00	Acme Coporation	Hardware	14



String methods

```
In [4]: sales['Company'].str.upper()
```

```
Out[4]:
```

```
0          HOOLI
1      MEDIACORE
2      INITECH
3      STREEPLEX
4  ACME COPORATION
5  ACME COPORATION
6          HOOLI
7  ACME COPORATION
8      STREEPLEX
9      MEDIACORE
10         INITECH
11         HOOLI
12         HOOLI
13      MEDIACORE
14      MEDIACORE
15      MEDIACORE
```

```
...
```



Substring matching

```
In [5]: sales['Product'].str.contains('ware')
```

```
Out[5]:
```

```
0      True
1      True
2      True
3      True
4      True
5      True
6     False
7      True
8     False
9      True
10     True
11     True
12     True
13     True
14     False
...
```



Boolean arithmetic

```
In [6]: True + False  
Out[6]: 1
```

```
In [7]: True + True  
Out[7]: 2
```

```
In [8]: False + False  
Out[8]: 0
```



Boolean reduction

```
In [9]: sales['Product'].str.contains('ware').sum()  
Out[9]: 14
```




Datetime methods

```
In [9]: sales['Date'].dt.hour
```

```
Out[9]:
```

```
0      8
1     21
2     14
3     15
4     22
5      2
6     22
7     23
8      9
9     13
10    20
11    23
12    12
13    11
14    16
...
```



Set timezone

```
In [10]: central = sales['Date'].dt.tz_localize('US/Central')
```

```
In [11]: central
```

```
Out[11]:
```

```
0      2015-02-02 08:30:00-06:00
1      2015-02-02 21:00:00-06:00
2      2015-02-03 14:00:00-06:00
3      2015-02-04 15:30:00-06:00
4      2015-02-04 22:00:00-06:00
5      2015-02-05 02:00:00-06:00
6      2015-02-05 22:00:00-06:00
7      2015-02-07 23:00:00-06:00
8      2015-02-09 09:00:00-06:00
9      2015-02-09 13:00:00-06:00
10     2015-02-11 20:00:00-06:00
11     2015-02-11 23:00:00-06:00
12     2015-02-16 12:00:00-06:00
```

```
...
```

```
Name: Date, dtype: datetime64[ns, US/Central]
```



Convert timezone

```
In [12]: central.dt.tz_convert('US/Eastern')
```

```
Out[12]:
```

```
0      2015-02-02 09:30:00-05:00
1      2015-02-02 22:00:00-05:00
2      2015-02-03 15:00:00-05:00
3      2015-02-04 16:30:00-05:00
4      2015-02-04 23:00:00-05:00
5      2015-02-05 03:00:00-05:00
6      2015-02-05 23:00:00-05:00
7      2015-02-08 00:00:00-05:00
8      2015-02-09 10:00:00-05:00
9      2015-02-09 14:00:00-05:00
10     2015-02-11 21:00:00-05:00
11     2015-02-12 00:00:00-05:00
12     2015-02-16 13:00:00-05:00
13     2015-02-19 12:00:00-05:00
14     2015-02-19 17:00:00-05:00
```

```
...
```

```
Name: Date, dtype: datetime64[ns, US/Eastern]
```



Method chaining

```
In [13]: sales['Date'].dt.tz_localize('US/Central').  
        ...: dt.tz_convert('US/Eastern')
```

```
Out[13]:
```

```
0      2015-02-02 09:30:00-05:00  
1      2015-02-02 22:00:00-05:00  
2      2015-02-03 15:00:00-05:00  
3      2015-02-04 16:30:00-05:00  
4      2015-02-04 23:00:00-05:00  
5      2015-02-05 03:00:00-05:00  
6      2015-02-05 23:00:00-05:00  
7      2015-02-08 00:00:00-05:00  
8      2015-02-09 10:00:00-05:00  
9      2015-02-09 14:00:00-05:00  
10     2015-02-11 21:00:00-05:00  
11     2015-02-12 00:00:00-05:00  
12     2015-02-16 13:00:00-05:00  
13     2015-02-19 12:00:00-05:00  
14     2015-02-19 17:00:00-05:00
```

```
...
```

```
Name: Date, dtype: datetime64[ns, US/Eastern]
```



World Population

```
In [14]: population = pd.read_csv('world_population.csv',  
    ...: parse_dates=True, index_col= 'Date')
```

```
In [15]: population
```

```
Out[15]:
```

Date	Population
1960-12-31	2.087485e+10
1970-12-31	2.536513e+10
1980-12-31	3.057186e+10
1990-12-31	3.644928e+10
2000-12-31	4.228550e+10
2010-12-31	4.802217e+10



Upsample population

```
In [16]: population.resample('A').first()  
Out[16]:
```

	Population
Date	
1960-12-31	2.087485e+10
1961-12-31	NaN
1962-12-31	NaN
1963-12-31	NaN
1964-12-31	NaN
1965-12-31	NaN
1966-12-31	NaN
1967-12-31	NaN
1968-12-31	NaN
1969-12-31	NaN
1970-12-31	2.536513e+10
1971-12-31	NaN
1972-12-31	NaN



Interpolate missing data

```
In [17]: population.resample('A').first().interpolate('linear')  
Out[17]:
```

Date	Population
1960-12-31	2.087485e+10
1961-12-31	2.132388e+10
1962-12-31	2.177290e+10
1963-12-31	2.222193e+10
1964-12-31	2.267096e+10
1965-12-31	2.311999e+10
1966-12-31	2.356902e+10
1967-12-31	2.401805e+10
1968-12-31	2.446707e+10
1969-12-31	2.491610e+10
1970-12-31	2.536513e+10
1971-12-31	2.588580e+10
1972-12-31	2.640648e+10



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Time series visualization



Topics

- Line types
- Plot types
- Subplots



S&P 500 Data

```
In [1]: import pandas as pd
```

```
In [2]: import matplotlib.pyplot as plt
```

```
In [3]: sp500 = pd.read_csv('sp500.csv', parse_dates=True,  
    ....:                    index_col= 'Date')
```

```
In [4]: sp500.head()
```

```
Out[4]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2010-01-04	1116.560059	1133.869995	1116.560059	1132.989990	3991400000	1132.989990
2010-01-05	1132.660034	1136.630005	1129.660034	1136.520020	2491020000	1136.520020
2010-01-06	1135.709961	1139.189941	1133.949951	1137.140015	4972660000	1137.140015
2010-01-07	1136.270020	1142.459961	1131.319946	1141.689941	5270680000	1141.689941
2010-01-08	1140.520020	1145.390015	1136.219971	1144.979980	4389590000	1144.979980



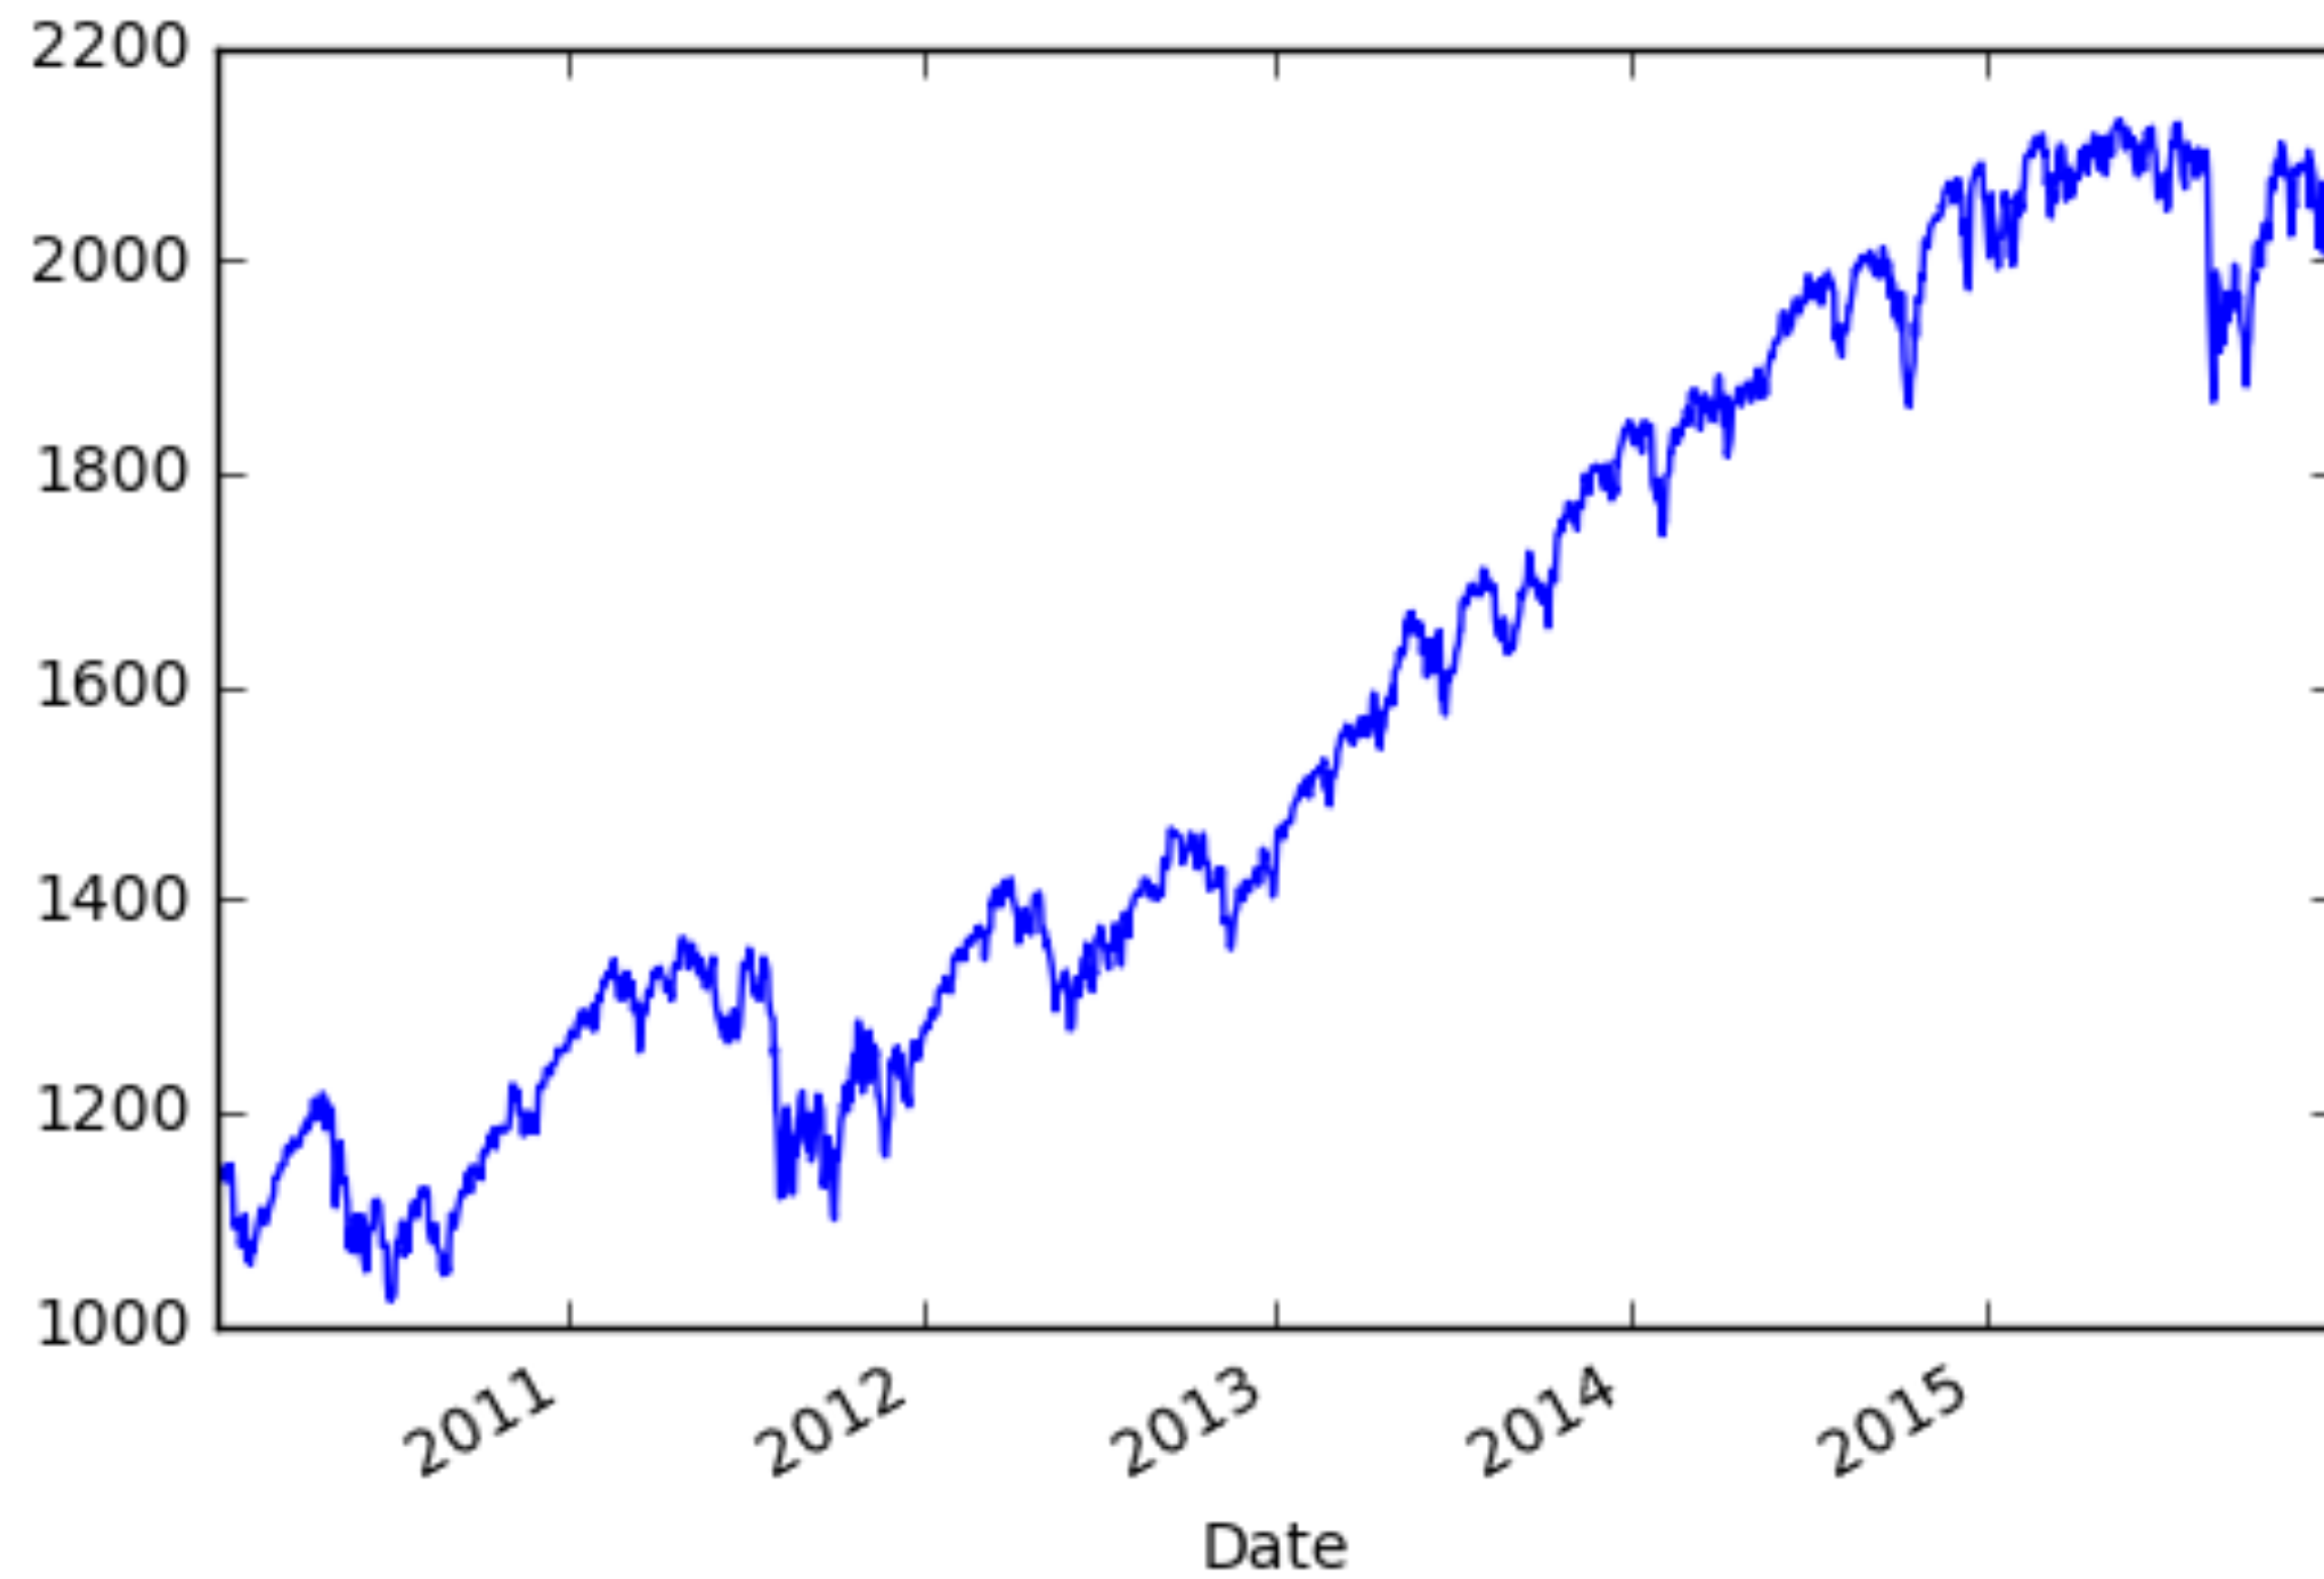
Pandas plot

```
In [5]: sp500['Close'].plot()
```

```
In [6]: plt.show()
```



Default plot





Labels and title

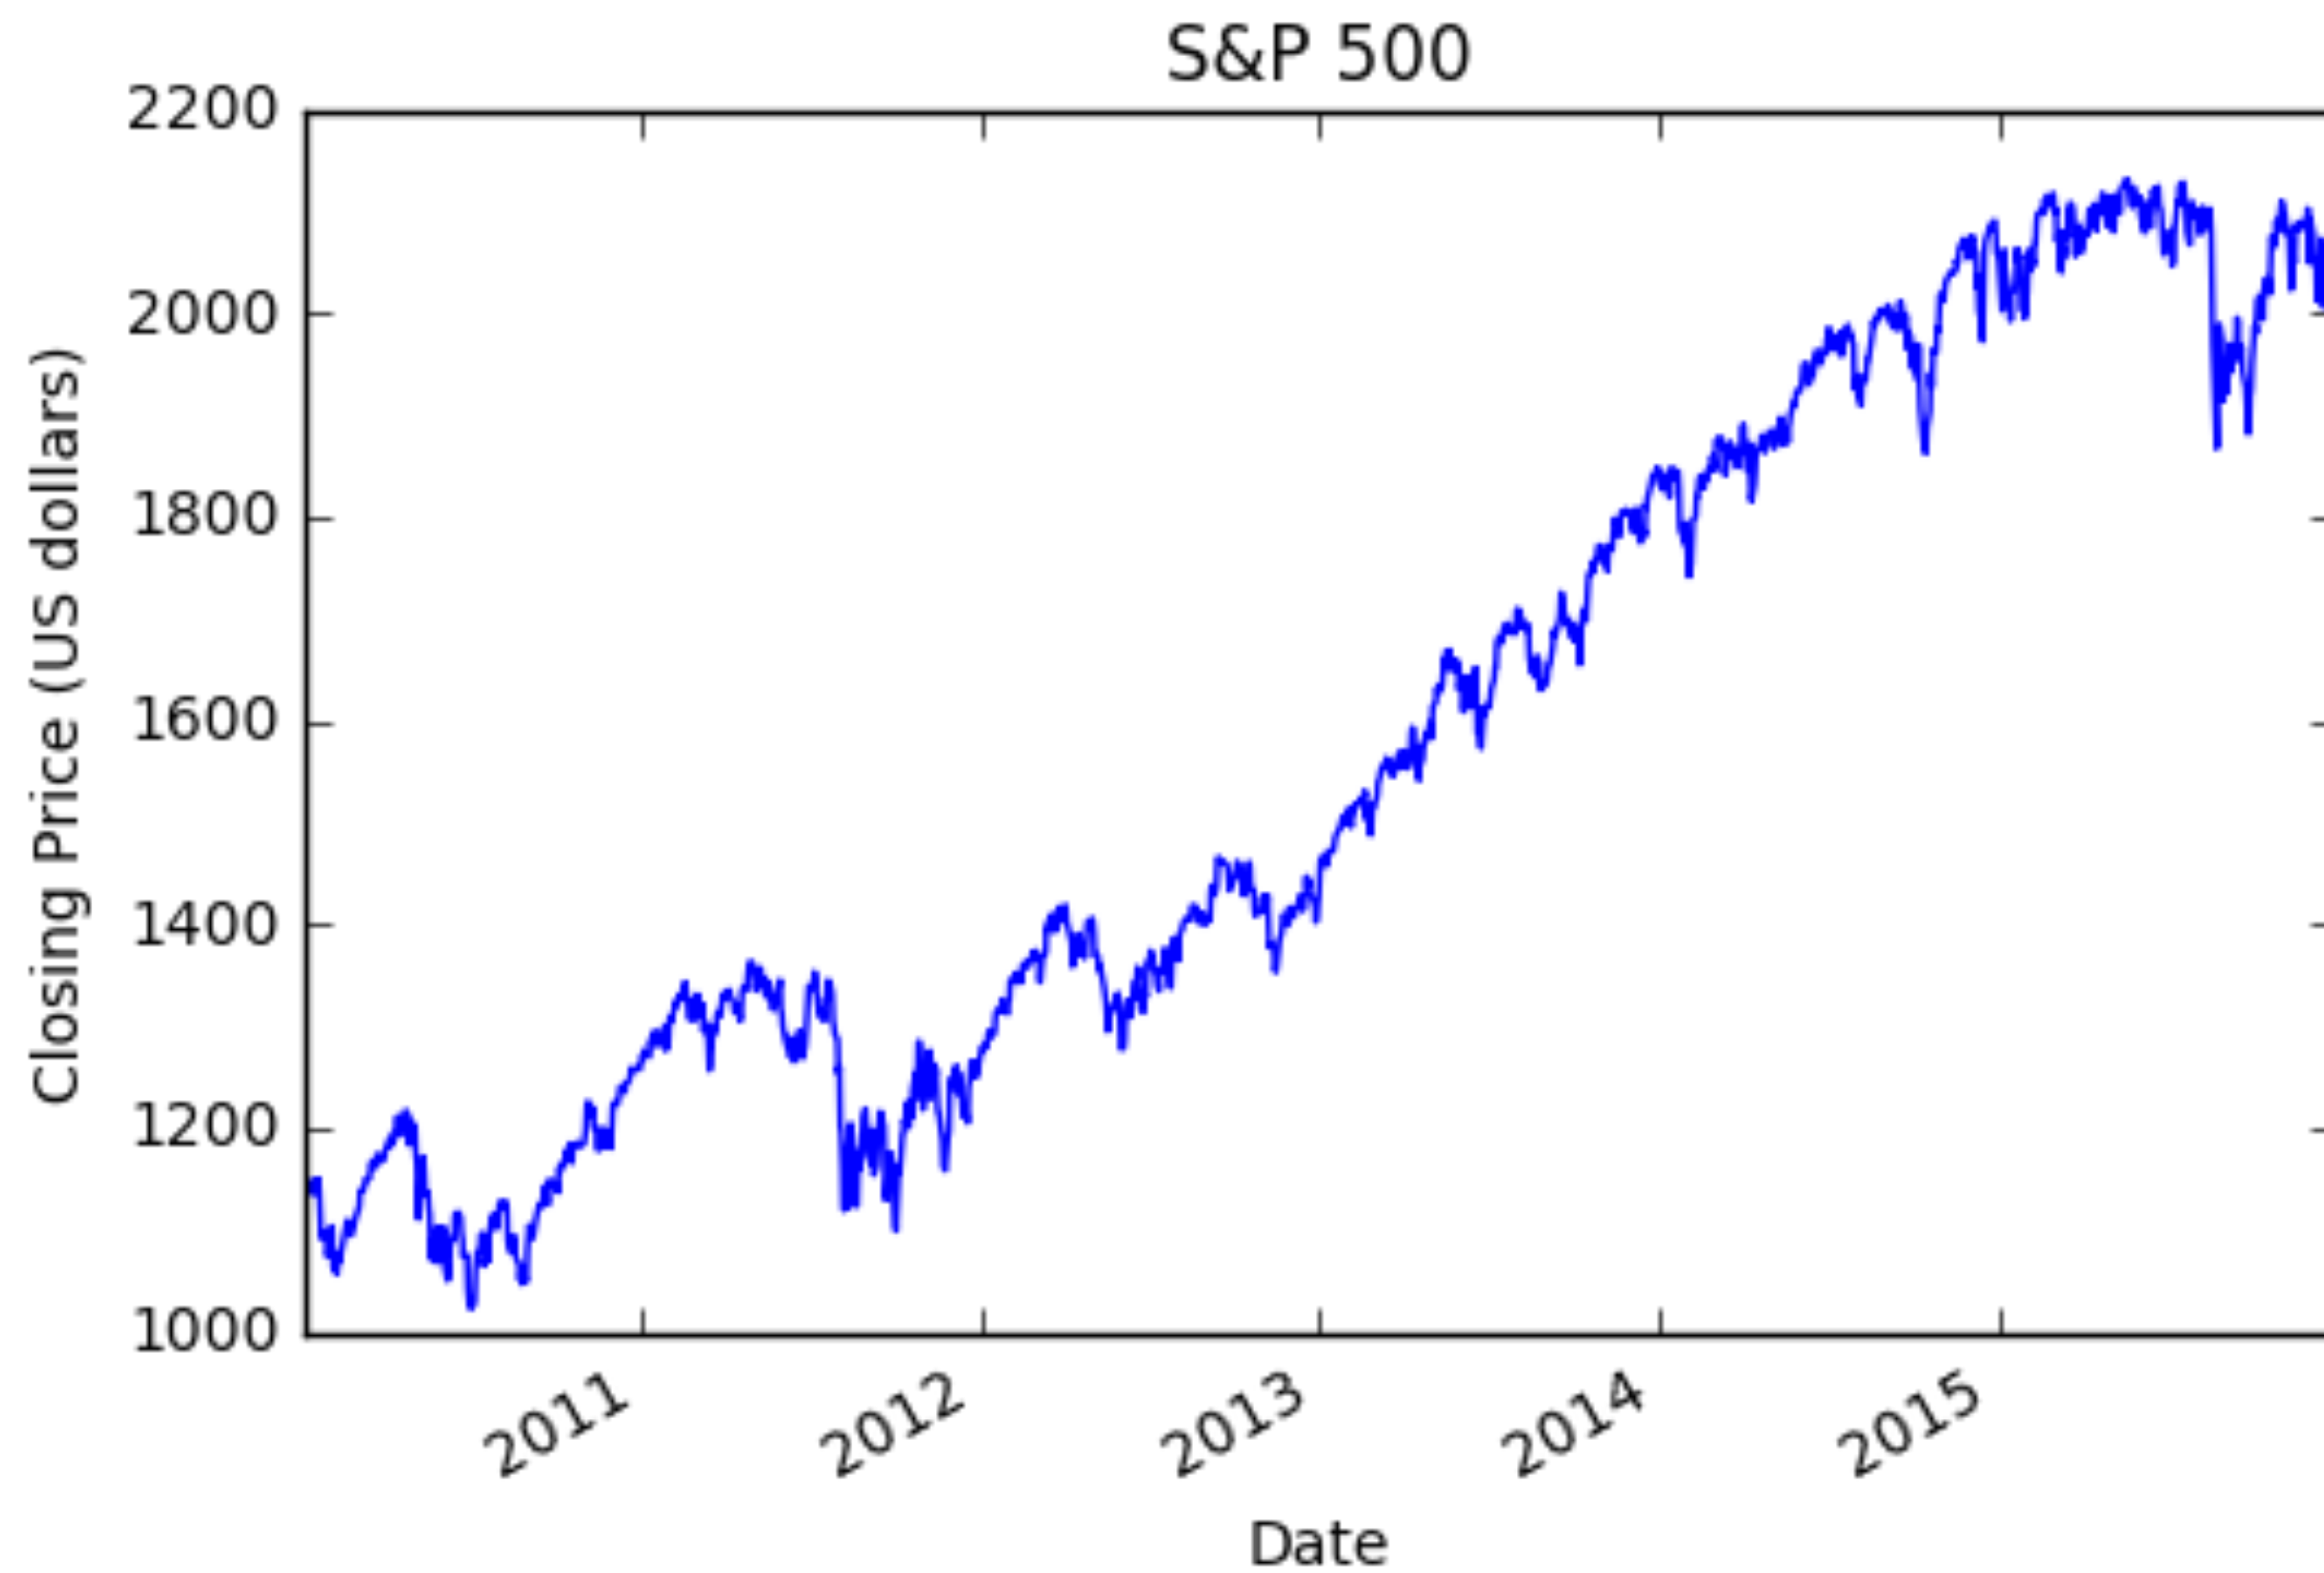
```
In [7]: sp500['Close'].plot(title='S&P 500')
```

```
In [8]: plt.ylabel('Closing Price (US Dollars)')
```

```
In [9]: plt.show()
```



Labels and title





One week

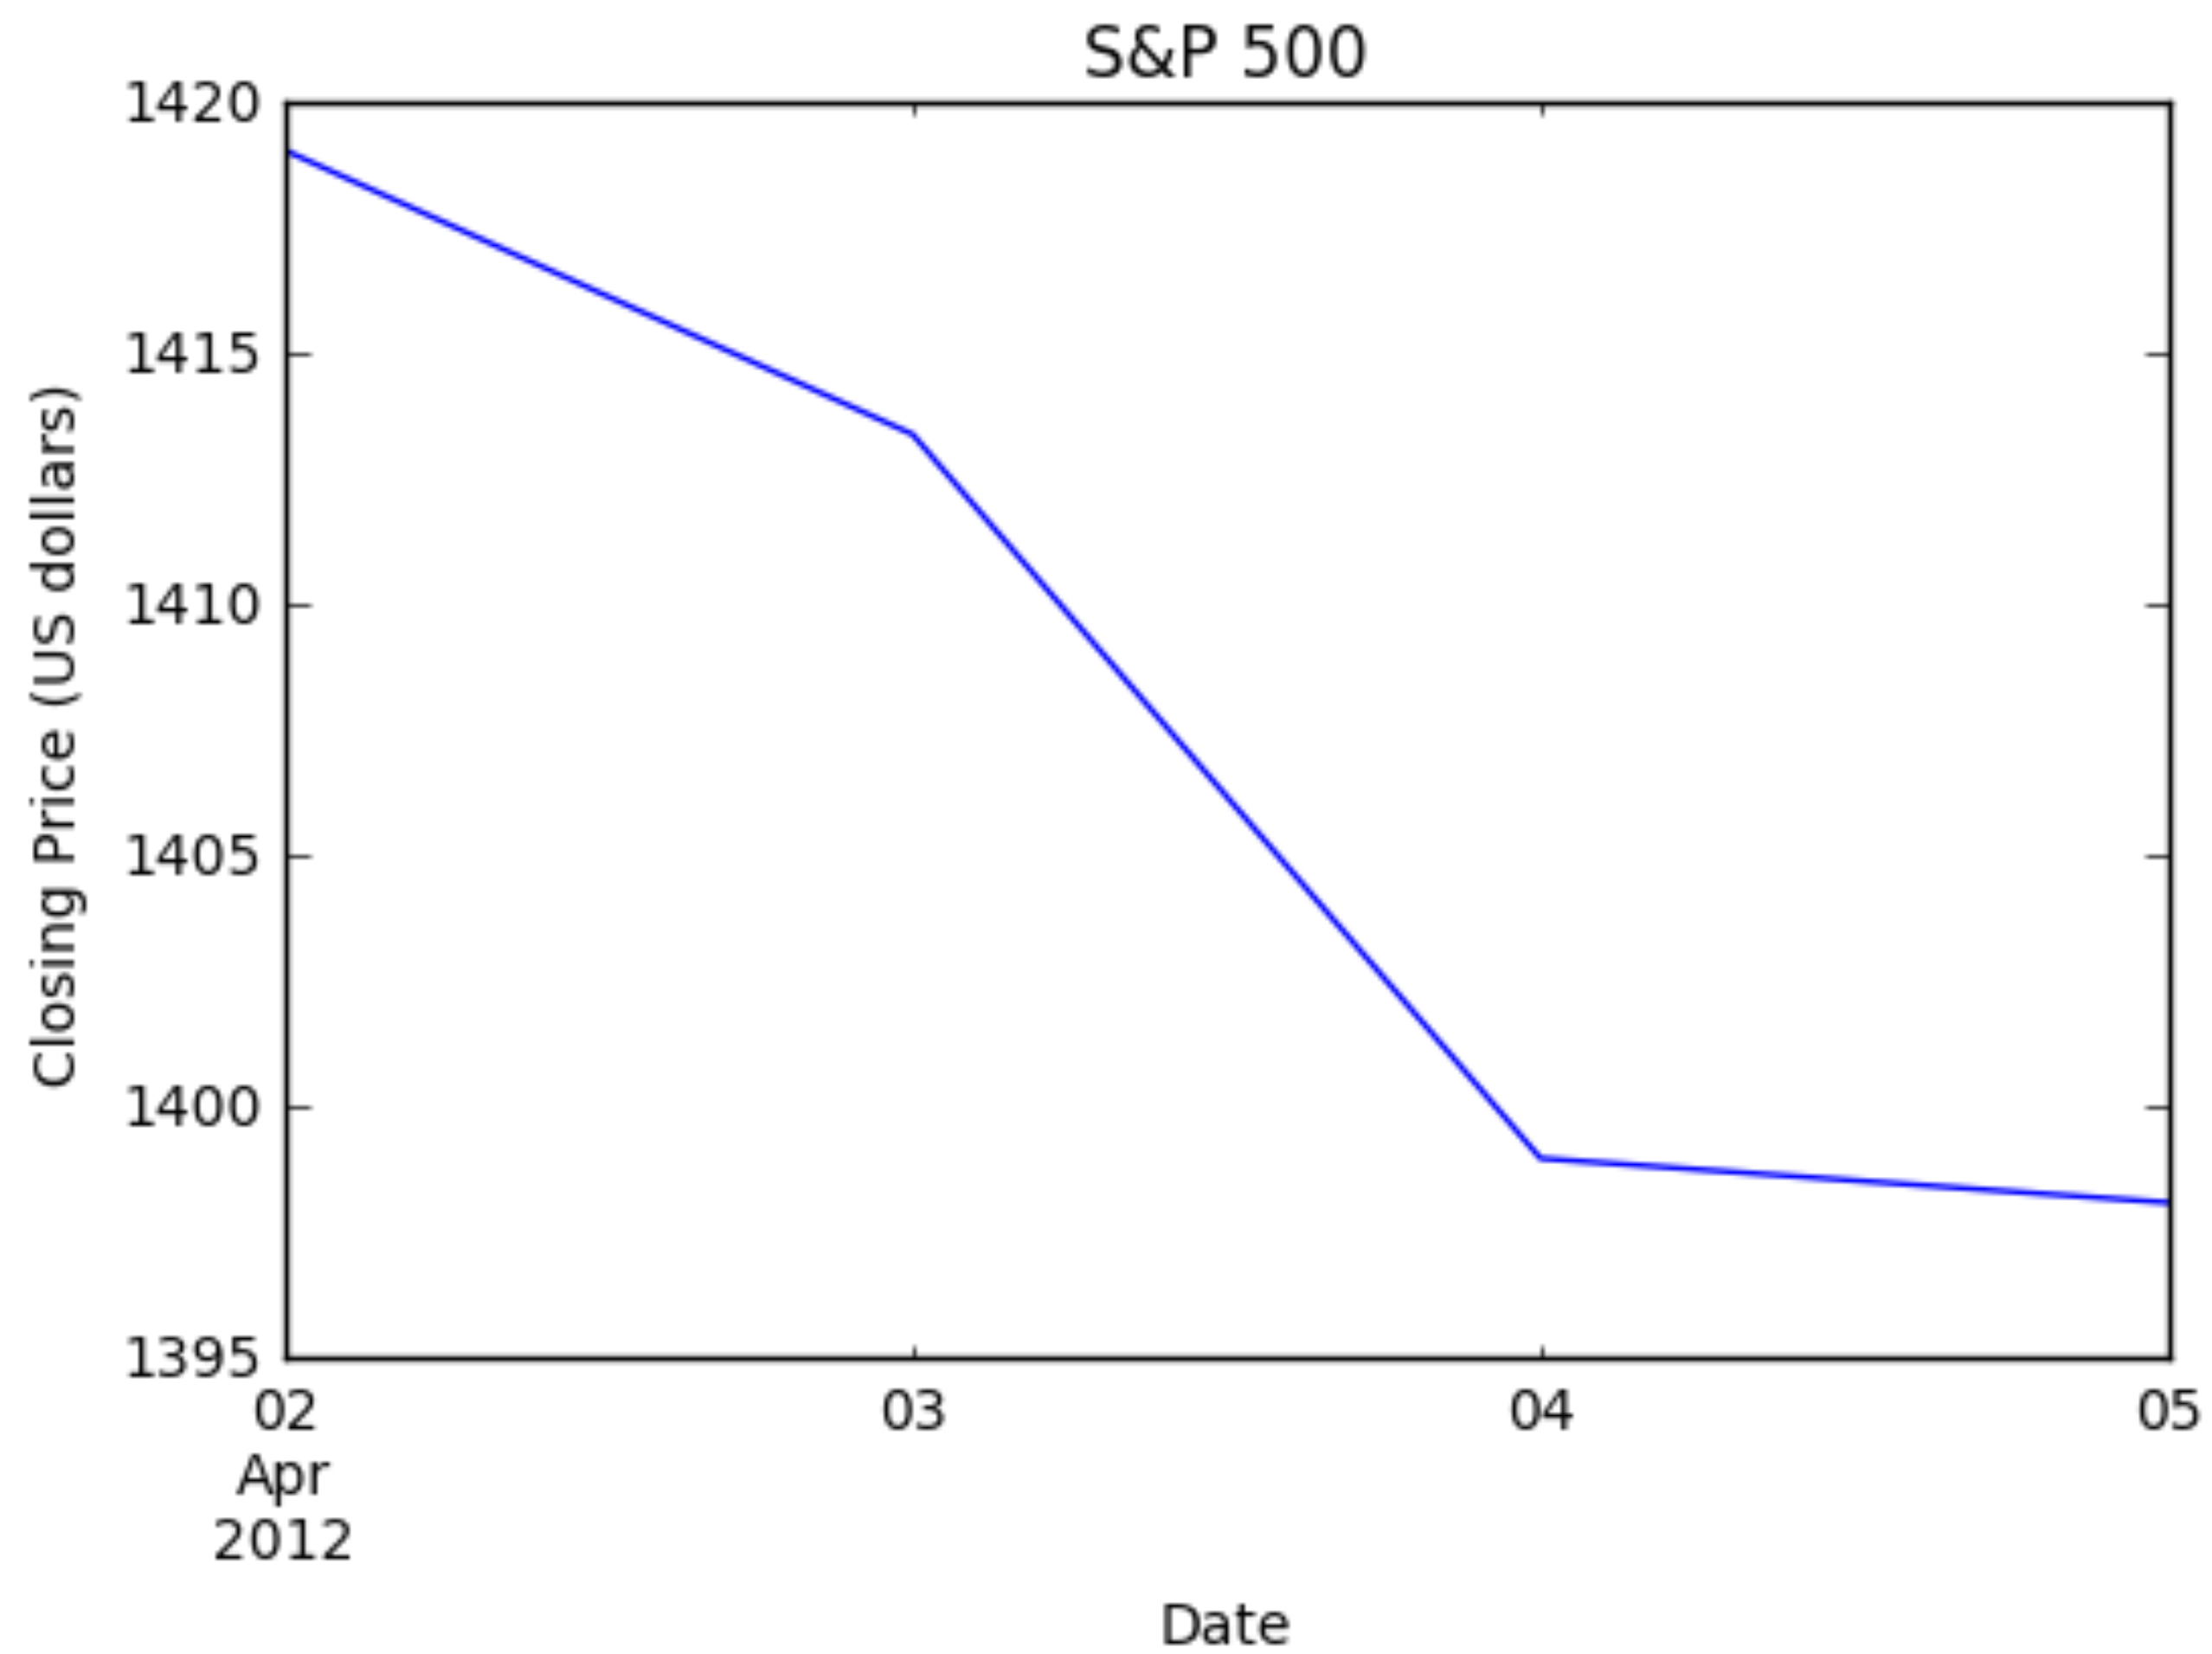
```
In [10]: sp500.loc['2012-4-1':'2012-4-7', 'Close'].plot(title='S&P  
...: 500')
```

```
In [11]: plt.ylabel('Closing Price (US Dollars)')
```

```
In [12]: plt.show()
```




One week





Plot styles

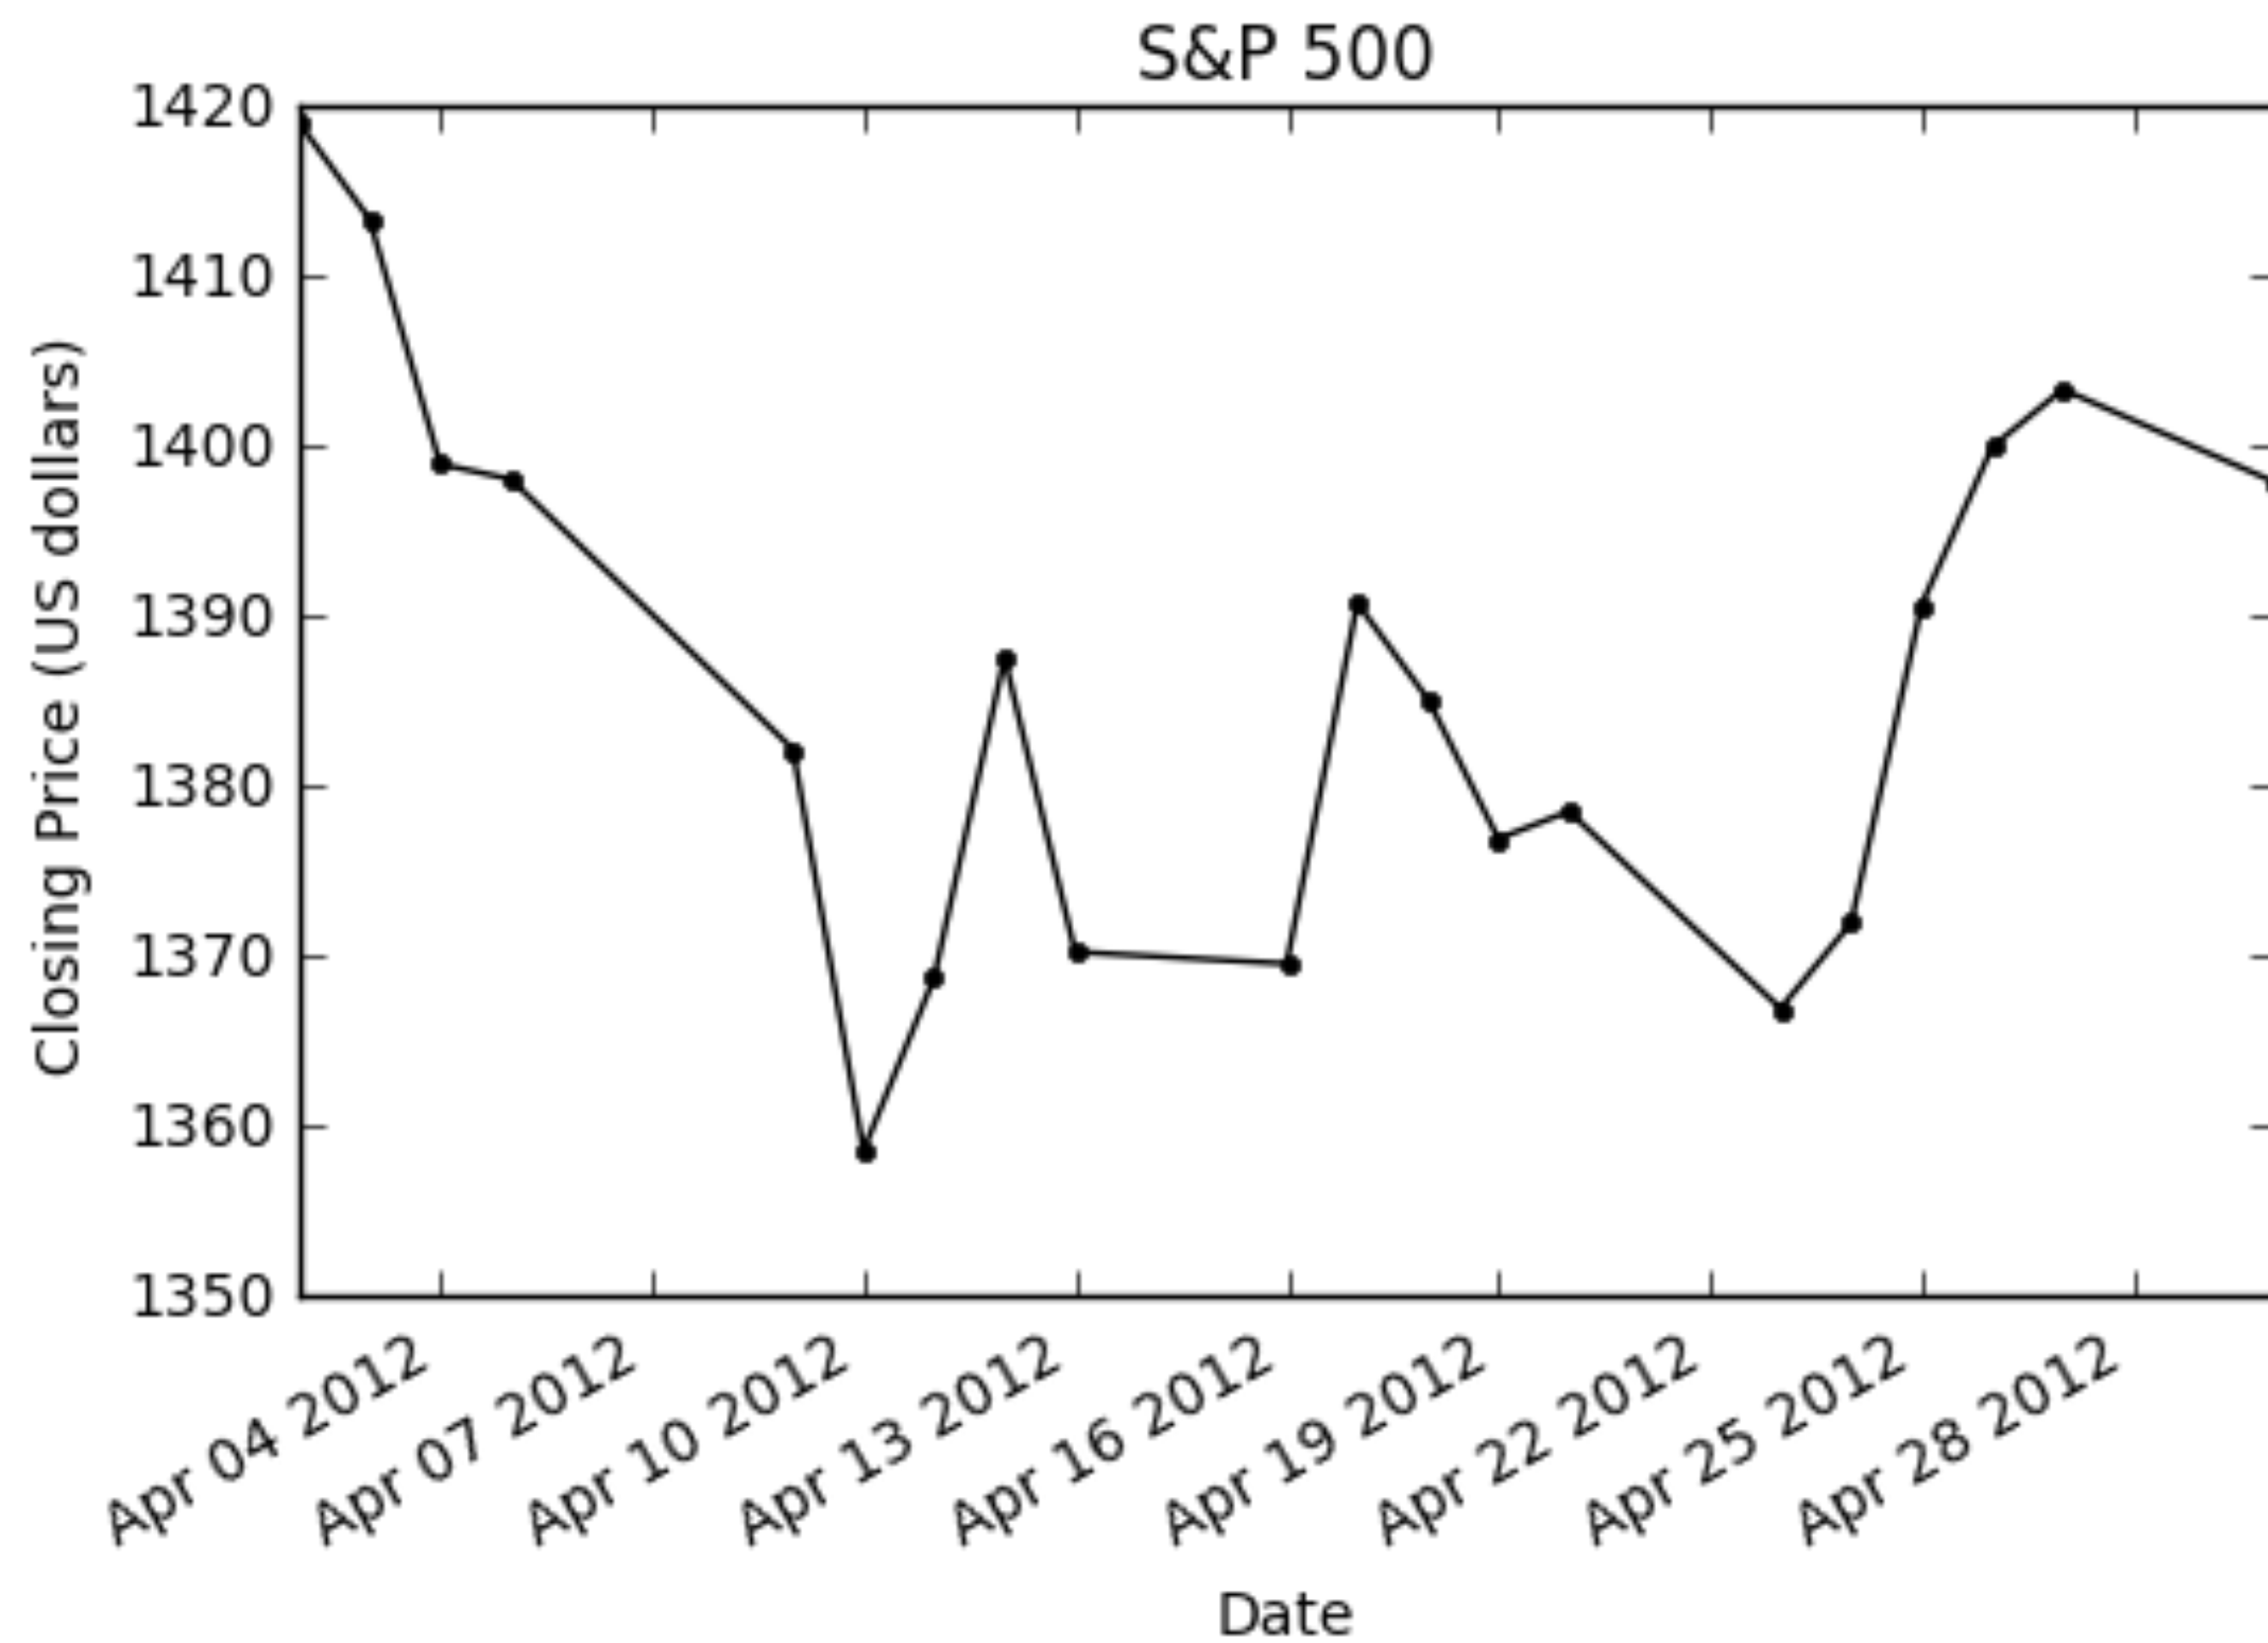
```
In [13]: sp500.loc['2012-4', 'Close'].plot(style='k.-',  
      ....:                                     title='S&P500')
```

```
In [14]: plt.ylabel('Closing Price (US Dollars)')
```

```
In [15]: plt.show()
```



One week



More plot styles

- Style format string
 - color (k: black)
 - marker (. : dot)
 - line type (-: solid)



More plot styles

Color	Marker	Line
b: blue	o: circle	: dotted
g: green	*: star	–: dashed
r: red	s: square	
c: cyan	+: plus	



Area plot

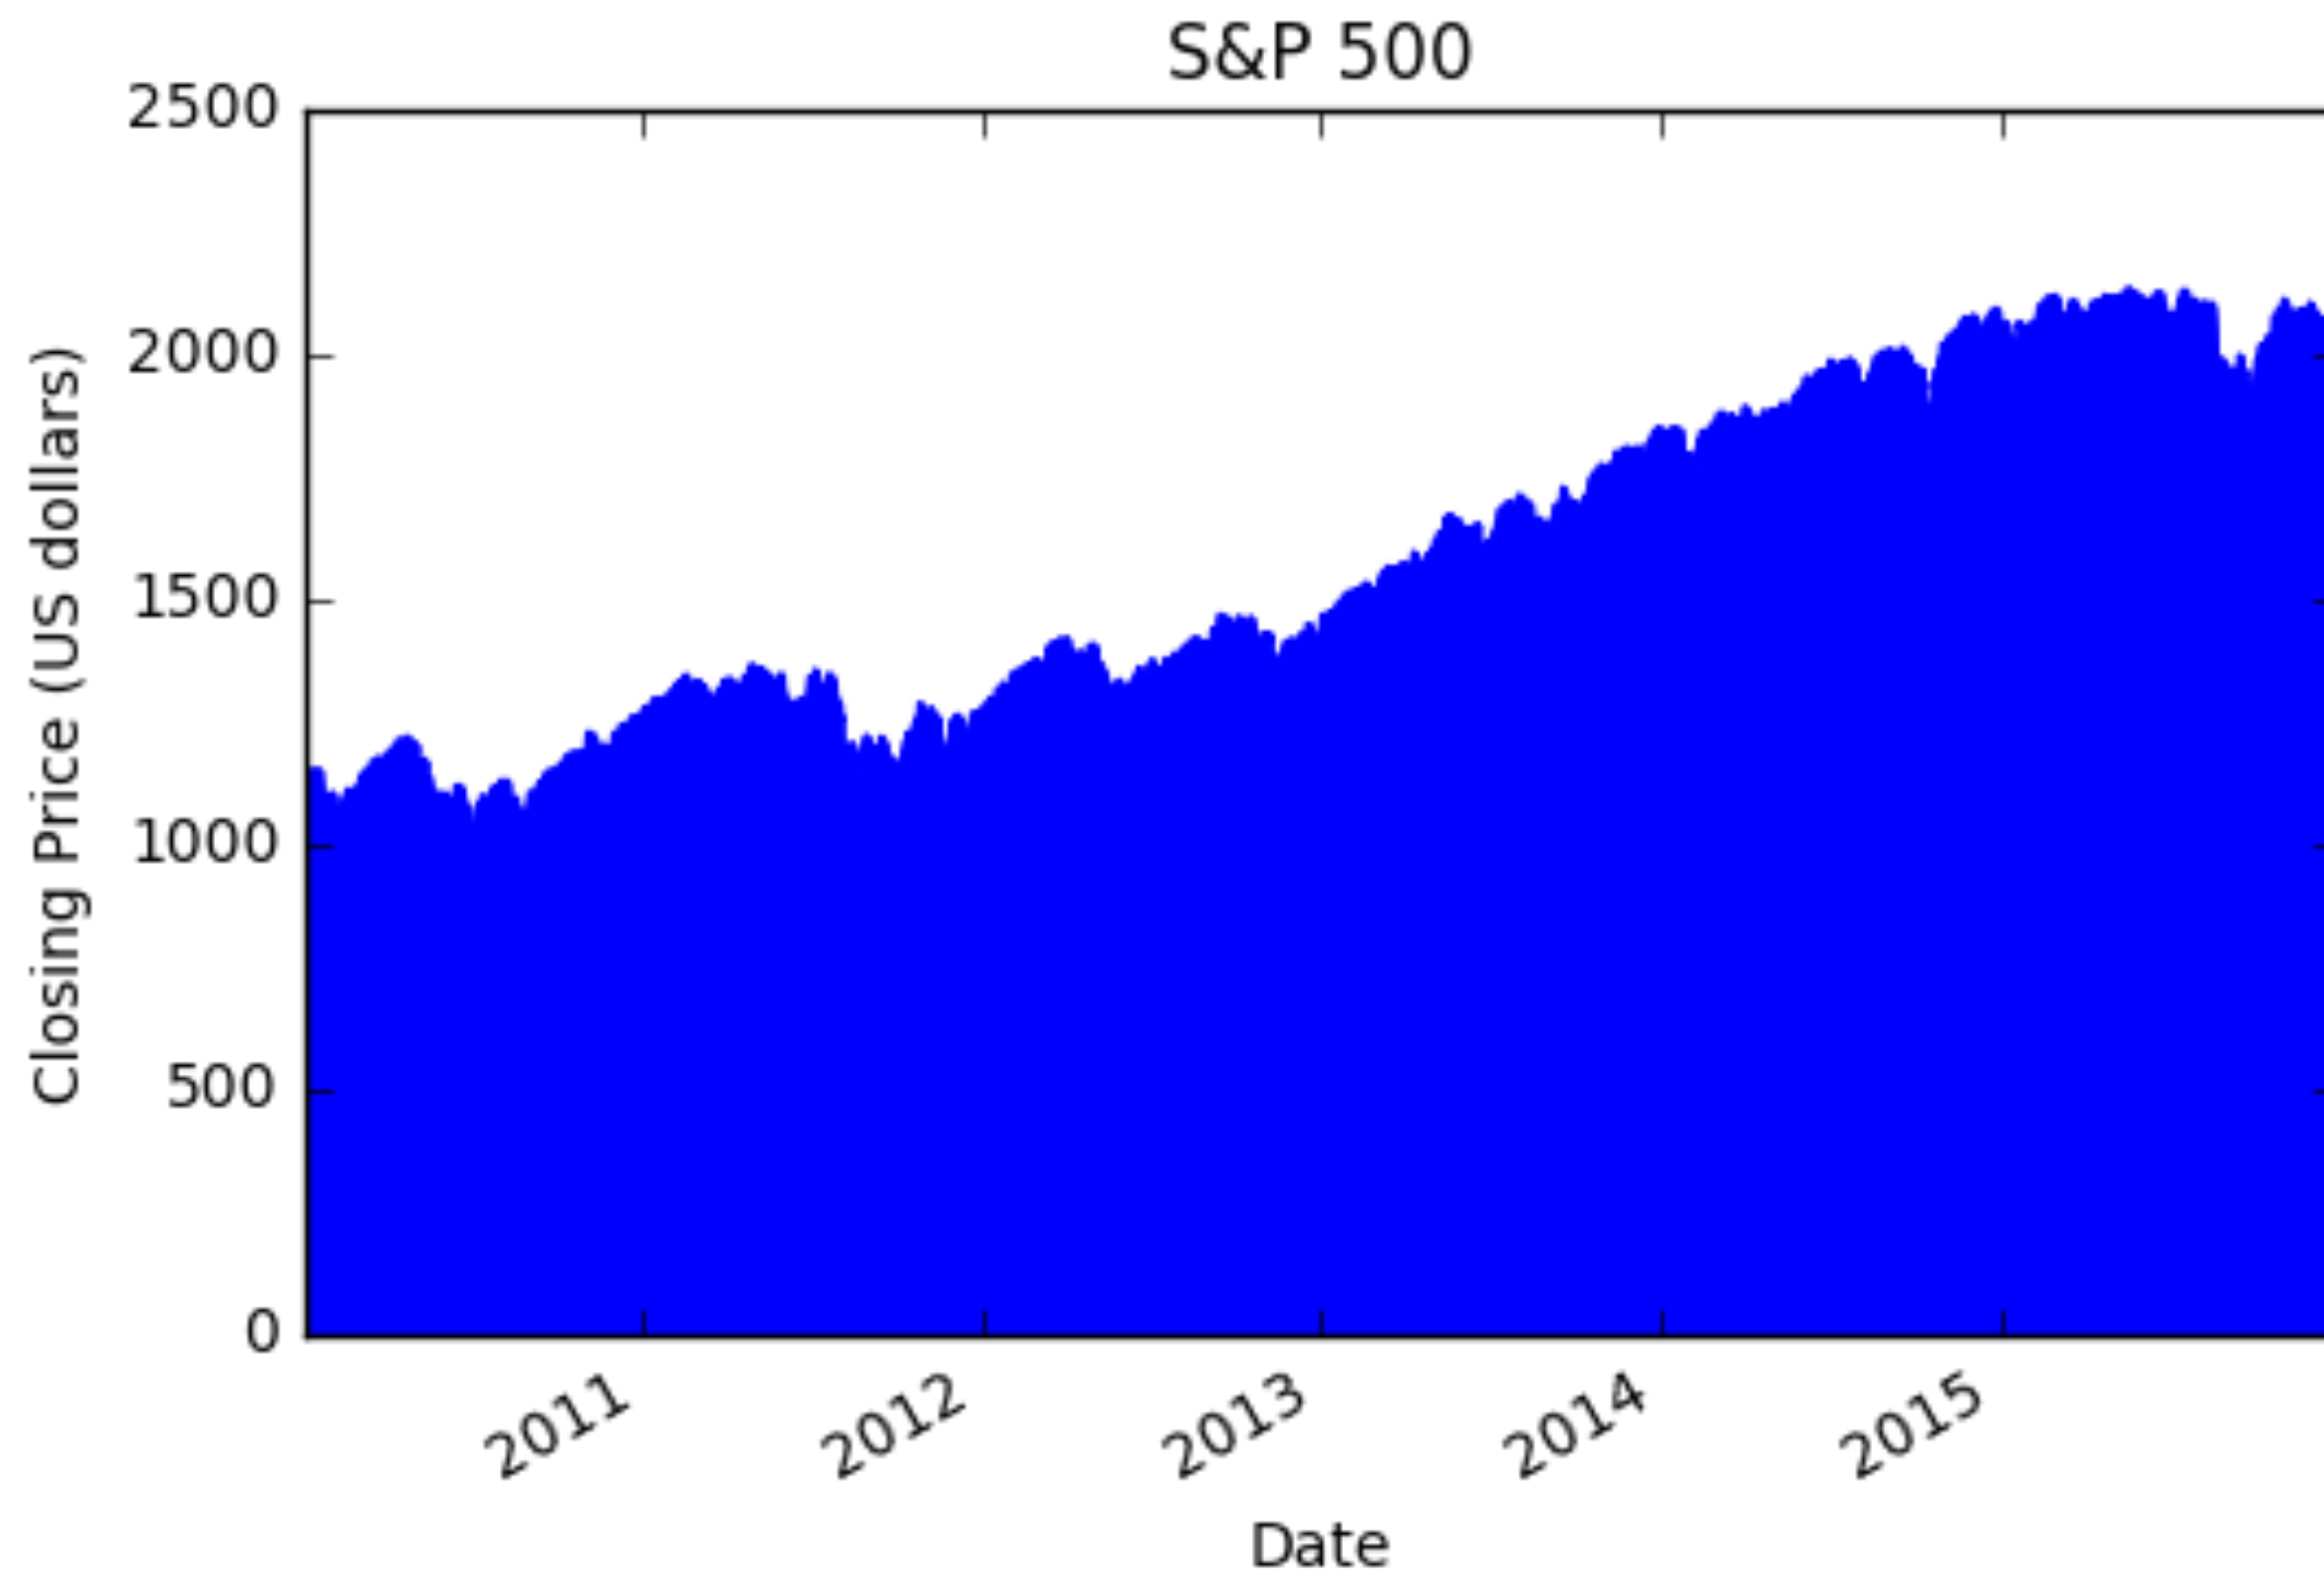
```
In [16]: sp500['Close'].plot(kind='area', title='S&P 500')
```

```
In [17]: plt.ylabel('Closing Price (US Dollars)')
```

```
In [18]: plt.show()
```



Area plot





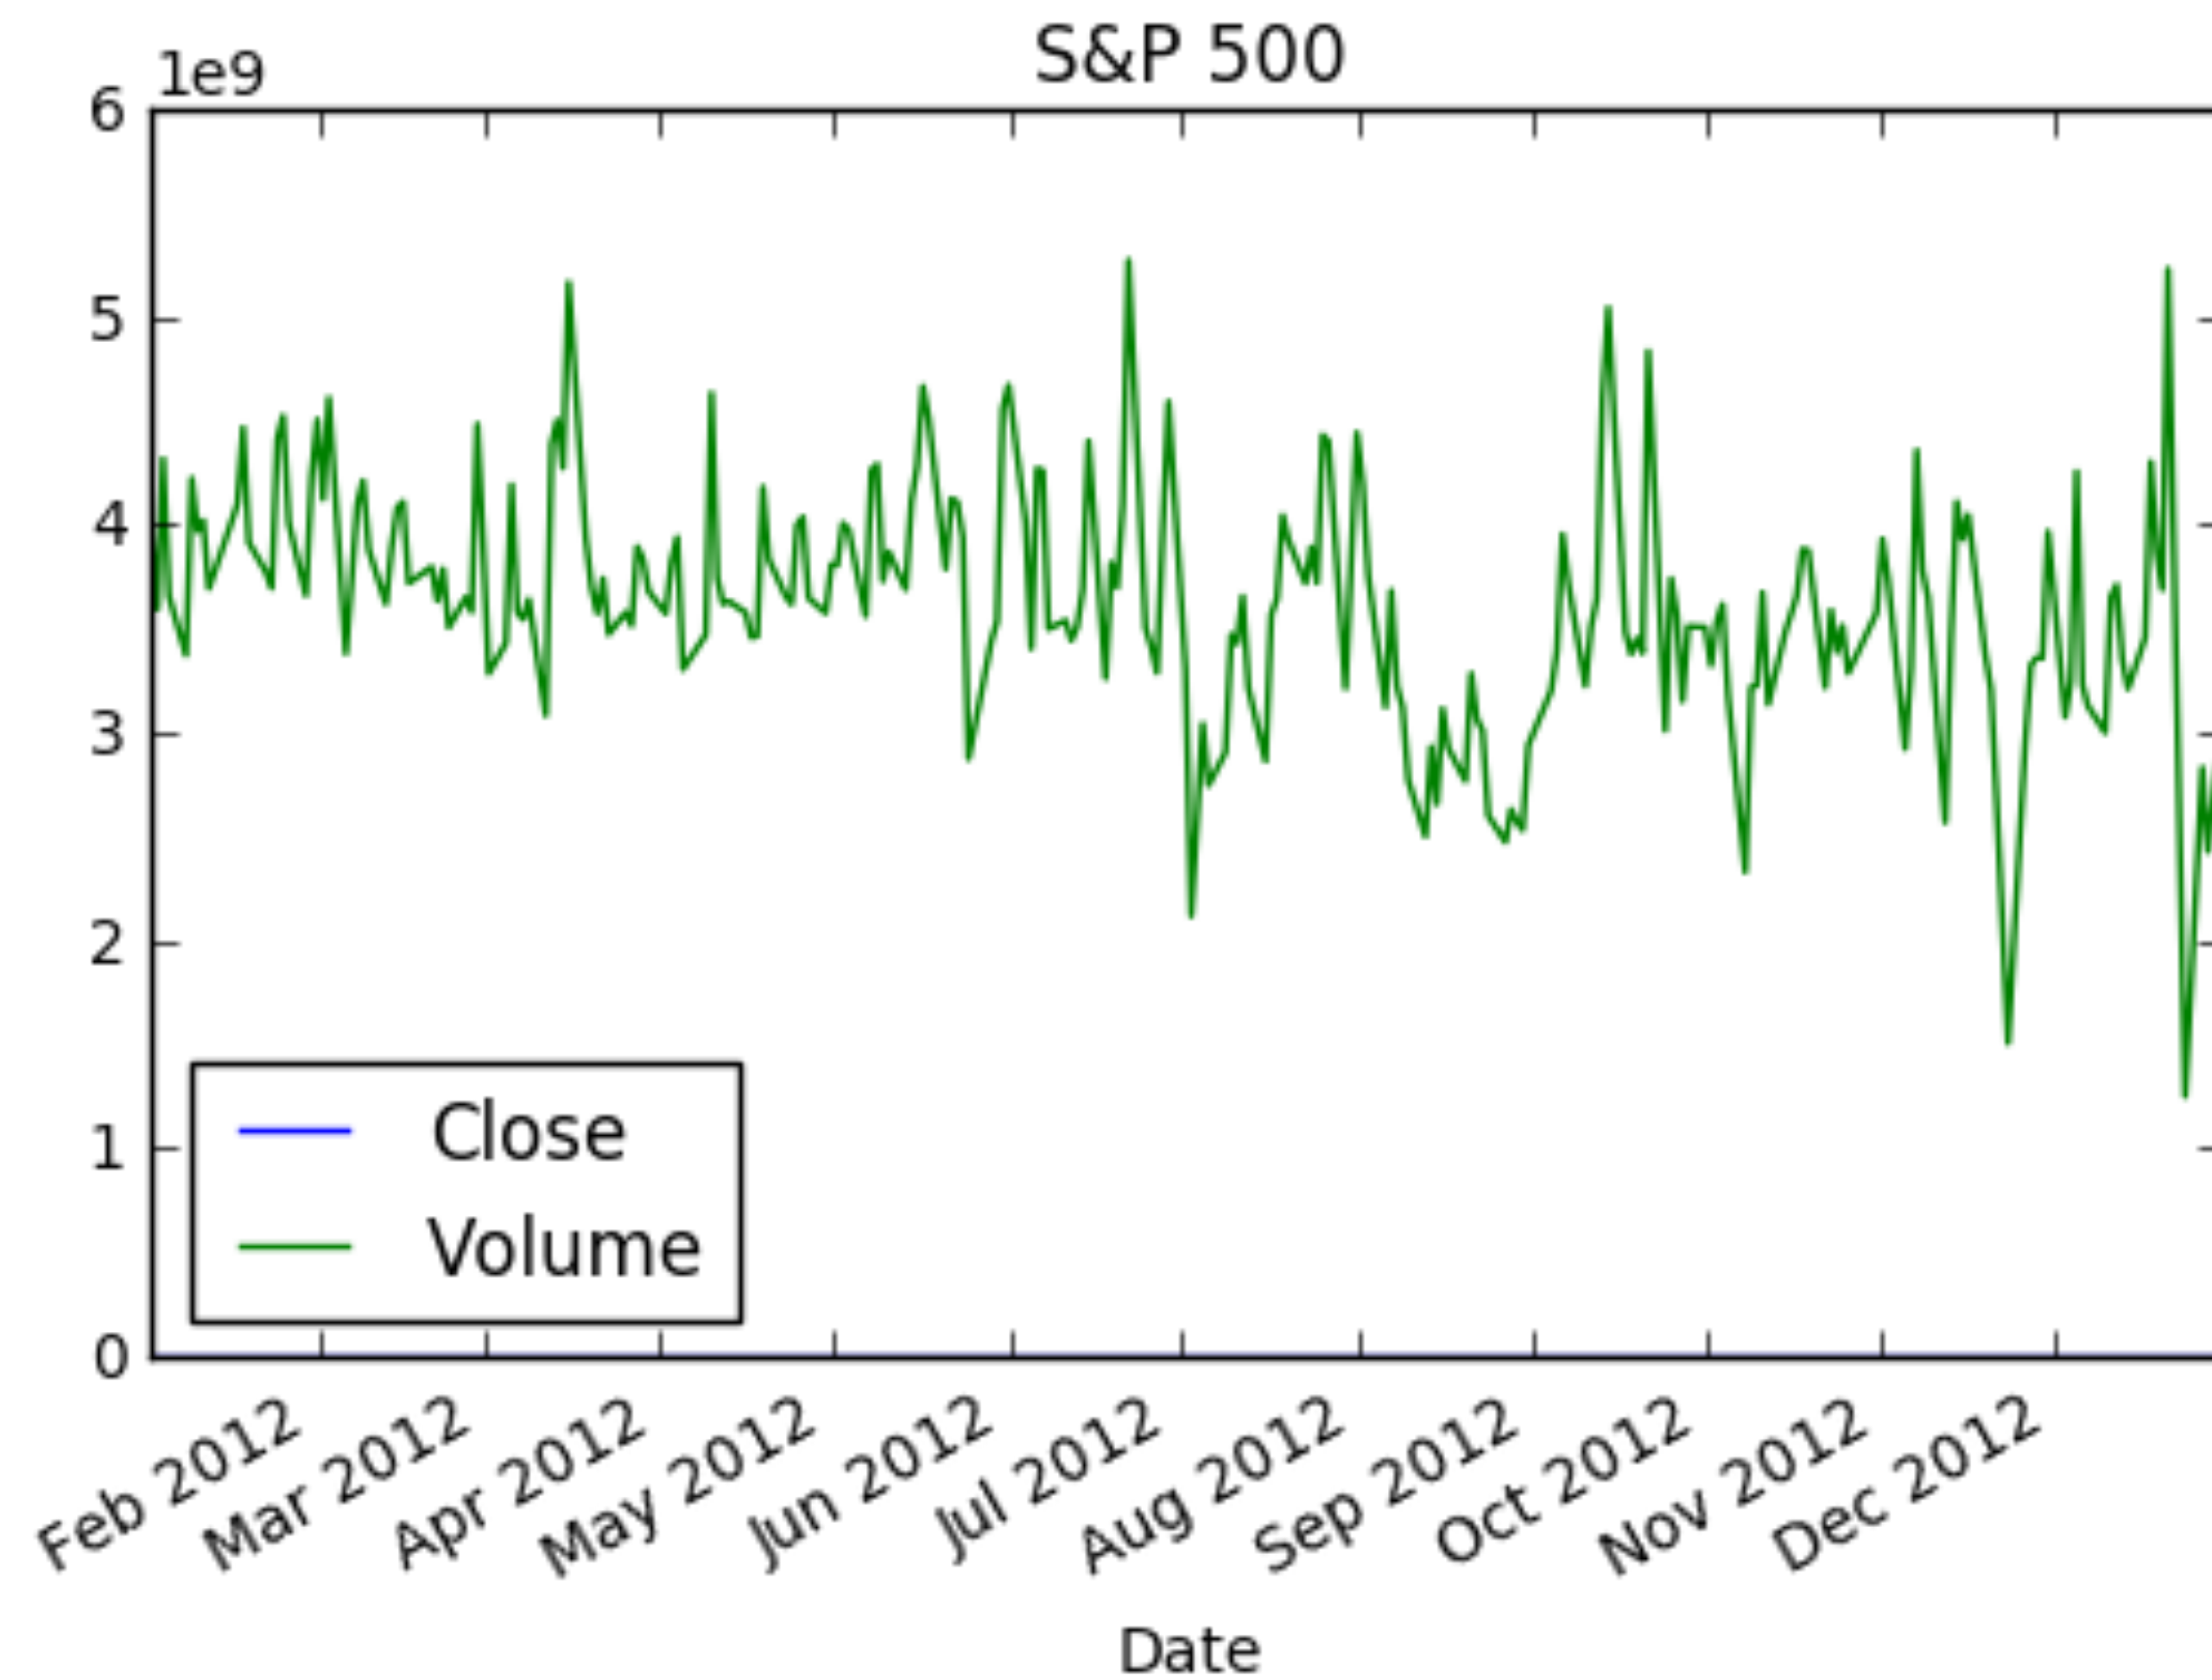
Multiple columns

```
In [19]: sp500.loc['2012', ['Close', 'Volume']].plot(title='S&P  
...: 500')
```

```
In [20]: plt.show()
```




Multiple columns






Subplots

```
In [21]: sp500.loc['2012', ['Close', 'Volume']].plot(subplots=True)  
  
In [22]: plt.show()
```



Subplots





PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Reading and cleaning the data



Case study

- Comparing observed weather data from two sources

	Temperature	DewPoint	Pressure	Date
0	46.2	37.5	1.0	20100101 00:00
1	44.6	37.1	1.0	20100101 01:00
2	44.1	36.9	1.0	20100101 02:00
3	43.8	36.9	1.0	20100101 03:00
4	43.5	36.8	1.0	20100101 04:00

	Date	Wban	...	station_pressure	sea_level_pressure
0	2011-01-01 00:53:00	13904	...	29.42	29.95
1	2011-01-01 01:53:00	13904	...	29.49	30.01
2	2011-01-01 02:53:00	13904	...	29.49	30.01
3	2011-01-01 03:53:00	13904	...	29.51	30.03
4	2011-01-01 04:53:00	13904	...	29.51	30.04



Climate normals of Austin, TX from 1981-2010

	Temperature	DewPoint	Pressure	Date
0	46.2	37.5	1.0	20100101 00:00
1	44.6	37.1	1.0	20100101 01:00
2	44.1	36.9	1.0	20100101 02:00
3	43.8	36.9	1.0	20100101 03:00
4	43.5	36.8	1.0	20100101 04:00
5	43.0	36.5	1.0	20100101 05:00
6	43.1	36.3	1.0	20100101 06:00
7	42.3	35.9	1.0	20100101 07:00
8	42.5	36.2	1.0	20100101 08:00
9	45.9	37.8	1.0	20100101 09:00



Weather data of Austin, TX from 2011

	Date	Wban	date	Time	StationType	...	relative_humidity	wind_speed	wind_direction	station_pressure	sea_level_pressure
0	2011-01-01 00:53:00	13904	20110101	5300	12	...	24.0	15.0	360	29.42	29.95
1	2011-01-01 01:53:00	13904	20110101	15300	12	...	23.0	10.0	340	29.49	30.01
2	2011-01-01 02:53:00	13904	20110101	25300	12	...	22.0	15.0	010	29.49	30.01
3	2011-01-01 03:53:00	13904	20110101	35300	12	...	27.0	7.0	350	29.51	30.03
4	2011-01-01 04:53:00	13904	20110101	45300	12	...	25.0	11.0	020	29.51	30.04
5	2011-01-01 05:53:00	13904	20110101	55300	12	...	28.0	6.0	010	29.53	30.06
6	2011-01-01 06:53:00	13904	20110101	65300	12	...	29.0	7.0	360	29.57	30.10
7	2011-01-01 07:53:00	13904	20110101	75300	12	...	29.0	11.0	020	29.59	30.12
8	2011-01-01 08:53:00	13904	20110101	85300	12	...	25.0	15.0	020	29.62	30.16
9	2011-01-01 09:53:00	13904	20110101	95300	12	...	22.0	18.0	010	29.65	30.19

Reminder: read_csv()

- Useful keyword options
 - names: assigning column labels
 - index_col: assigning index
 - parse_dates: parsing datetimes
 - na_values: parsing NaNs



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Statistical exploratory data analysis



Reminder: time series

- Index selection by date time
- Partial datetime selection
- Slicing ranges of datetimes

```
In [1]: climate2010['2010-05-31 22:00:00'] # datetime
```

```
In [2]: climate2010['2010-06-01'] # Entire day
```

```
In [3]: climate2010['2010-04'] # Entire month
```

```
In [4]: climate2010['2010-09':'2010-10'] # 2 months
```

Reminder: statistics methods

- Methods for computing statistics:
 - `describe()`: summary
 - `mean()`: average
 - `count()`: counting entries
 - `median()`: median
 - `std()`: standard deviation



PANDAS FOUNDATIONS

Let's practice!




PANDAS FOUNDATIONS

Visual exploratory data analysis



Line plots in pandas






Line plots in pandas

```
In [1]: import matplotlib.pyplot as plt  
  
In [2]: climate2010.Temperature['2010-07'].plot()  
  
In [3]: plt.title('Temperature (July 2010)')  
  
In [4]: plt.show()
```



Histograms in pandas





Histograms in pandas


```
In [5]: climate2010['DewPoint'].plot(kind= 'hist', bins=30)
```

```
In [6]: plt.title('Dew Point distribution (2010)')
```

```
In [7]: plt.show()
```



Box plots in pandas





Box plots in pandas


```
In [8]: climate2010['DewPoint'].plot(kind='box')
```

```
In [9]: plt.title('Dew Point distribution (2010)')
```

```
In [10]: plt.show()
```



Subplots in pandas





Subplots in pandas

```
In [11]: climate2010.plot(kind='hist', normed=True, subplots=True)  
  
In [12]: plt.show()
```



PANDAS FOUNDATIONS

Let's practice!



PANDAS FOUNDATIONS

Final thoughts



You can now...

- Import many types of datasets and deal with import issues
- Export data to facilitate collaborative data science
- Perform statistical and visual EDA natively in pandas



PANDAS FOUNDATIONS

**See you in the
next course!**

```
In [1]: dict = {
"country":["Brazil", "Russia", "India", "China", "South Africa"],
"capital":["Brasilia", "Moscow", "New Delhi", "Beijing", "Pretoria"],
"area":[8.516, 17.10, 3.286, 9.597, 1.221],
"population":[200.4, 143.5, 1252, 1357, 52.98] }
print(dict)
```

```
{'country': ['Brazil', 'Russia', 'India', 'China', 'South Africa'], 'capital': ['Brasilia', 'Moscow', 'New Delhi', 'Beijing', 'Pretoria'], 'area': [8.516, 17.1, 3.286, 9.597, 1.221], 'population': [200.4, 143.5, 1252, 1357, 52.98]}
```

```
In [2]: import pandas as pd
x=pd.DataFrame(dict)
print(x)
```

	country	capital	area	population
0	Brazil	Brasilia	8.516	200.40
1	Russia	Moscow	17.100	143.50
2	India	New Delhi	3.286	1252.00
3	China	Beijing	9.597	1357.00
4	South Africa	Pretoria	1.221	52.98

```
In [3]: x.index = ["BR", "RU", "IN", "CH", "SA"]
print(x)
```

	country	capital	area	population
BR	Brazil	Brasilia	8.516	200.40
RU	Russia	Moscow	17.100	143.50
IN	India	New Delhi	3.286	1252.00
CH	China	Beijing	9.597	1357.00
SA	South Africa	Pretoria	1.221	52.98

```
In [44]: x= pd.read_csv("H:\std.csv")
print(x)
```

	rollno	sub1	sub2	sub3
0	90	90	95.0	100
1	91	98	94.0	63
2	92	63	NaN	47
3	93	95	94.0	75
4	94	65	NaN	69
5	95	82	89.0	85

```
In [5]: x['country']
```

```
Out[5]: BR      Brazil
RU      Russia
IN      India
CH      China
SA      South Africa
Name: country, dtype: object
```

```
In [6]: type(x["country"])
```

```
Out[6]: pandas.core.series.Series
```

Column Access []

In [7]: `x[['country']]`

Out[7]:

	country
BR	Brazil
RU	Russia
IN	India
CH	China
SA	South Africa

In [8]: `x[["country", "capital"]]`

Out[8]:

	country	capital
BR	Brazil	Brasilia
RU	Russia	Moscow
IN	India	New Delhi
CH	China	Beijing
SA	South Africa	Pretoria

Row Access

In [33]: `x.iloc[0:2,0:2]`

Out[33]:

	country	capital
BR	Brazil	Brasilia
RU	Russia	Moscow

loc - (label-based)

iloc- (integer position-based)

In [10]: `#Row as Pandas Series
print(x.loc['RU'])
#DataFrame
print(x.loc[['RU']])`

country	Russia			
capital	Moscow			
area	17.1			
population	143.5			
Name: RU, dtype: object				
country	capital	area	population	
RU	Russia	Moscow	17.1	143.5

In [11]:

x.loc[["RU", "IN", "CH"]]

Out[11]:

	country	capital	area	population
RU	Russia	Moscow	17.100	143.5
IN	India	New Delhi	3.286	1252.0
CH	China	Beijing	9.597	1357.0

In [12]:

x.loc[["RU", "IN", "CH"], ["country", "capital"]]

Out[12]:

	country	capital
RU	Russia	Moscow
IN	India	New Delhi
CH	China	Beijing

In [13]:

x.loc[:, ['country', 'capital']]

Out[13]:

	country	capital
BR	Brazil	Brasilia
RU	Russia	Moscow
IN	India	New Delhi
CH	China	Beijing
SA	South Africa	Pretoria

In [14]:

x.iloc[[0]]

Out[14]:

	country	capital	area	population
BR	Brazil	Brasilia	8.516	200.4

In [15]:

x.iloc[[1,2,3]]

Out[15]:

	country	capital	area	population
RU	Russia	Moscow	17.100	143.5
IN	India	New Delhi	3.286	1252.0
CH	China	Beijing	9.597	1357.0

In [16]:

x.iloc[:, [0,2]]

Out[16]:

	country	area
BR	Brazil	8.516
RU	Russia	17.100
IN	India	3.286
CH	China	9.597
SA	South Africa	1.221

```
In [17]: x[x["area"] > 8]
```

Out[17]:

	country	capital	area	population
BR	Brazil	Brasilia	8.516	200.4
RU	Russia	Moscow	17.100	143.5
CH	China	Beijing	9.597	1357.0

```
In [18]: x[x.iloc[:,2]>8]
```

Out[18]:

	country	capital	area	population
BR	Brazil	Brasilia	8.516	200.4
RU	Russia	Moscow	17.100	143.5
CH	China	Beijing	9.597	1357.0

```
In [34]: import numpy as np
print(x)
```

	country	capital	area	population	newcol	new
BR	Brazil	Brasilia	8.516	200.40	6.0	6
RU	Russia	Moscow	17.100	143.50	6.0	6
IN	India	New Delhi	3.286	1252.00	5.0	5
CH	China	Beijing	9.597	1357.00	5.0	5
SA	South Africa	Pretoria	1.221	52.98	12.0	12

```
In [20]: x[np.logical_and(x["area"] > 8, x["area"] < 10)]
```

Out[20]:

	country	capital	area	population
BR	Brazil	Brasilia	8.516	200.4
CH	China	Beijing	9.597	1357.0

```
In [21]: for i in x:
print(i)
```

country
capital
area
population

```
In [22]: for l,r in x.iterrows():
          print(l)
          print(r)
```

```
BR
country      Brazil
capital      Brasilia
area         8.516
population   200.4
Name: BR, dtype: object
```

```
RU
country      Russia
capital      Moscow
area         17.1
population   143.5
Name: RU, dtype: object
```

```
IN
country      India
capital      New Delhi
area         3.286
population   1252
Name: IN, dtype: object
```

```
CH
country      China
capital      Beijing
area         9.597
population   1357
Name: CH, dtype: object
```

```
SA
country      South Africa
capital      Pretoria
area         1.221
population   52.98
Name: SA, dtype: object
```

```
In [23]: for l,r in x.iterrows():
          print(l,":",r['country'])
```

```
BR : Brazil
RU : Russia
IN : India
CH : China
SA : South Africa
```

```
In [46]: for l,r in x.iterrows():
          x.loc[l,"totalmarks"]=r['sub1']+r['sub2']+r['sub3']
```

```
In [47]: print(x)
```

	rollno	sub1	sub2	sub3	totalmarks
0	90	90	95.0	100	285.0
1	91	98	94.0	63	255.0
2	92	63	NaN	47	NaN
3	93	95	94.0	75	264.0
4	94	65	NaN	69	NaN
5	95	82	89.0	85	256.0

```
In [26]: x['new']=x['country'].apply(len)
```


In [27]:

x

Out[27]:

	country	capital	area	population	newcol	new
BR	Brazil	Brasilia	8.516	200.40	6.0	6
RU	Russia	Moscow	17.100	143.50	6.0	6
IN	India	New Delhi	3.286	1252.00	5.0	5
CH	China	Beijing	9.597	1357.00	5.0	5
SA	South Africa	Pretoria	1.221	52.98	12.0	12

In [28]:

x.loc[['RU','IN','CH']]

Out[28]:

	country	capital	area	population	newcol	new
RU	Russia	Moscow	17.100	143.5	6.0	6
IN	India	New Delhi	3.286	1252.0	5.0	5
CH	China	Beijing	9.597	1357.0	5.0	5

In [45]:

x.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6 entries, 0 to 5
Data columns (total 4 columns):
rollno 6 non-null int64
sub1 6 non-null int64
sub2 4 non-null float64
sub3 6 non-null int64
dtypes: float64(1), int64(3)
memory usage: 272.0 bytes

In [48]:

x.to_csv("abc.csv")

In [49]:

x.to_excel("dat.xlsx")

In [32]:

x.iloc[:3,1:3]

Out[32]:

	capital	area
BR	Brasilia	8.516
RU	Moscow	17.100
IN	New Delhi	3.286

In []:

```
In [5]: import pandas as pd
data=pd.read_csv('C:\\Users\\iraga\\Desktop\\python\\iris.csv')
print(data.tail())
```

	sepal.length	sepal.width	petal.length	petal.width	variety
145	6.7	3.0	5.2	2.3	Virginica
146	6.3	2.5	5.0	1.9	Virginica
147	6.5	3.0	5.2	2.0	Virginica
148	6.2	3.4	5.4	2.3	Virginica
149	5.9	3.0	5.1	1.8	Virginica

```
In [7]: edata=pd.ExcelFile('C:\\Users\\iraga\\Desktop\\python\\logins.xlsx')
print(edata.sheet_names)

['sheet', 'sheet1']
```

```
In [10]: df1=edata.parse(0)
print(df1.head(20))
```

	username	password	email
0	CHDFAC001	Infy@123	kamravikas@akgec.ac.in
1	CHDFAC002	Infy@123	rawatgaaurav@akgec.ac.in
2	CHDFAC003	Infy@123	seema.baghae@gmail.com
3	CHDFAC004	Infy@123	tannuchanana1990@gmail.com
4	CHDFAC005	Infy@123	sivajyothi.cse@anits.edu.in
5	CHDFAC006	Infy@123	meena.it@anits.edu.in
6	CHDFAC007	Infy@123	mnprasadu@gmail.com
7	CHDFAC008	Infy@123	paramesh.u@gmail.com
8	CHDFAC009	Infy@123	suvarnashirke@atharvacoe.ac.in
9	CHDFAC010	Infy@123	nidaparkar@atharvacoe.ac.in
10	CHDFAC011	Infy@123	geethu0517@gmail.com
11	CHDFAC012	Infy@123	ravindra.meegada99@gmail.com
12	CHDFAC013	Infy@123	promila.verma@baddiuniv.ac.in
13	CHDFAC014	Infy@123	priyanka.sharma@baddiuniv.ac.in
14	CHDFAC015	Infy@123	prashanthbabubandi@gmail.com
15	CHDFAC016	Infy@123	prasanth.k@becbapatla.ac.in
16	CHDFAC017	Infy@123	pardeep308@yahoo.co.in
17	CHDFAC018	Infy@123	ashiarya@gmail.com
18	CHDFAC019	Infy@123	vasanthasena_it@cbit.ac.in
19	CHDFAC020	Infy@123	gnyanadeepa_it@cbit.ac.in

```
In [6]: df1=edata.parse(1)
df1.head()
```

Out[6]:

	username	password	email
0	CHDFAC104	Infy@123	varsha469@gmail.com
1	CHDFAC105	Infy@123	jhdevare@mitaoe.ac.in
2	CHDFAC106	Infy@123	stwarpe@comp.maepune.ac.in
3	CHDFAC107	Infy@123	kanwalpreetsingh_pu@yahoo.com
4	CHDFAC108	Infy@123	manishkumar.3006@poornima.org

```
In [4]: df1=edata.parse('sheet')
df1.head()
```

Out[4]:

	username	password	email
0	CHDFAC001	Infy@123	kamravikas@akgec.ac.in
1	CHDFAC002	Infy@123	rawatgurav@akgec.ac.in
2	CHDFAC003	Infy@123	seema.baghae@gmail.com
3	CHDFAC004	Infy@123	tannuchanana1990@gmail.com
4	CHDFAC005	Infy@123	sivajyothi.cse@anits.edu.in

```
In [5]: df1=edata.parse('sheet1')
df1.head()
```

Out[5]:

	username	password	email
0	CHDFAC104	Infy@123	varsha469@gmail.com
1	CHDFAC105	Infy@123	jhdevare@mitaoe.ac.in
2	CHDFAC106	Infy@123	stwarpe@comp.maepune.ac.in
3	CHDFAC107	Infy@123	kanwalpreetsingh_pu@yahoo.com
4	CHDFAC108	Infy@123	manishkumar.3006@poornima.org

In []: